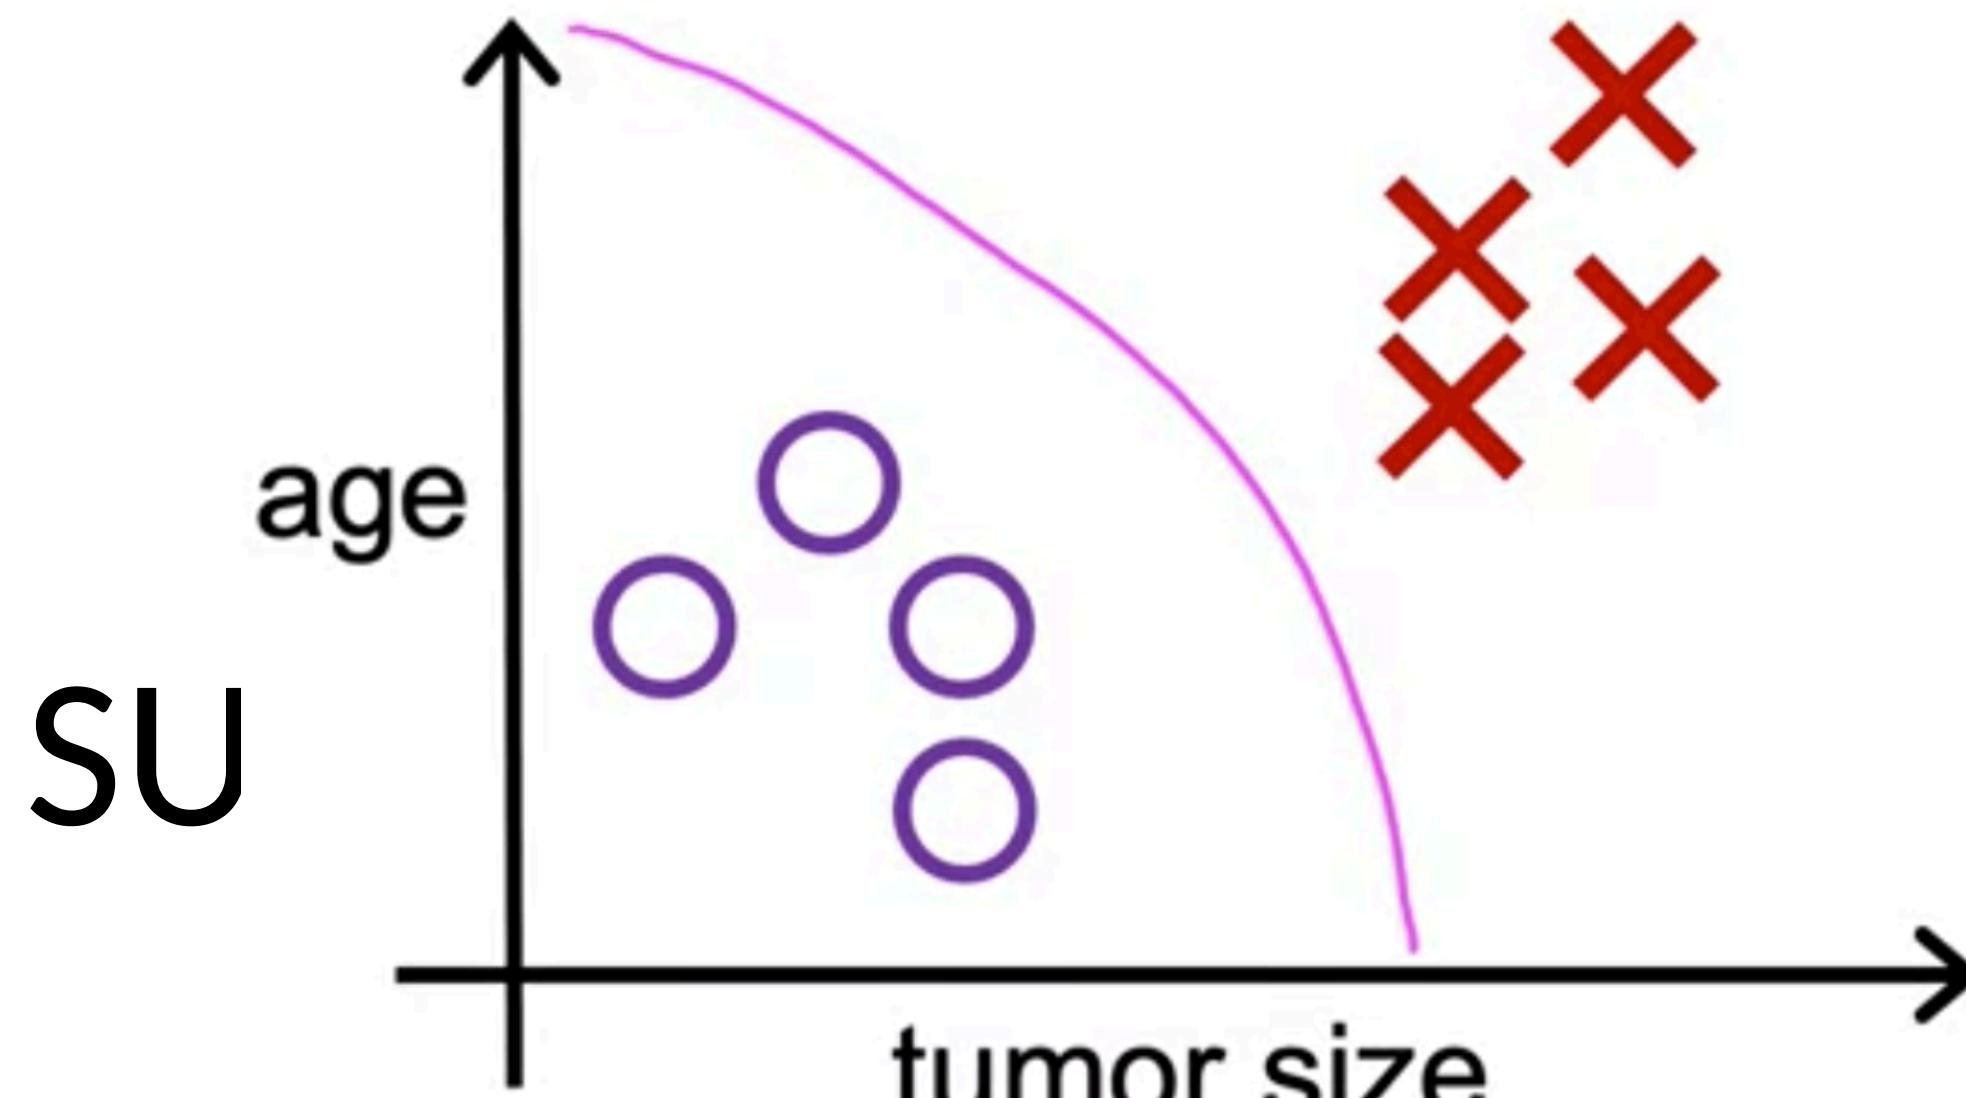


AIBridge

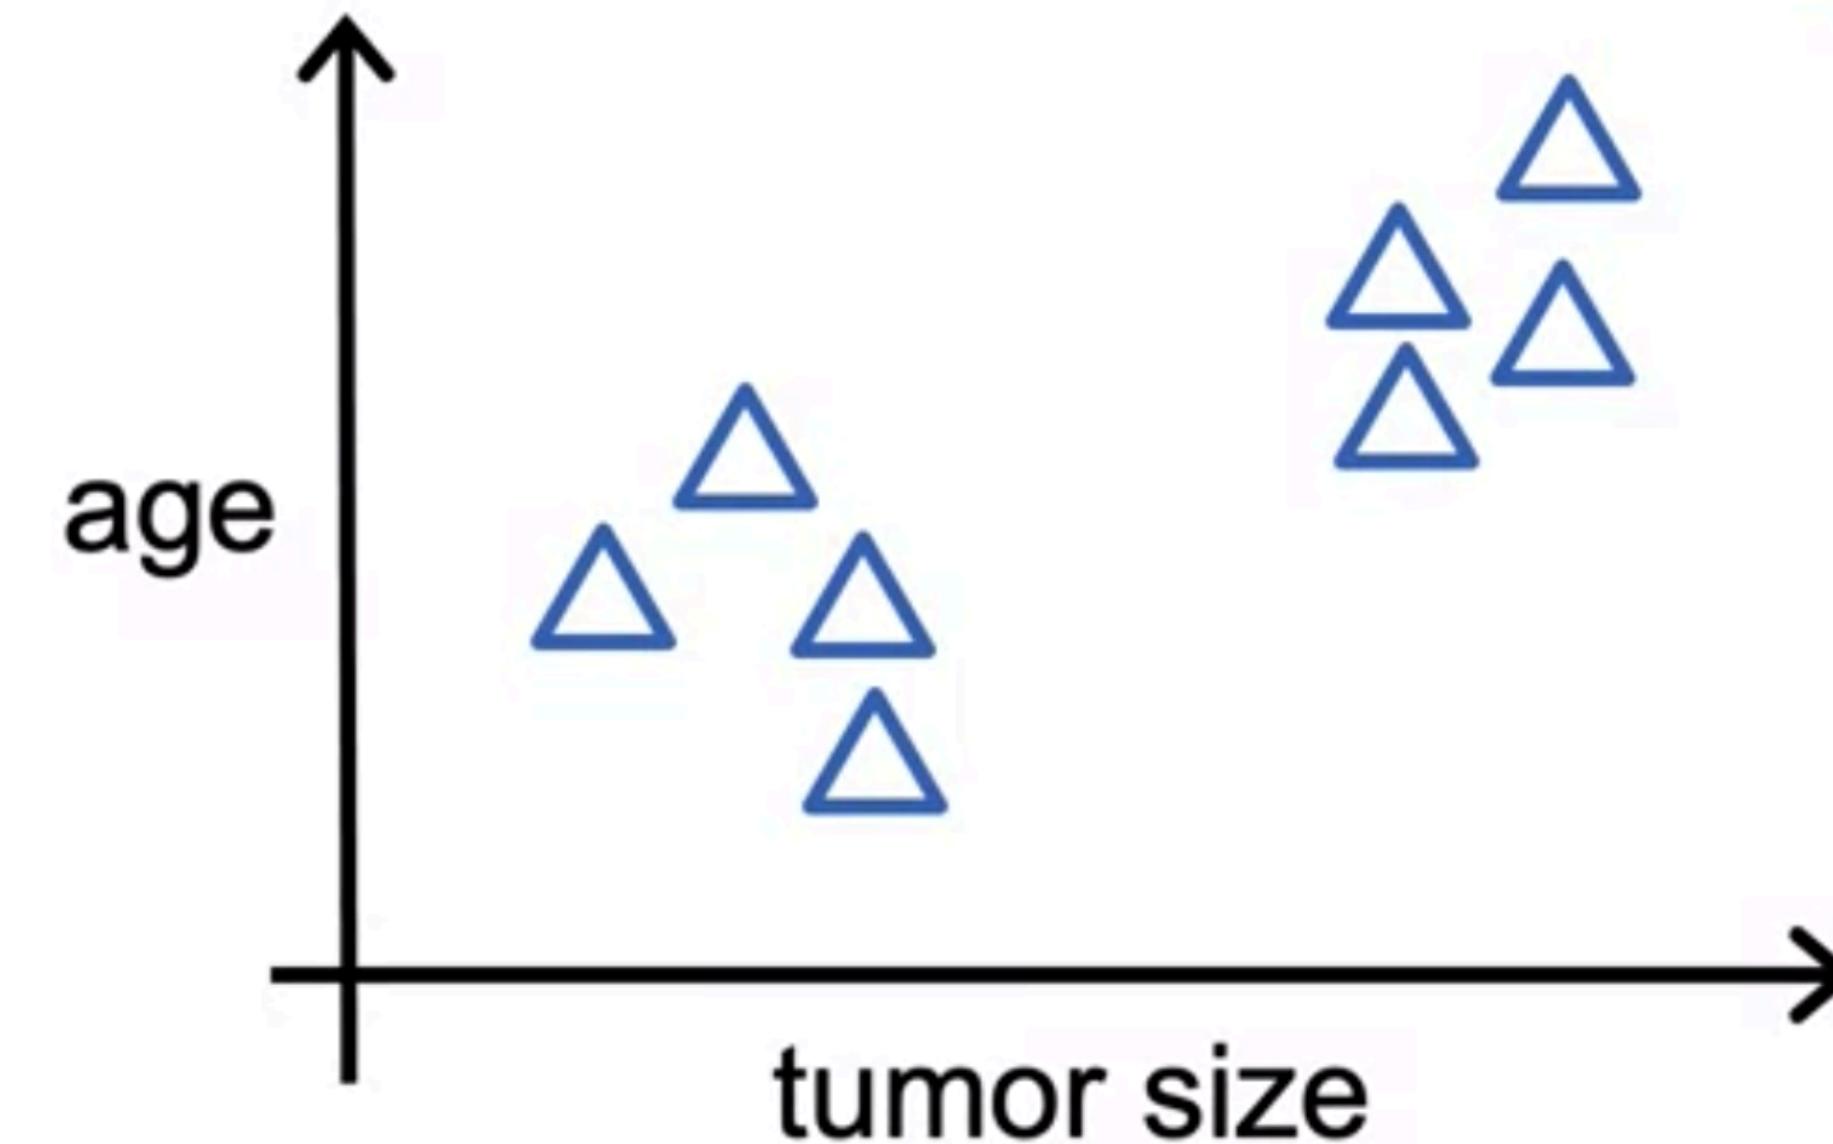
Lecture 8

UNSUPERVISED LEARNING

Supervised learning
Learn from data **labeled**
with the “**right answers**”



Unsupervised learning



TWO MAIN APPLICATIONS

Clustering

Dimension reduction

Clustering: Google news

Giant panda gives birth to rare twin cubs at Japan's oldest zoo

USA TODAY · 6 hours ago



- Giant panda gives birth to twin cubs at Japan's oldest zoo

CBS News · 7 hours ago

- Giant panda gives birth to twin cubs at Tokyo's Ueno Zoo

WHBL News · 16 hours ago

- A Joyful Surprise at Japan's Oldest Zoo: The Birth of Twin Pandas

The New York Times · 1 hour ago

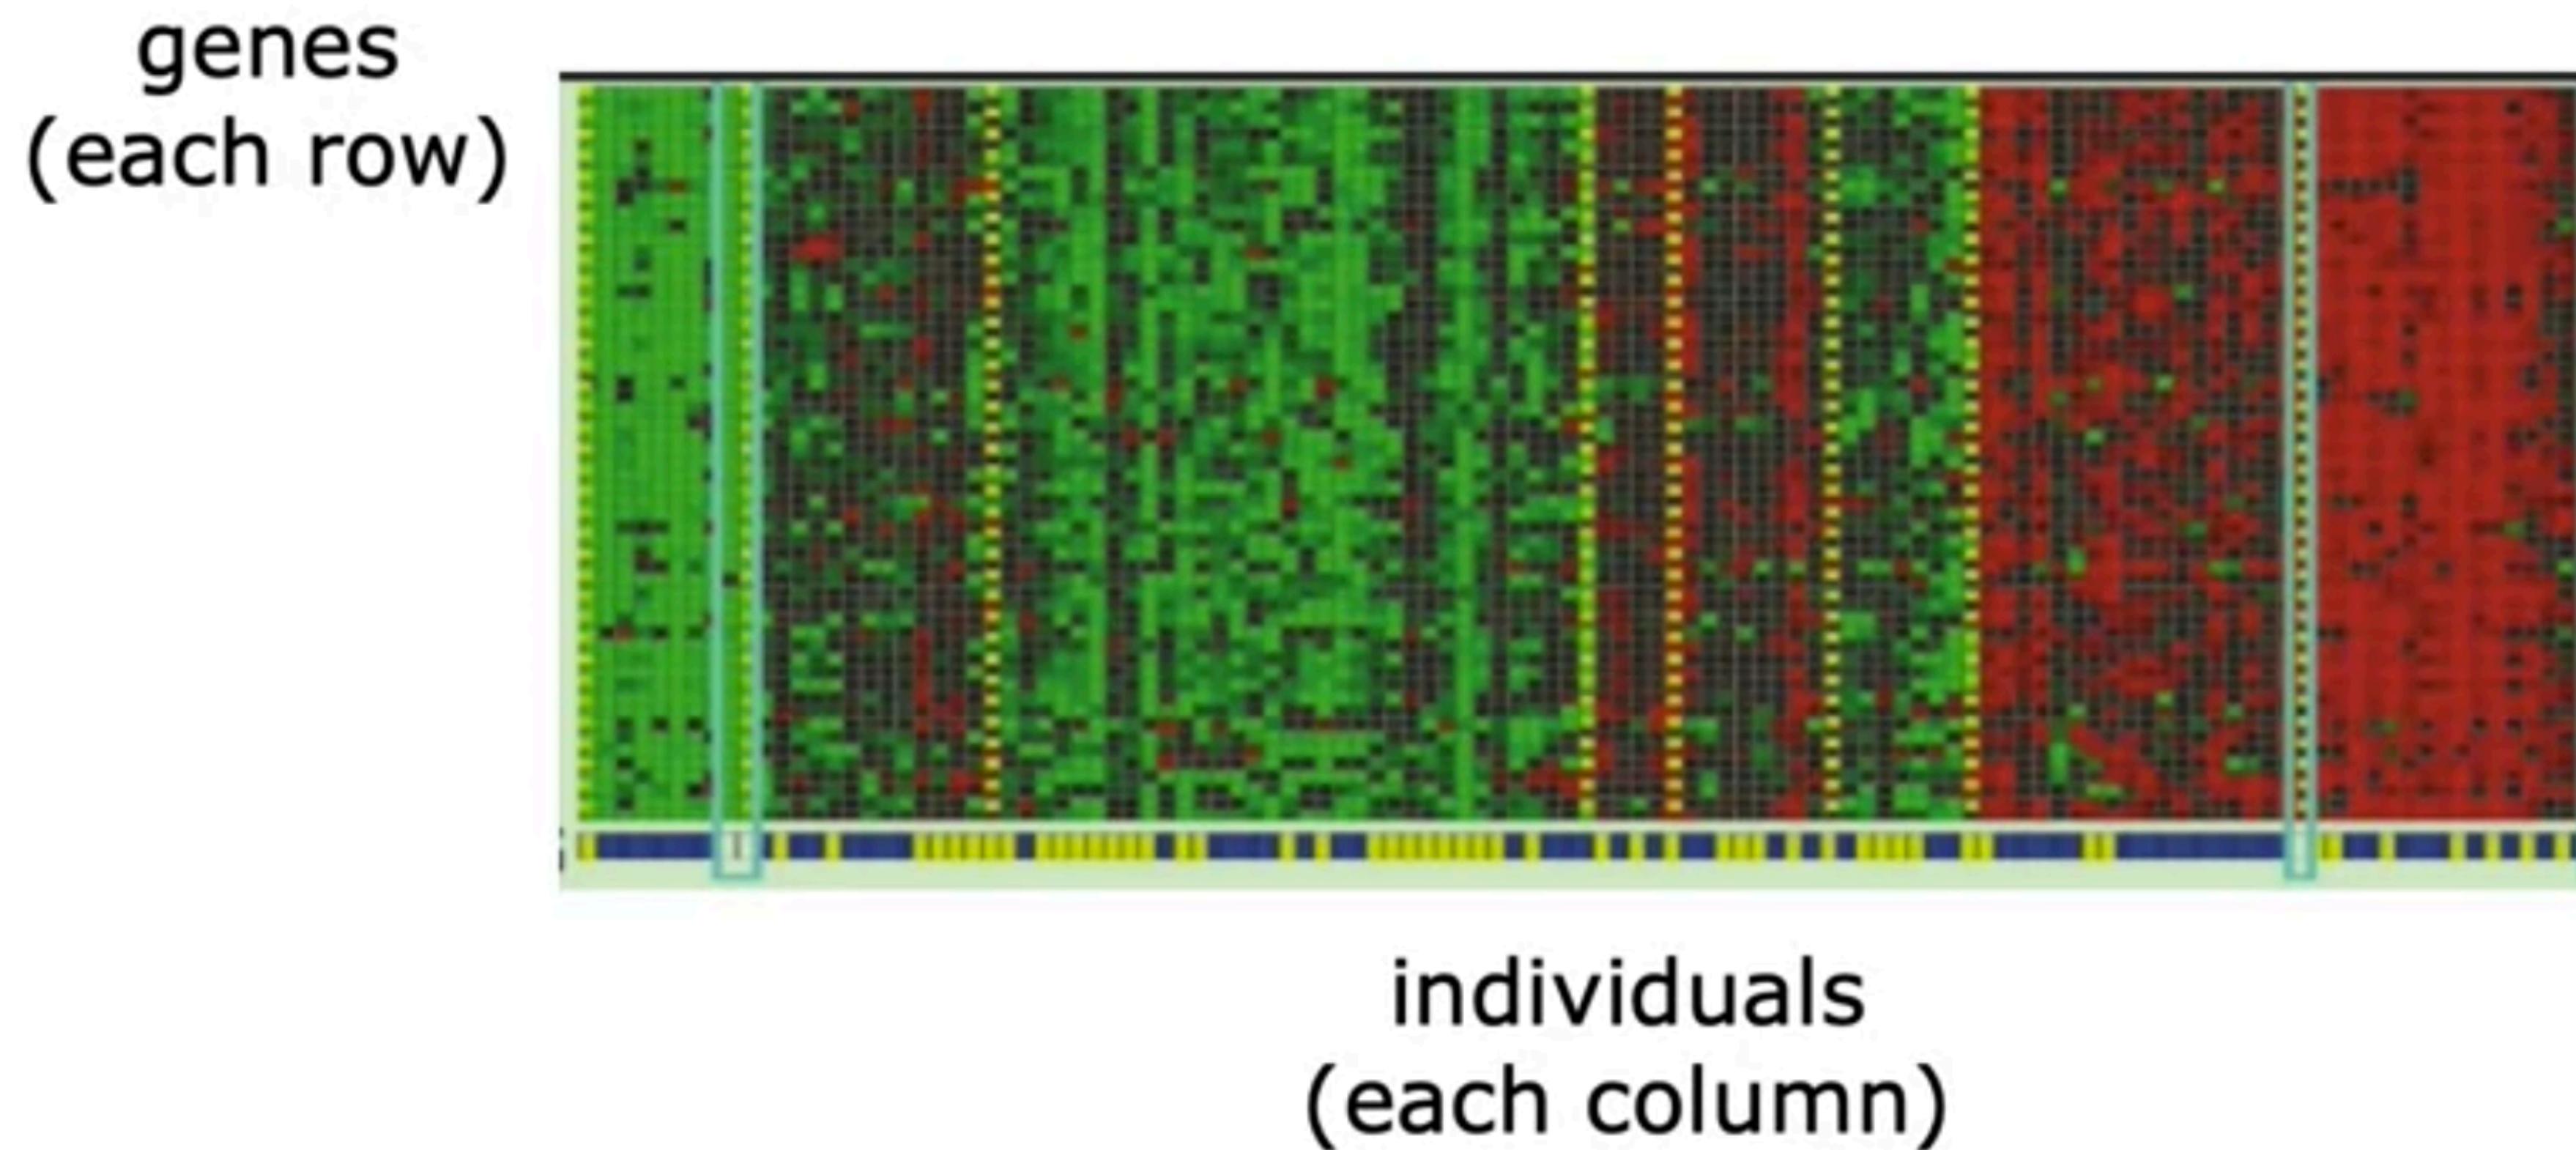
- Twin Panda Cubs Born at Tokyo's Ueno Zoo

PEOPLE · 6 hours ago

 View Full Coverage

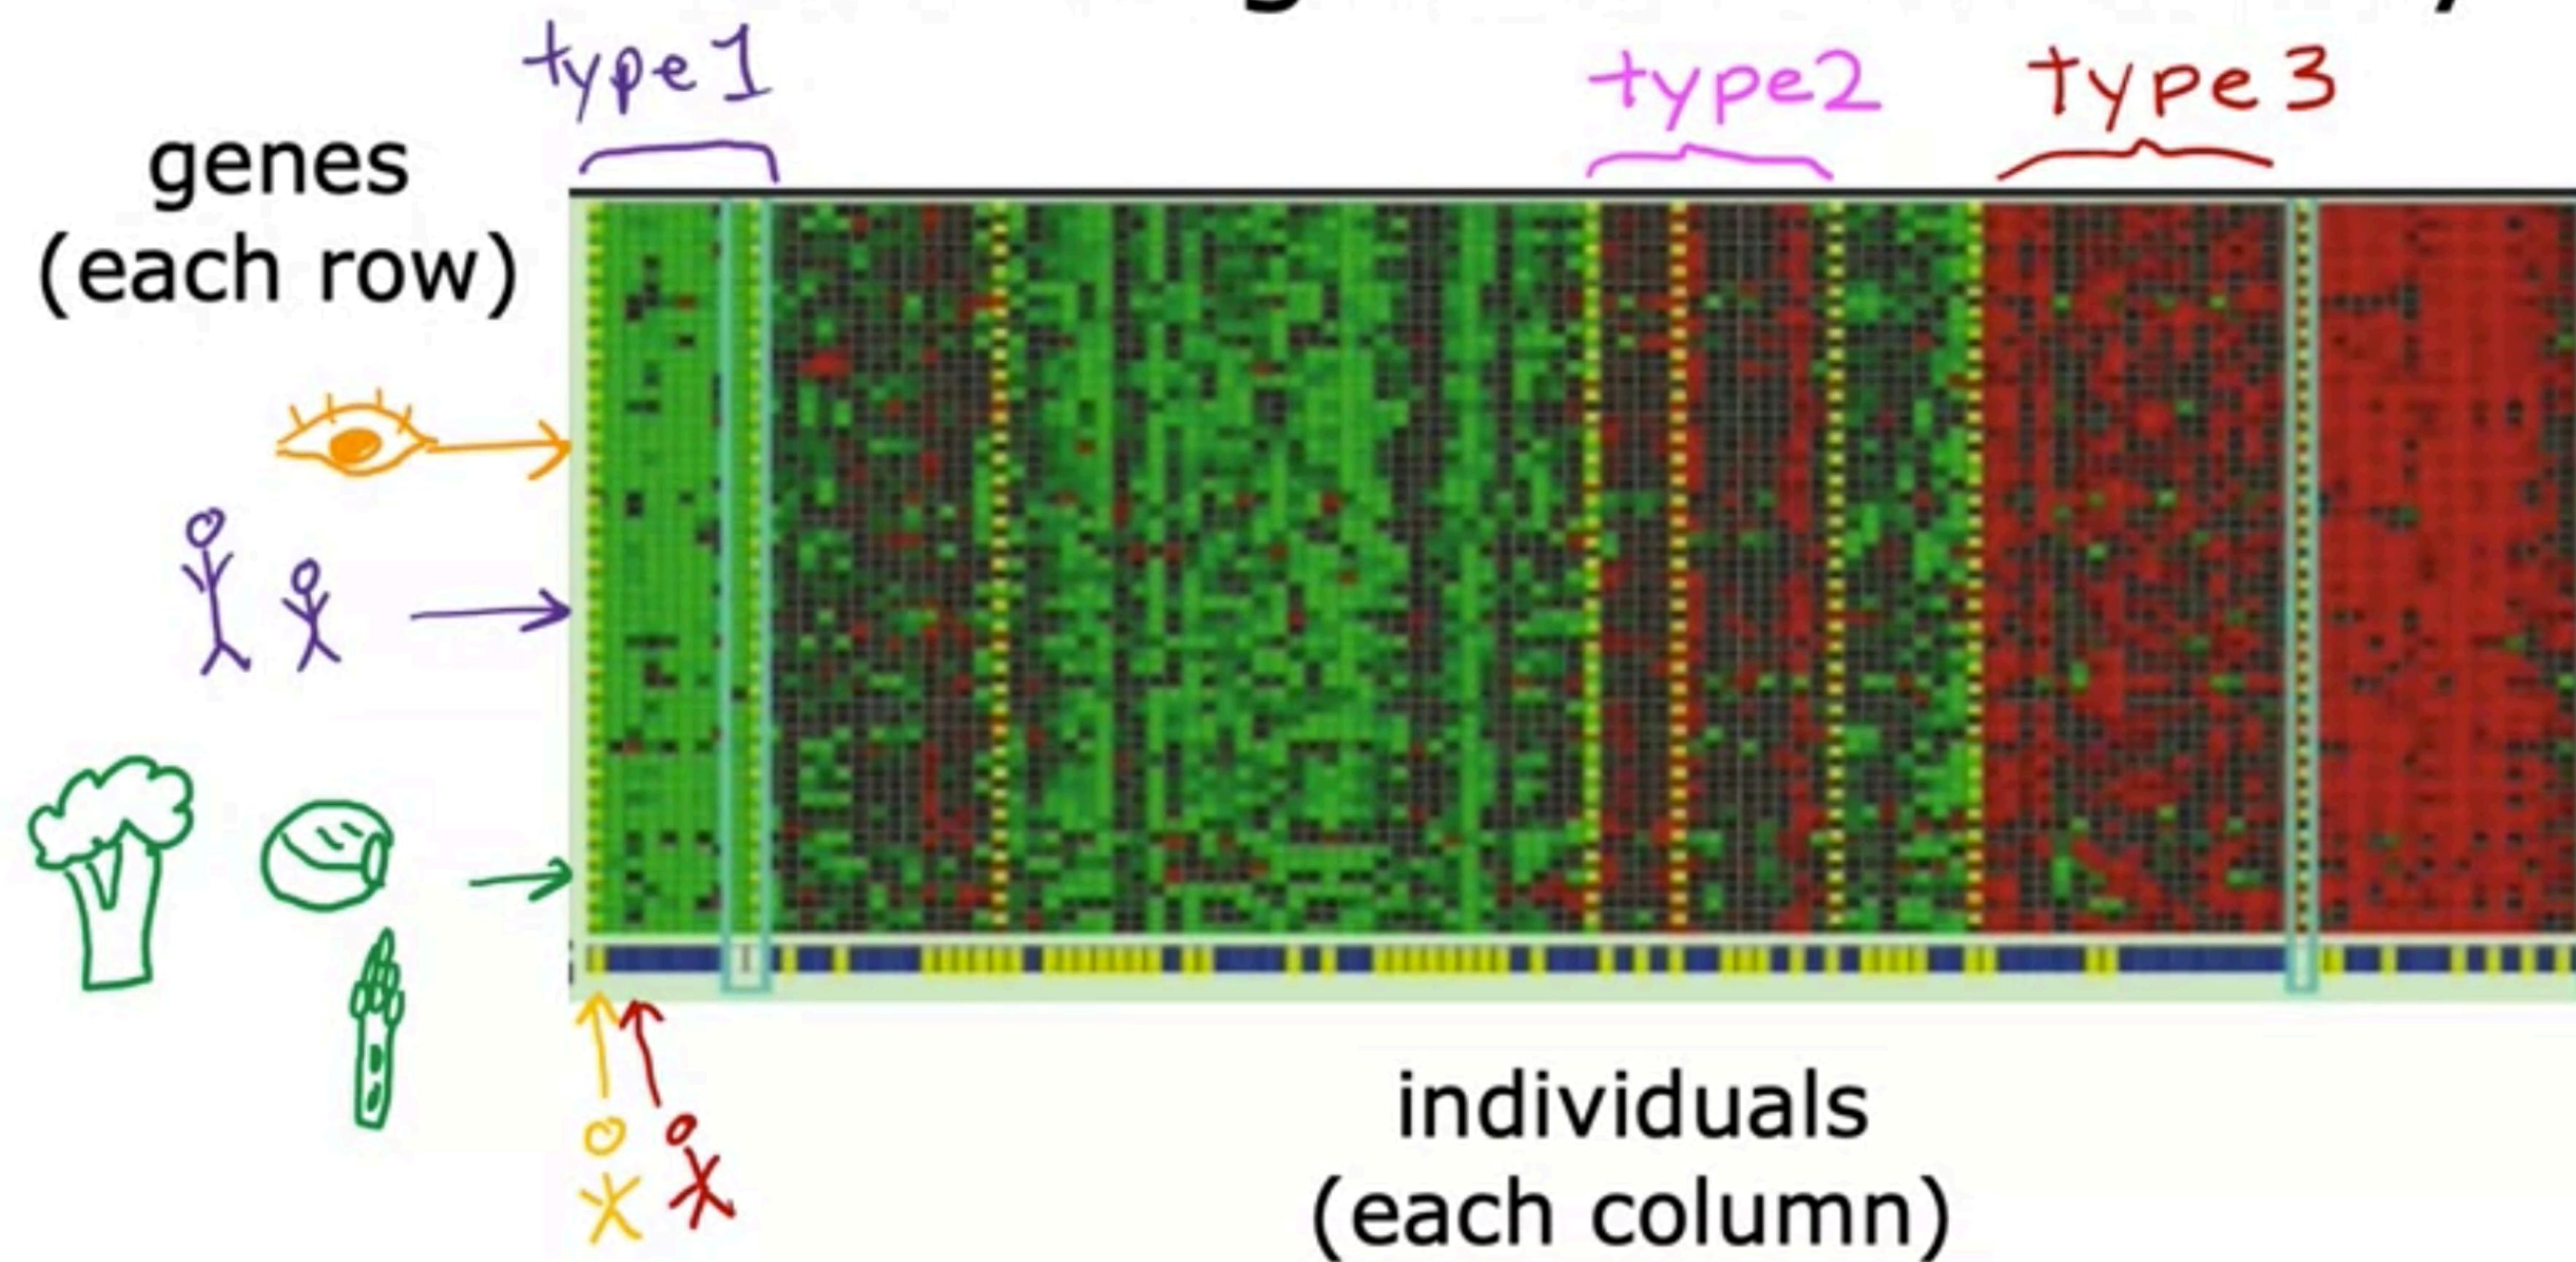


Clustering: DNA microarray



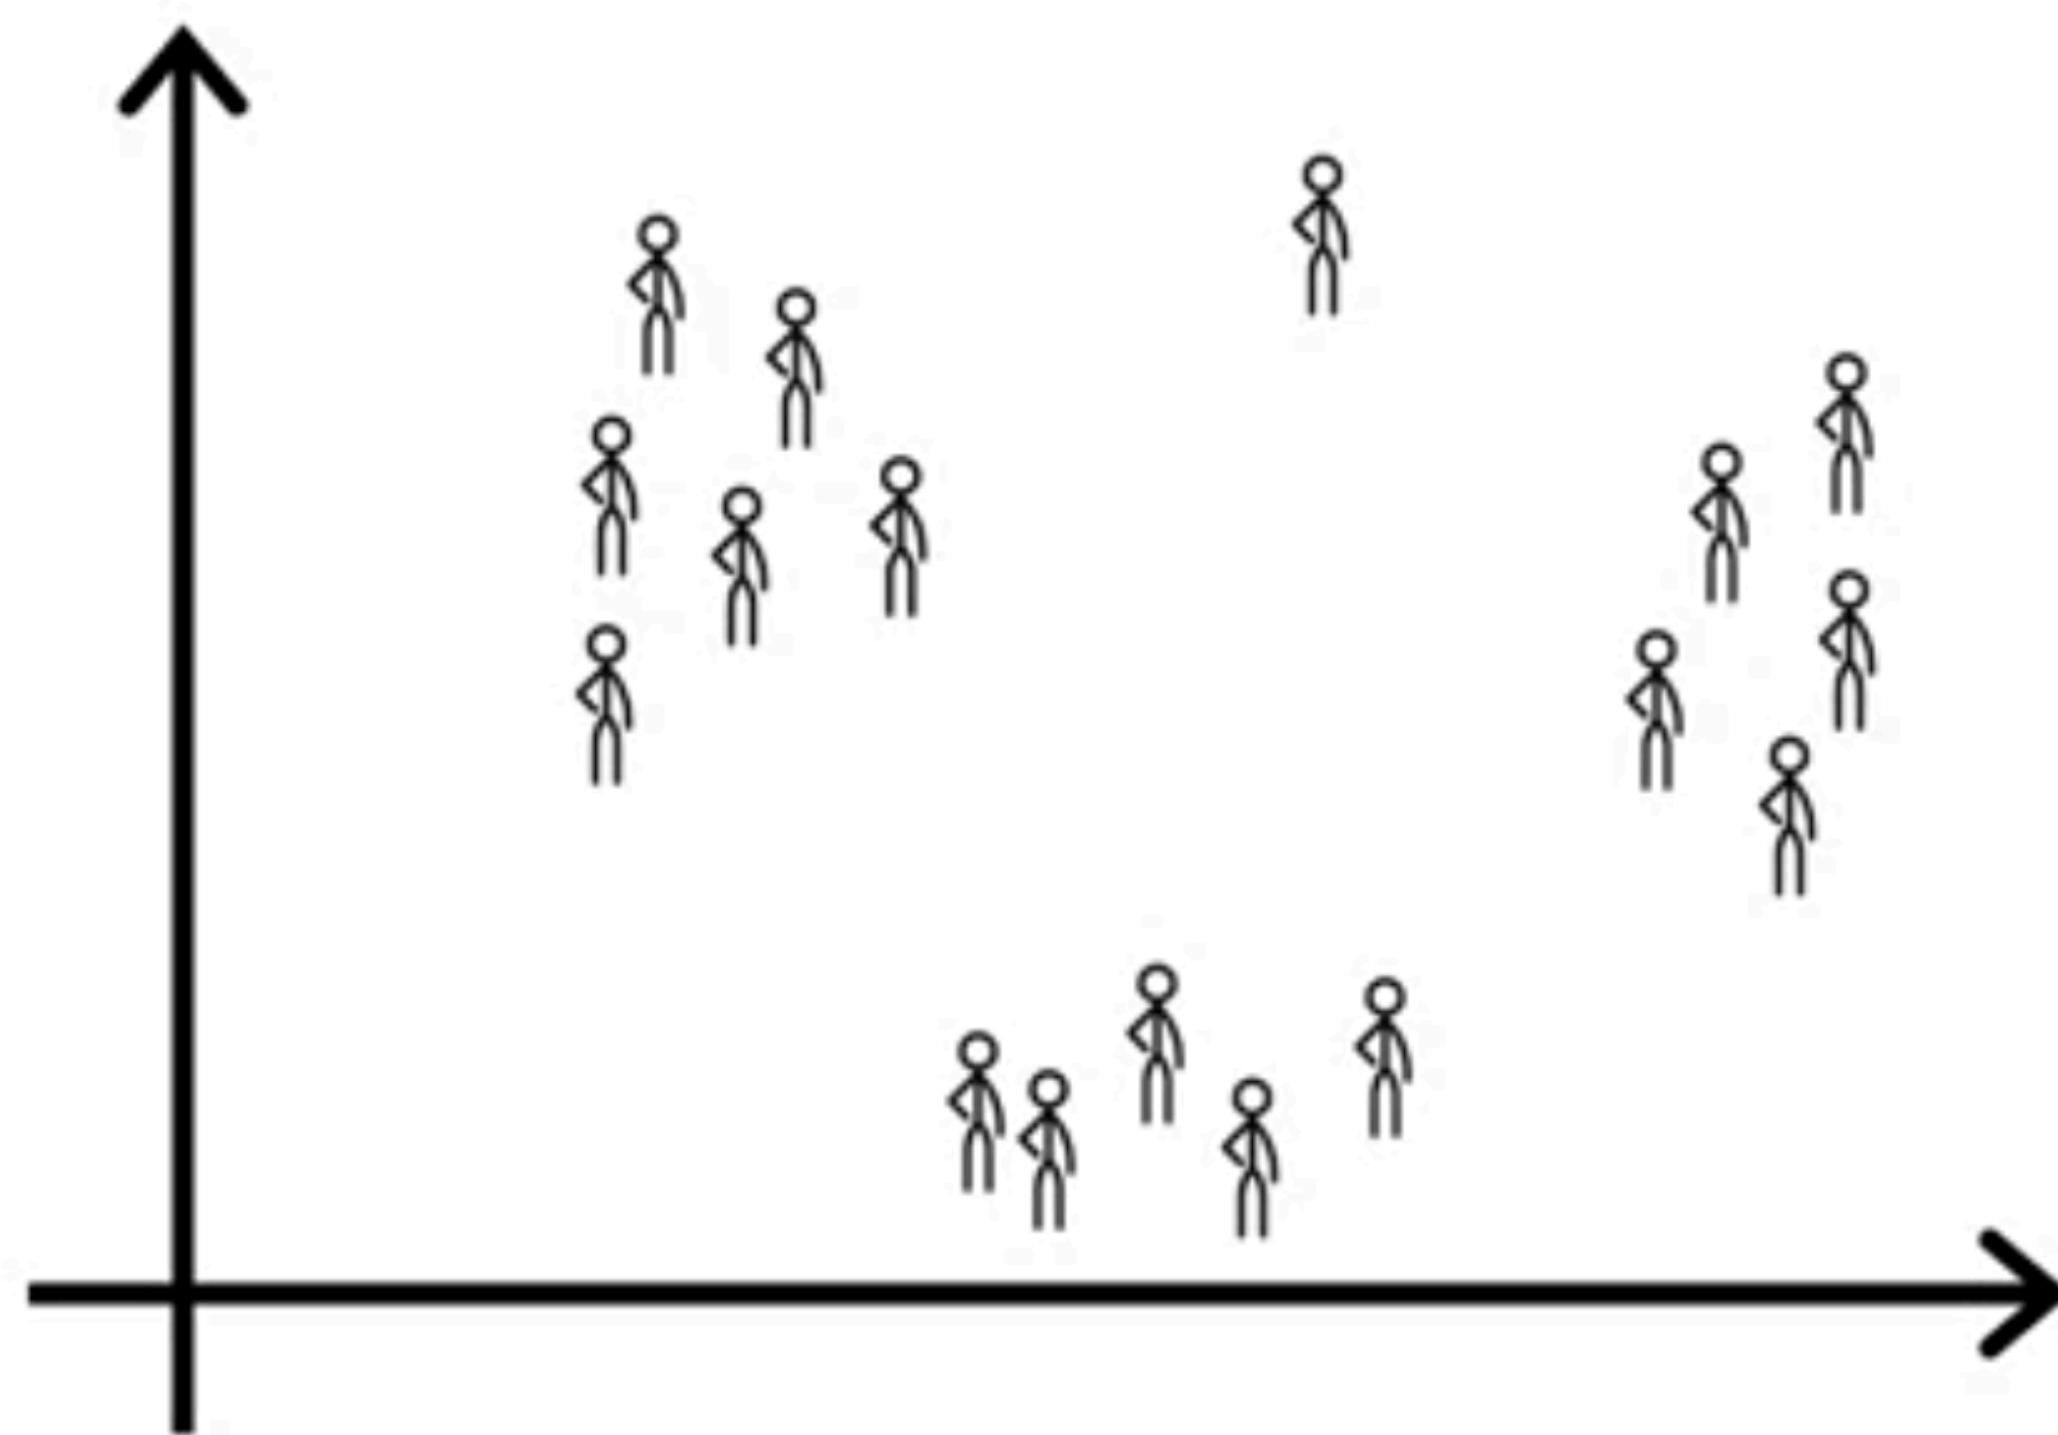
Credit: Andrew Ng, [Machine Learning](#)

Clustering: DNA microarray

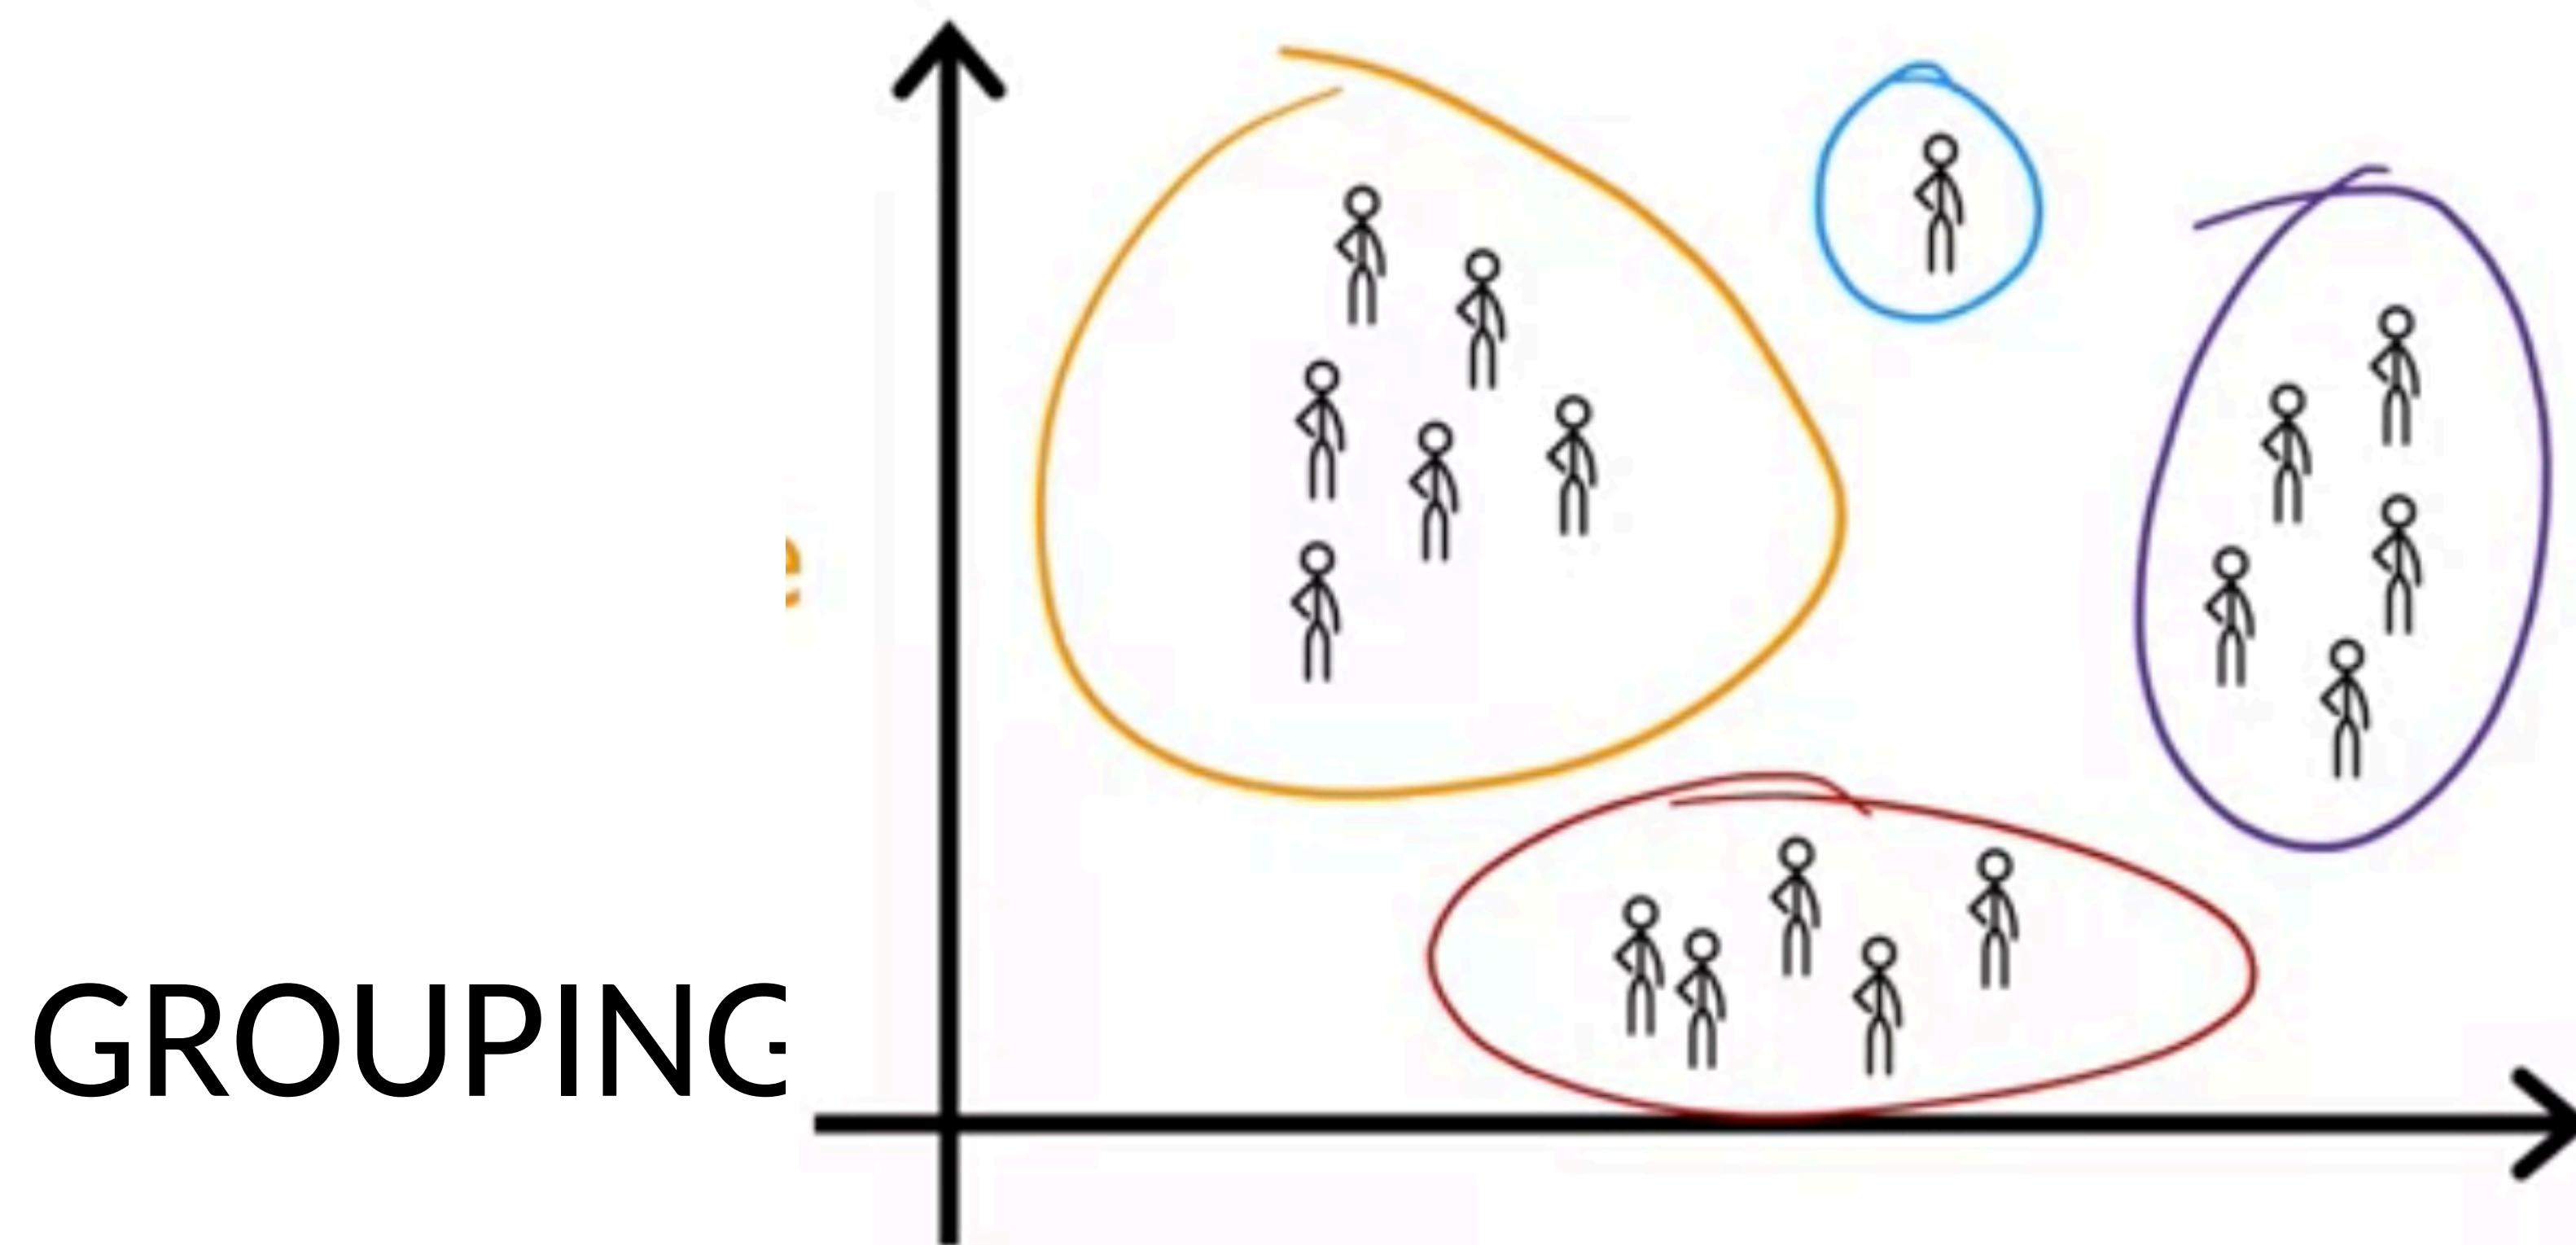


Credit: Andrew Ng, [Machine Learning](#)

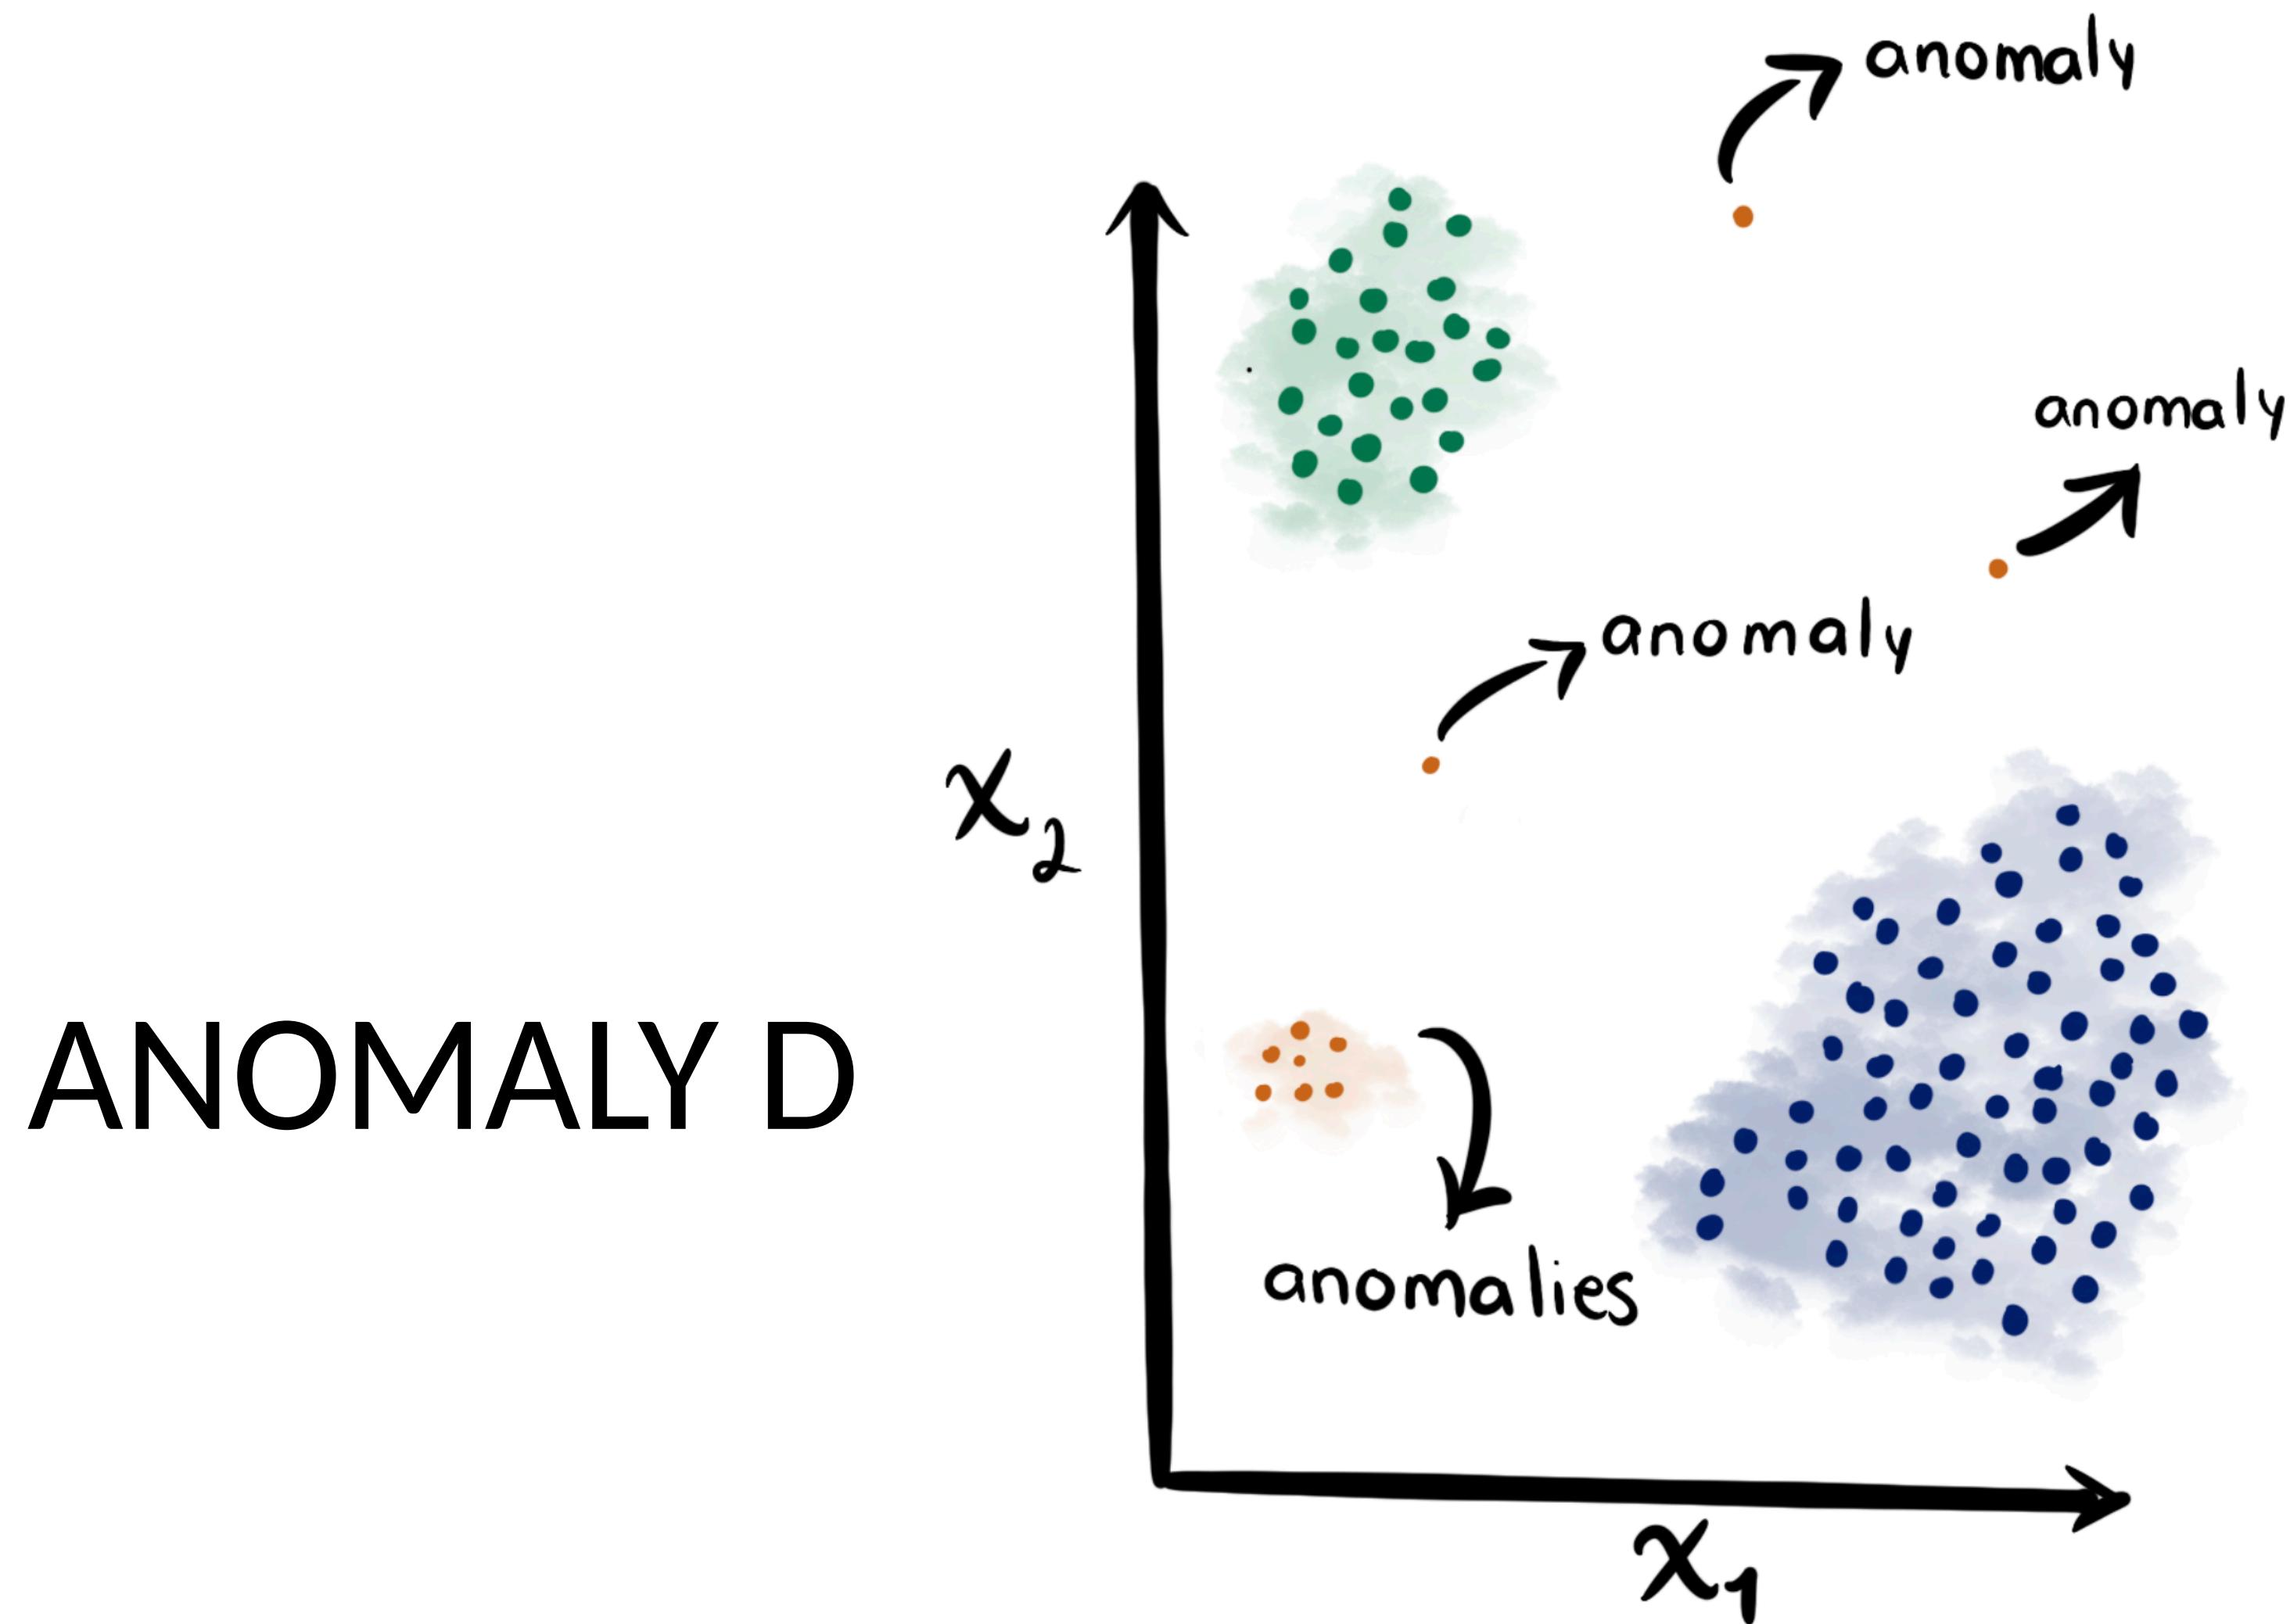
Clustering: Grouping customers



Credit: Andrew Ng, [Machine Learning](#)



Credit: Andrew Ng, [Machine Learning](#)



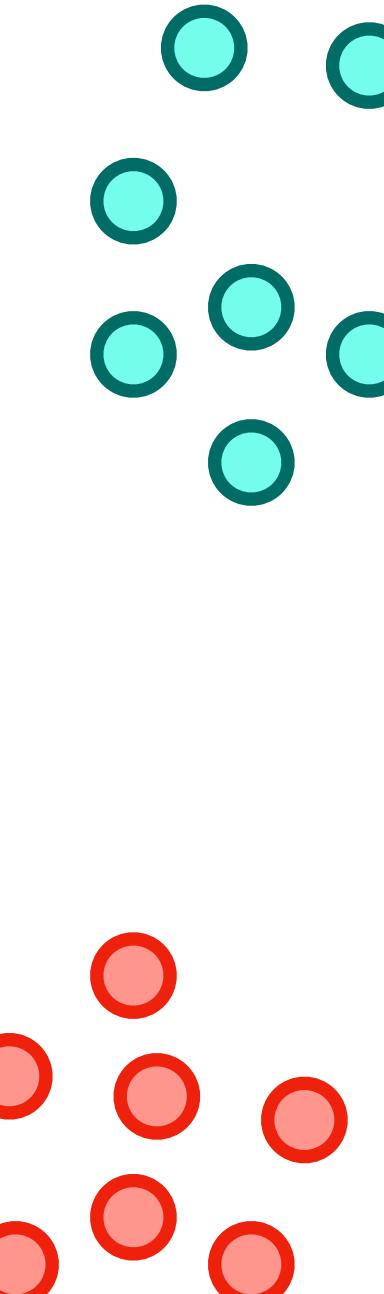
Credit: [Anomaly Detection](#)

how do we find the clusters?

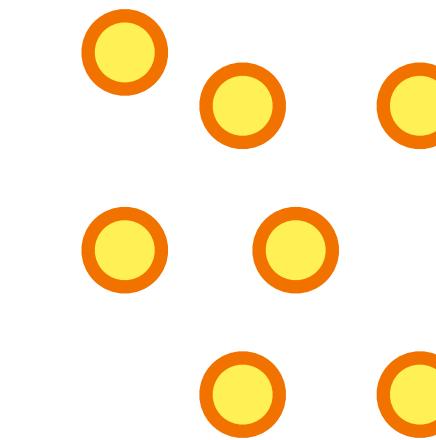


Comfort

super
comfortable

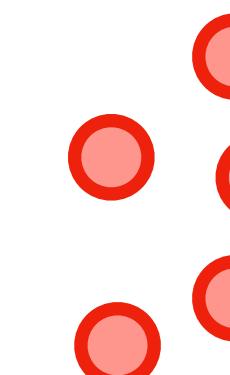


super
fashionable



all-rounder

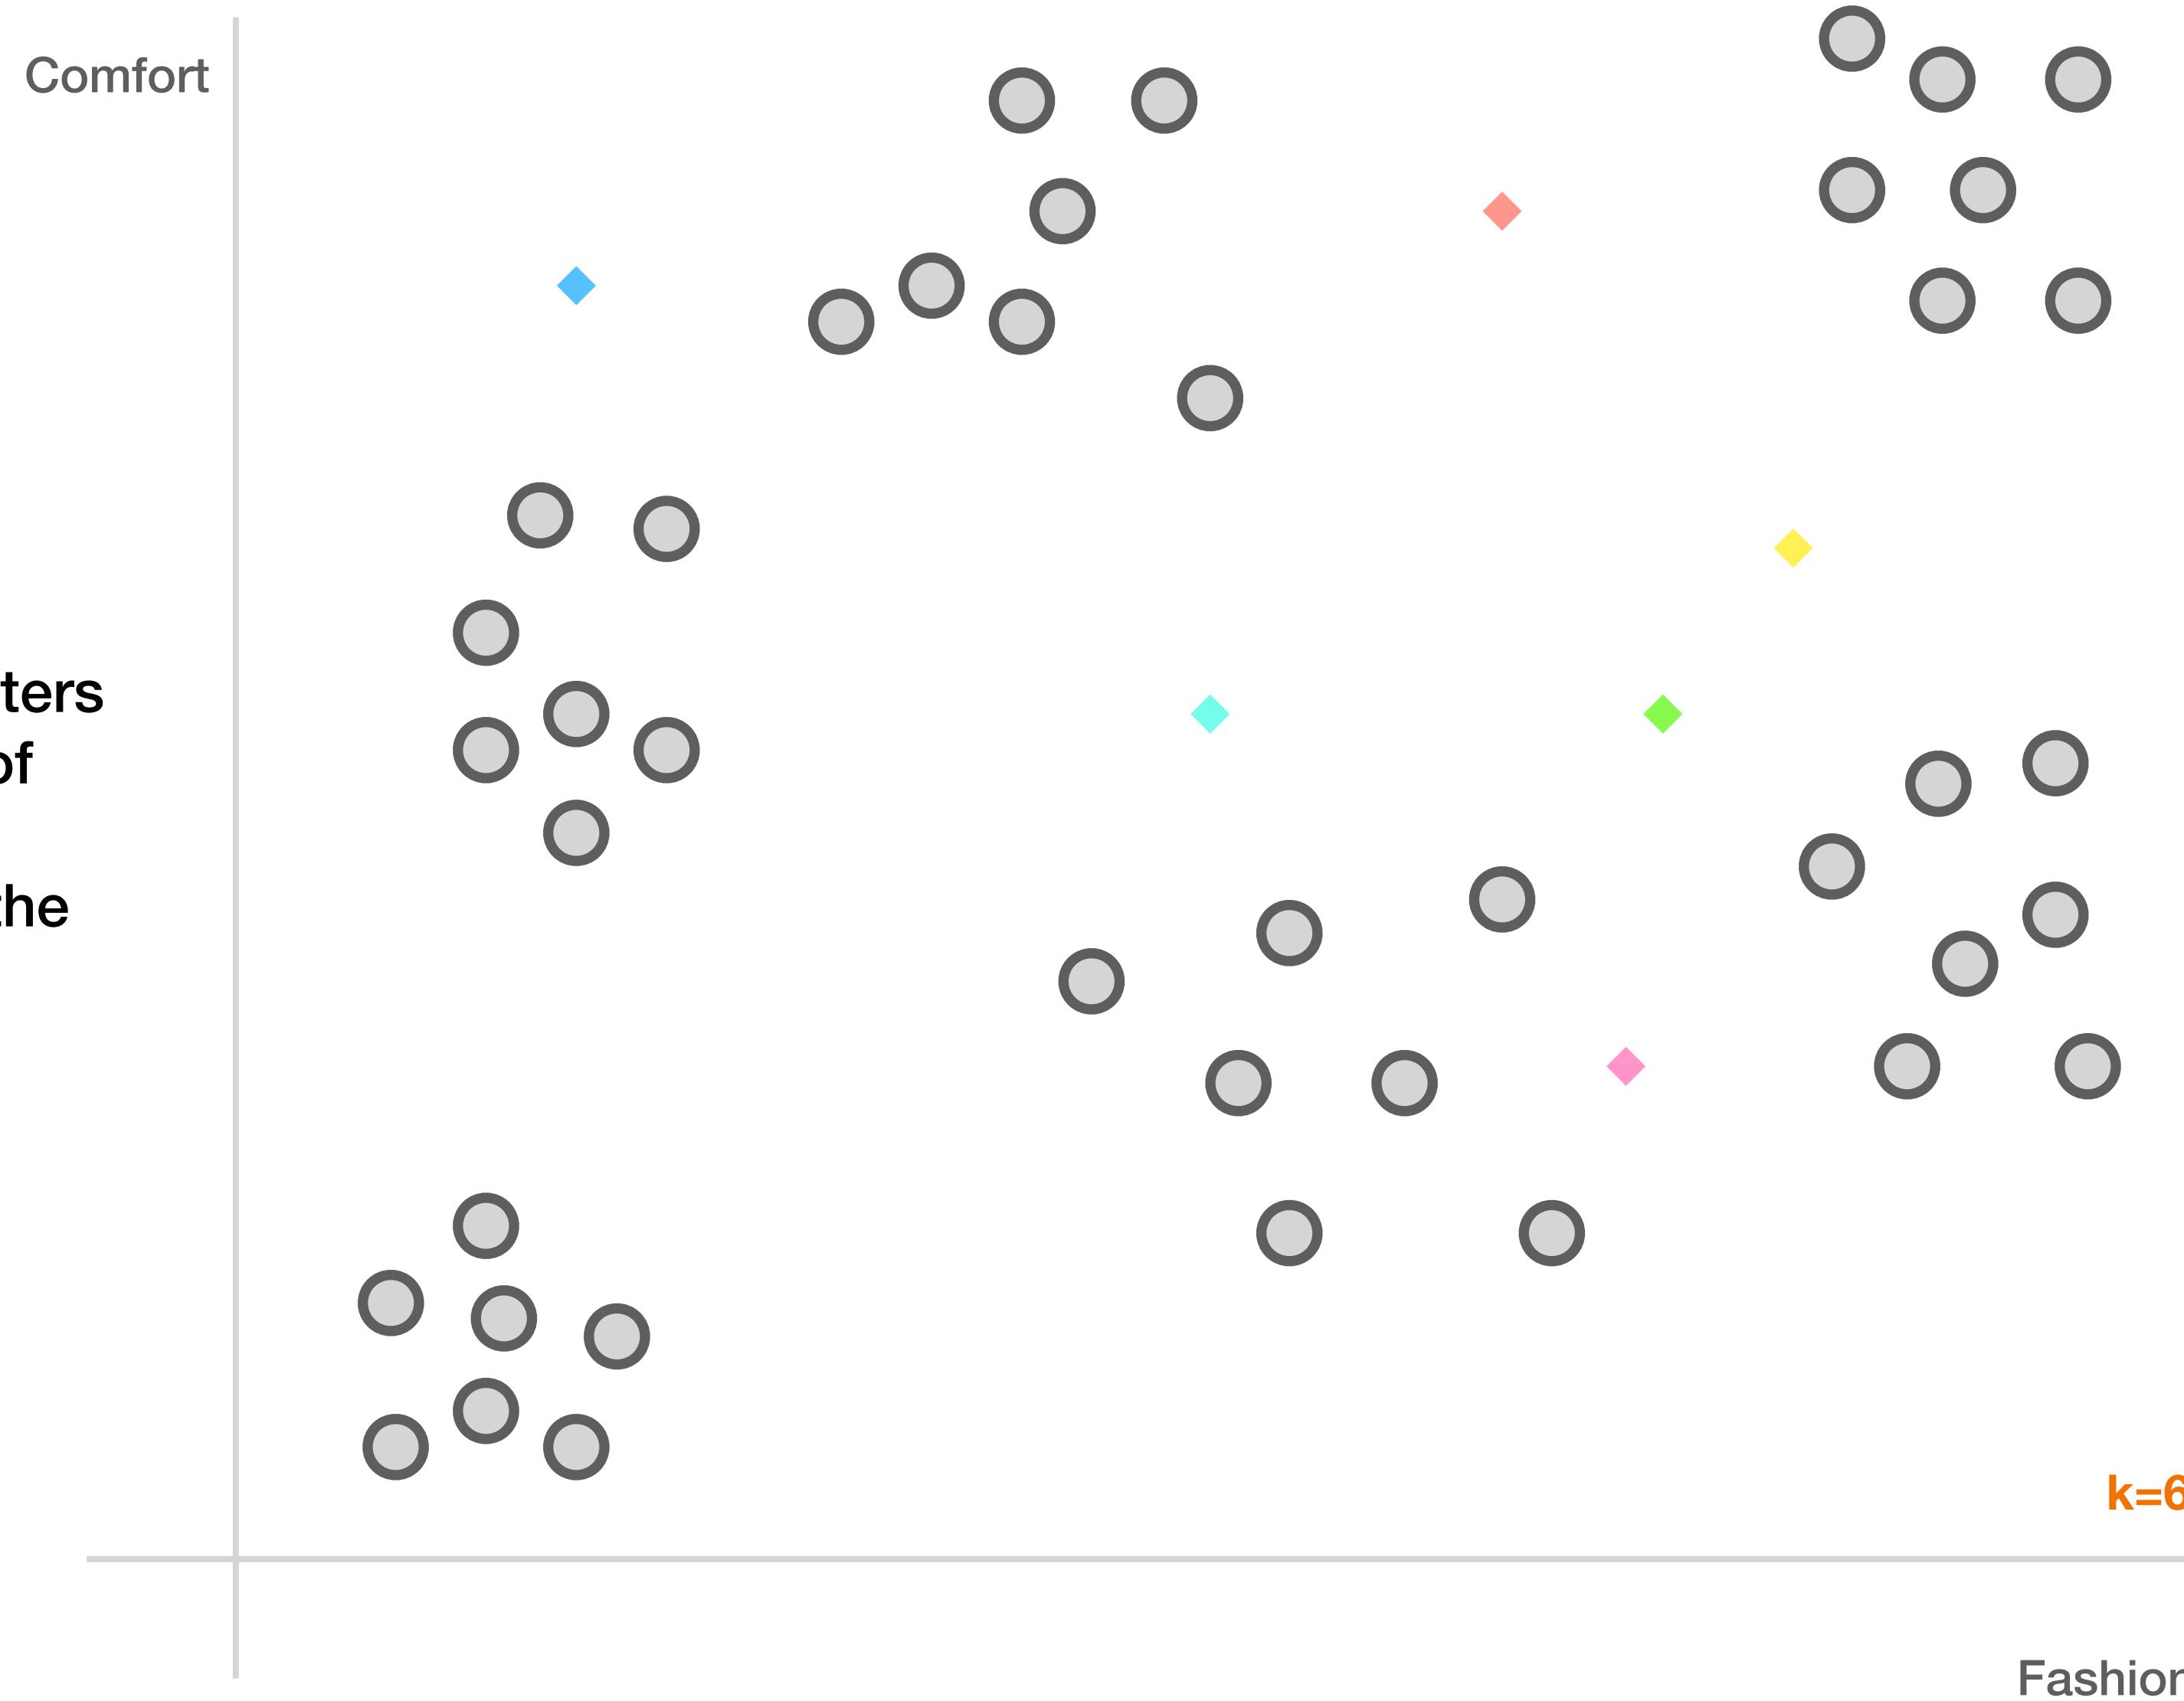
really
terrible



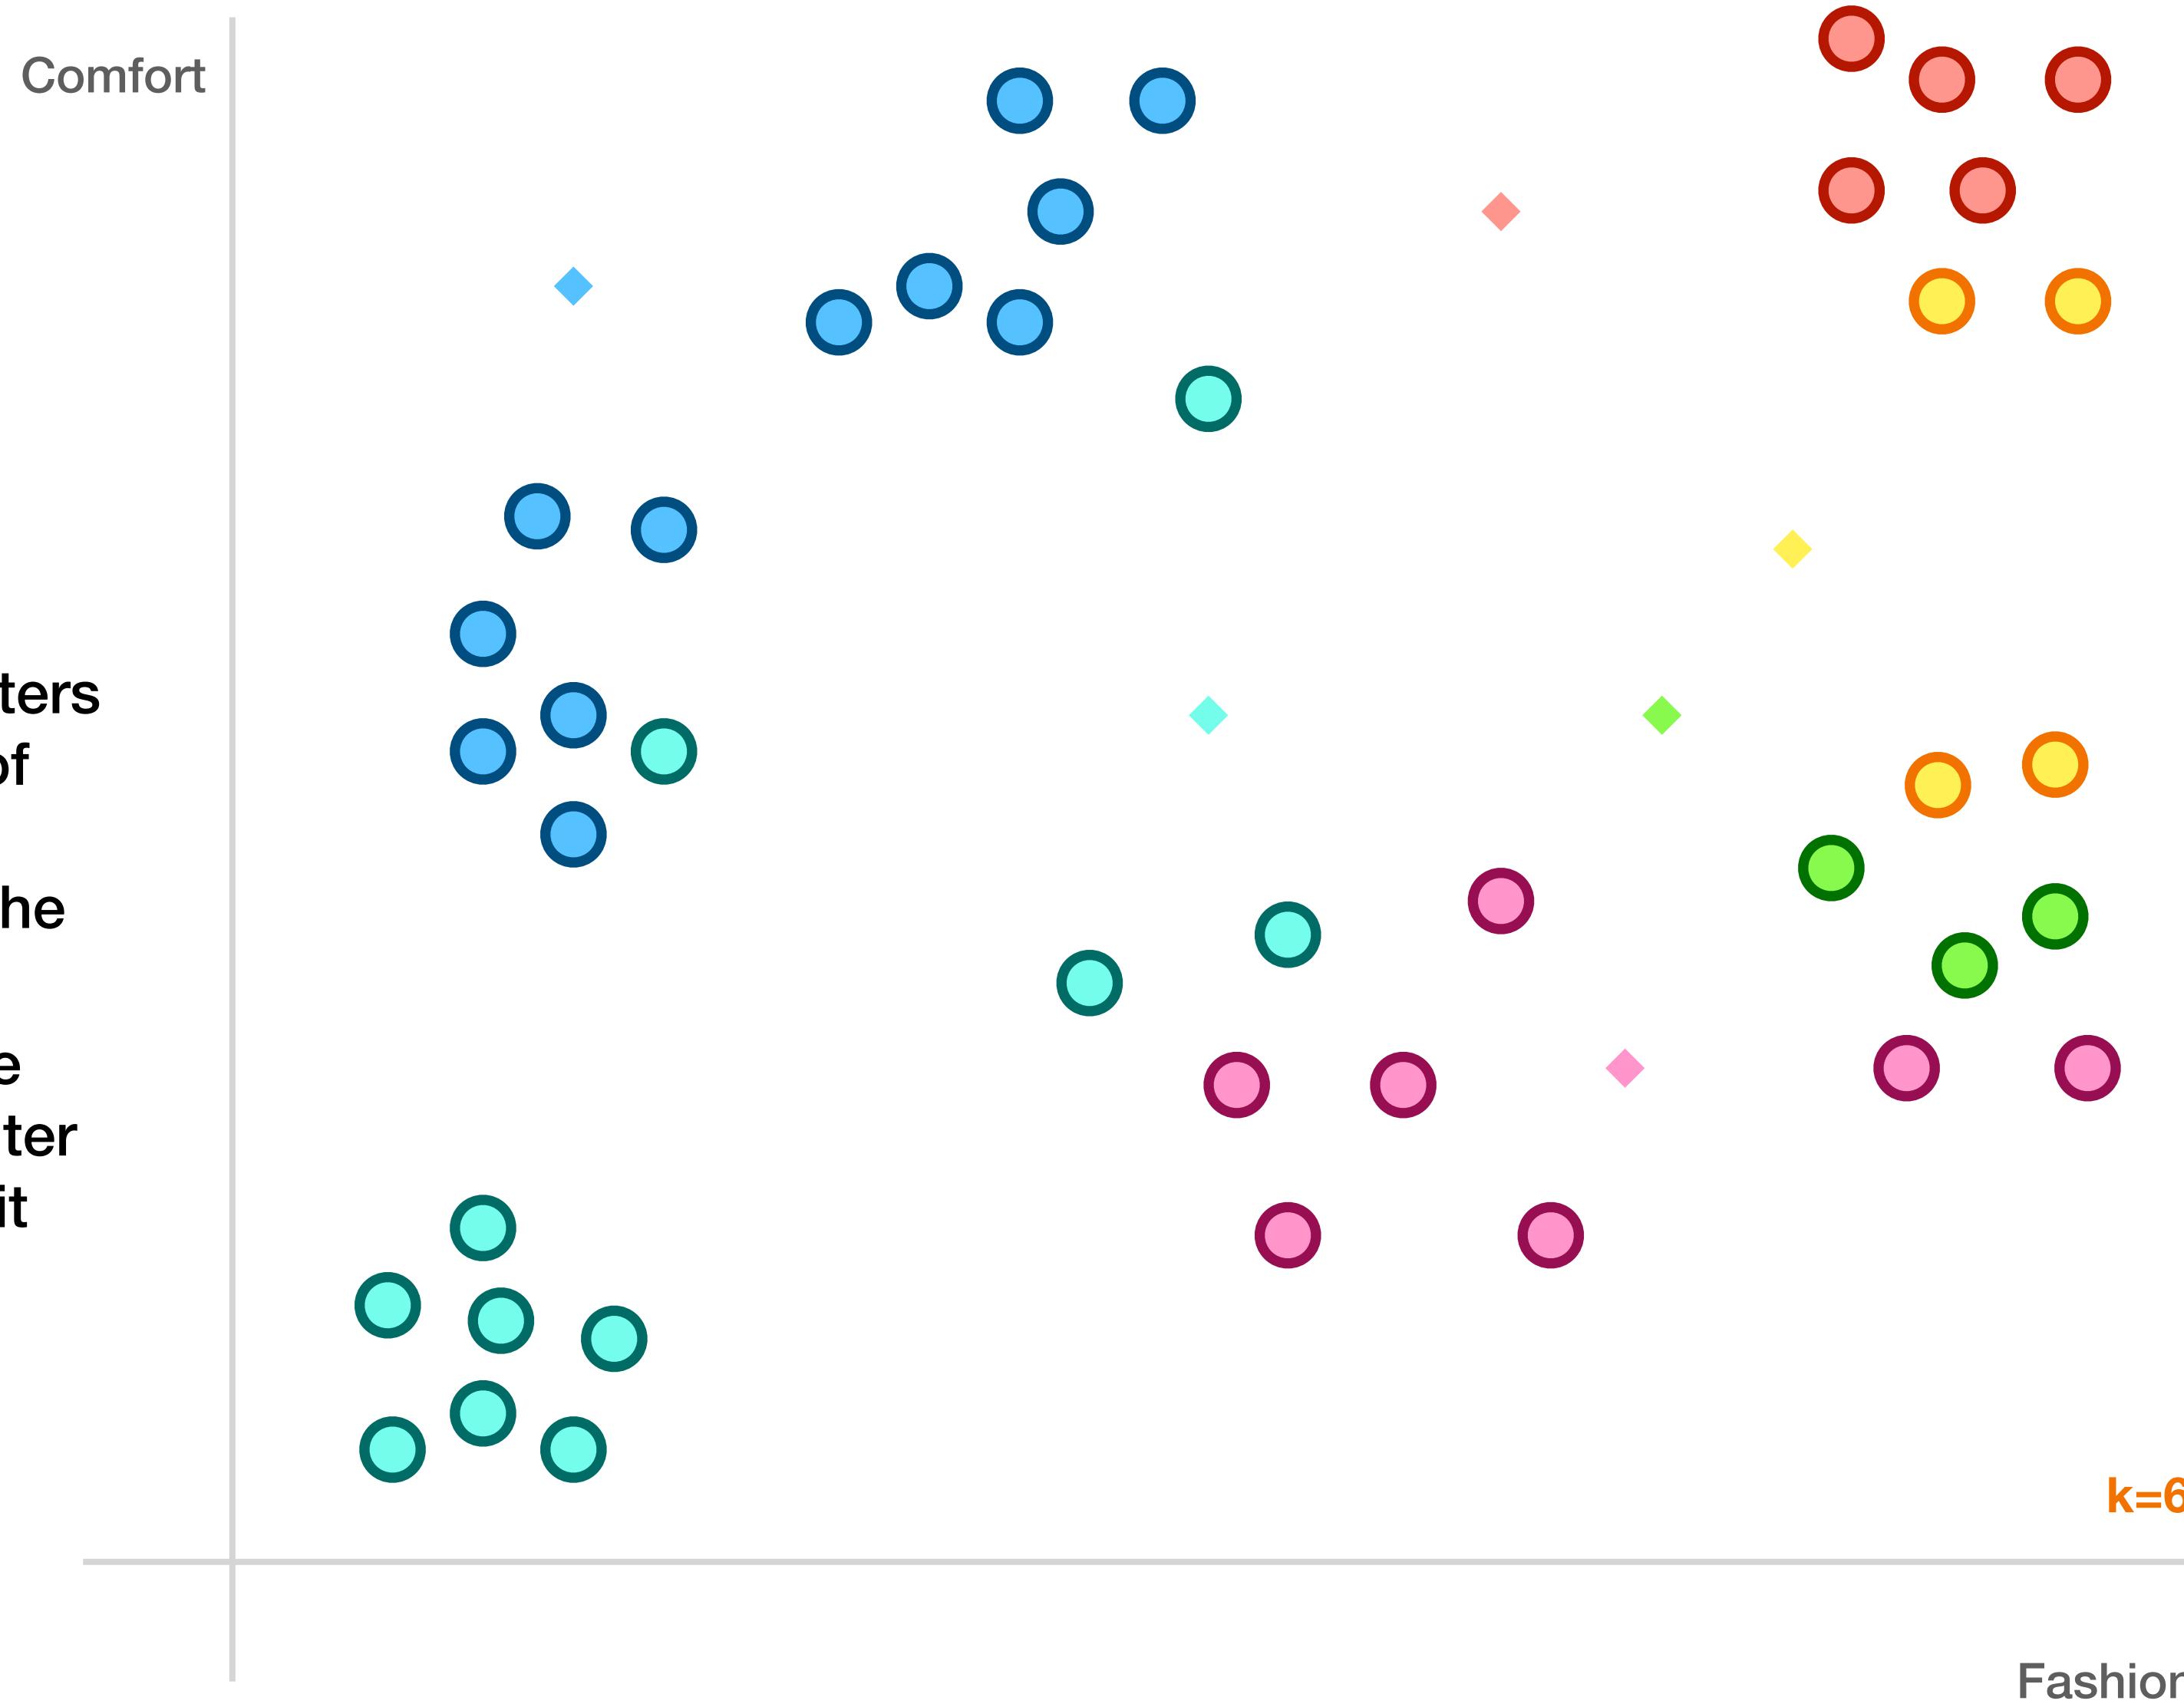
Fashion

- **clusters** can tell us specifics about the relationship of data
- ...even if they are unlabeled! → unsupervised learning!

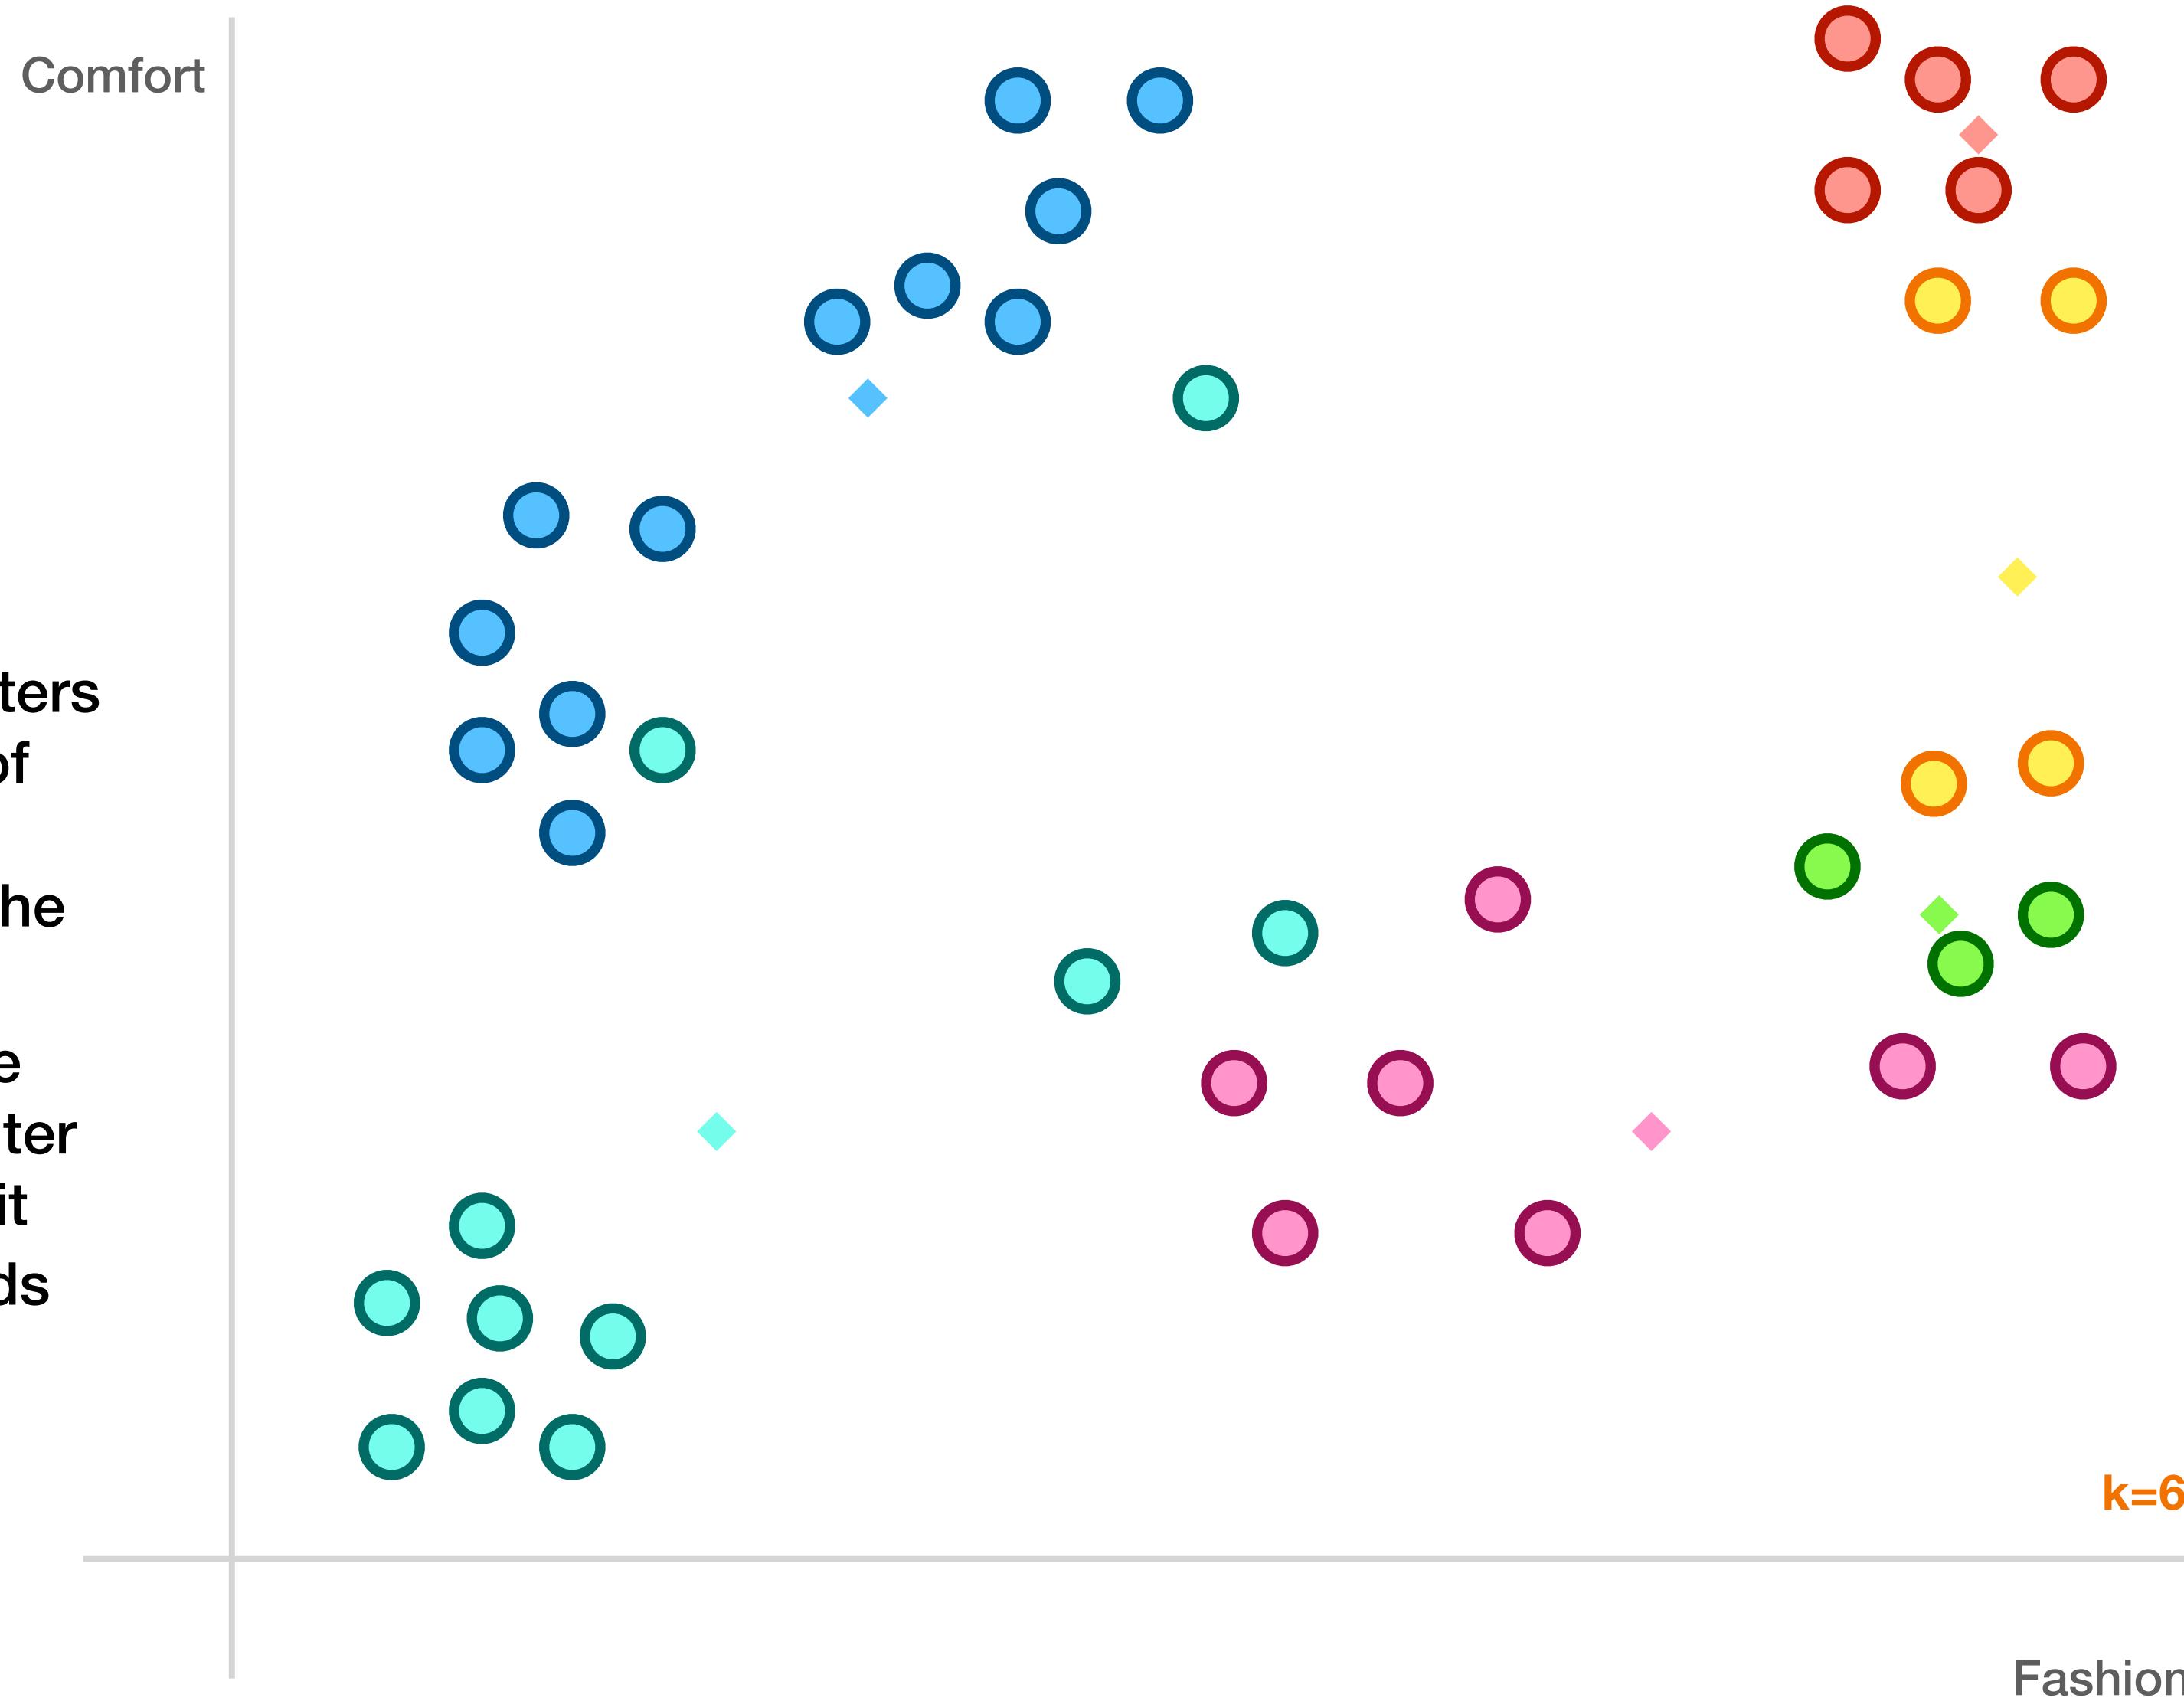
1. pick a K-number of clusters
2. randomly pick a series of "centroids"
3. assign each particle to the centroid closest to it



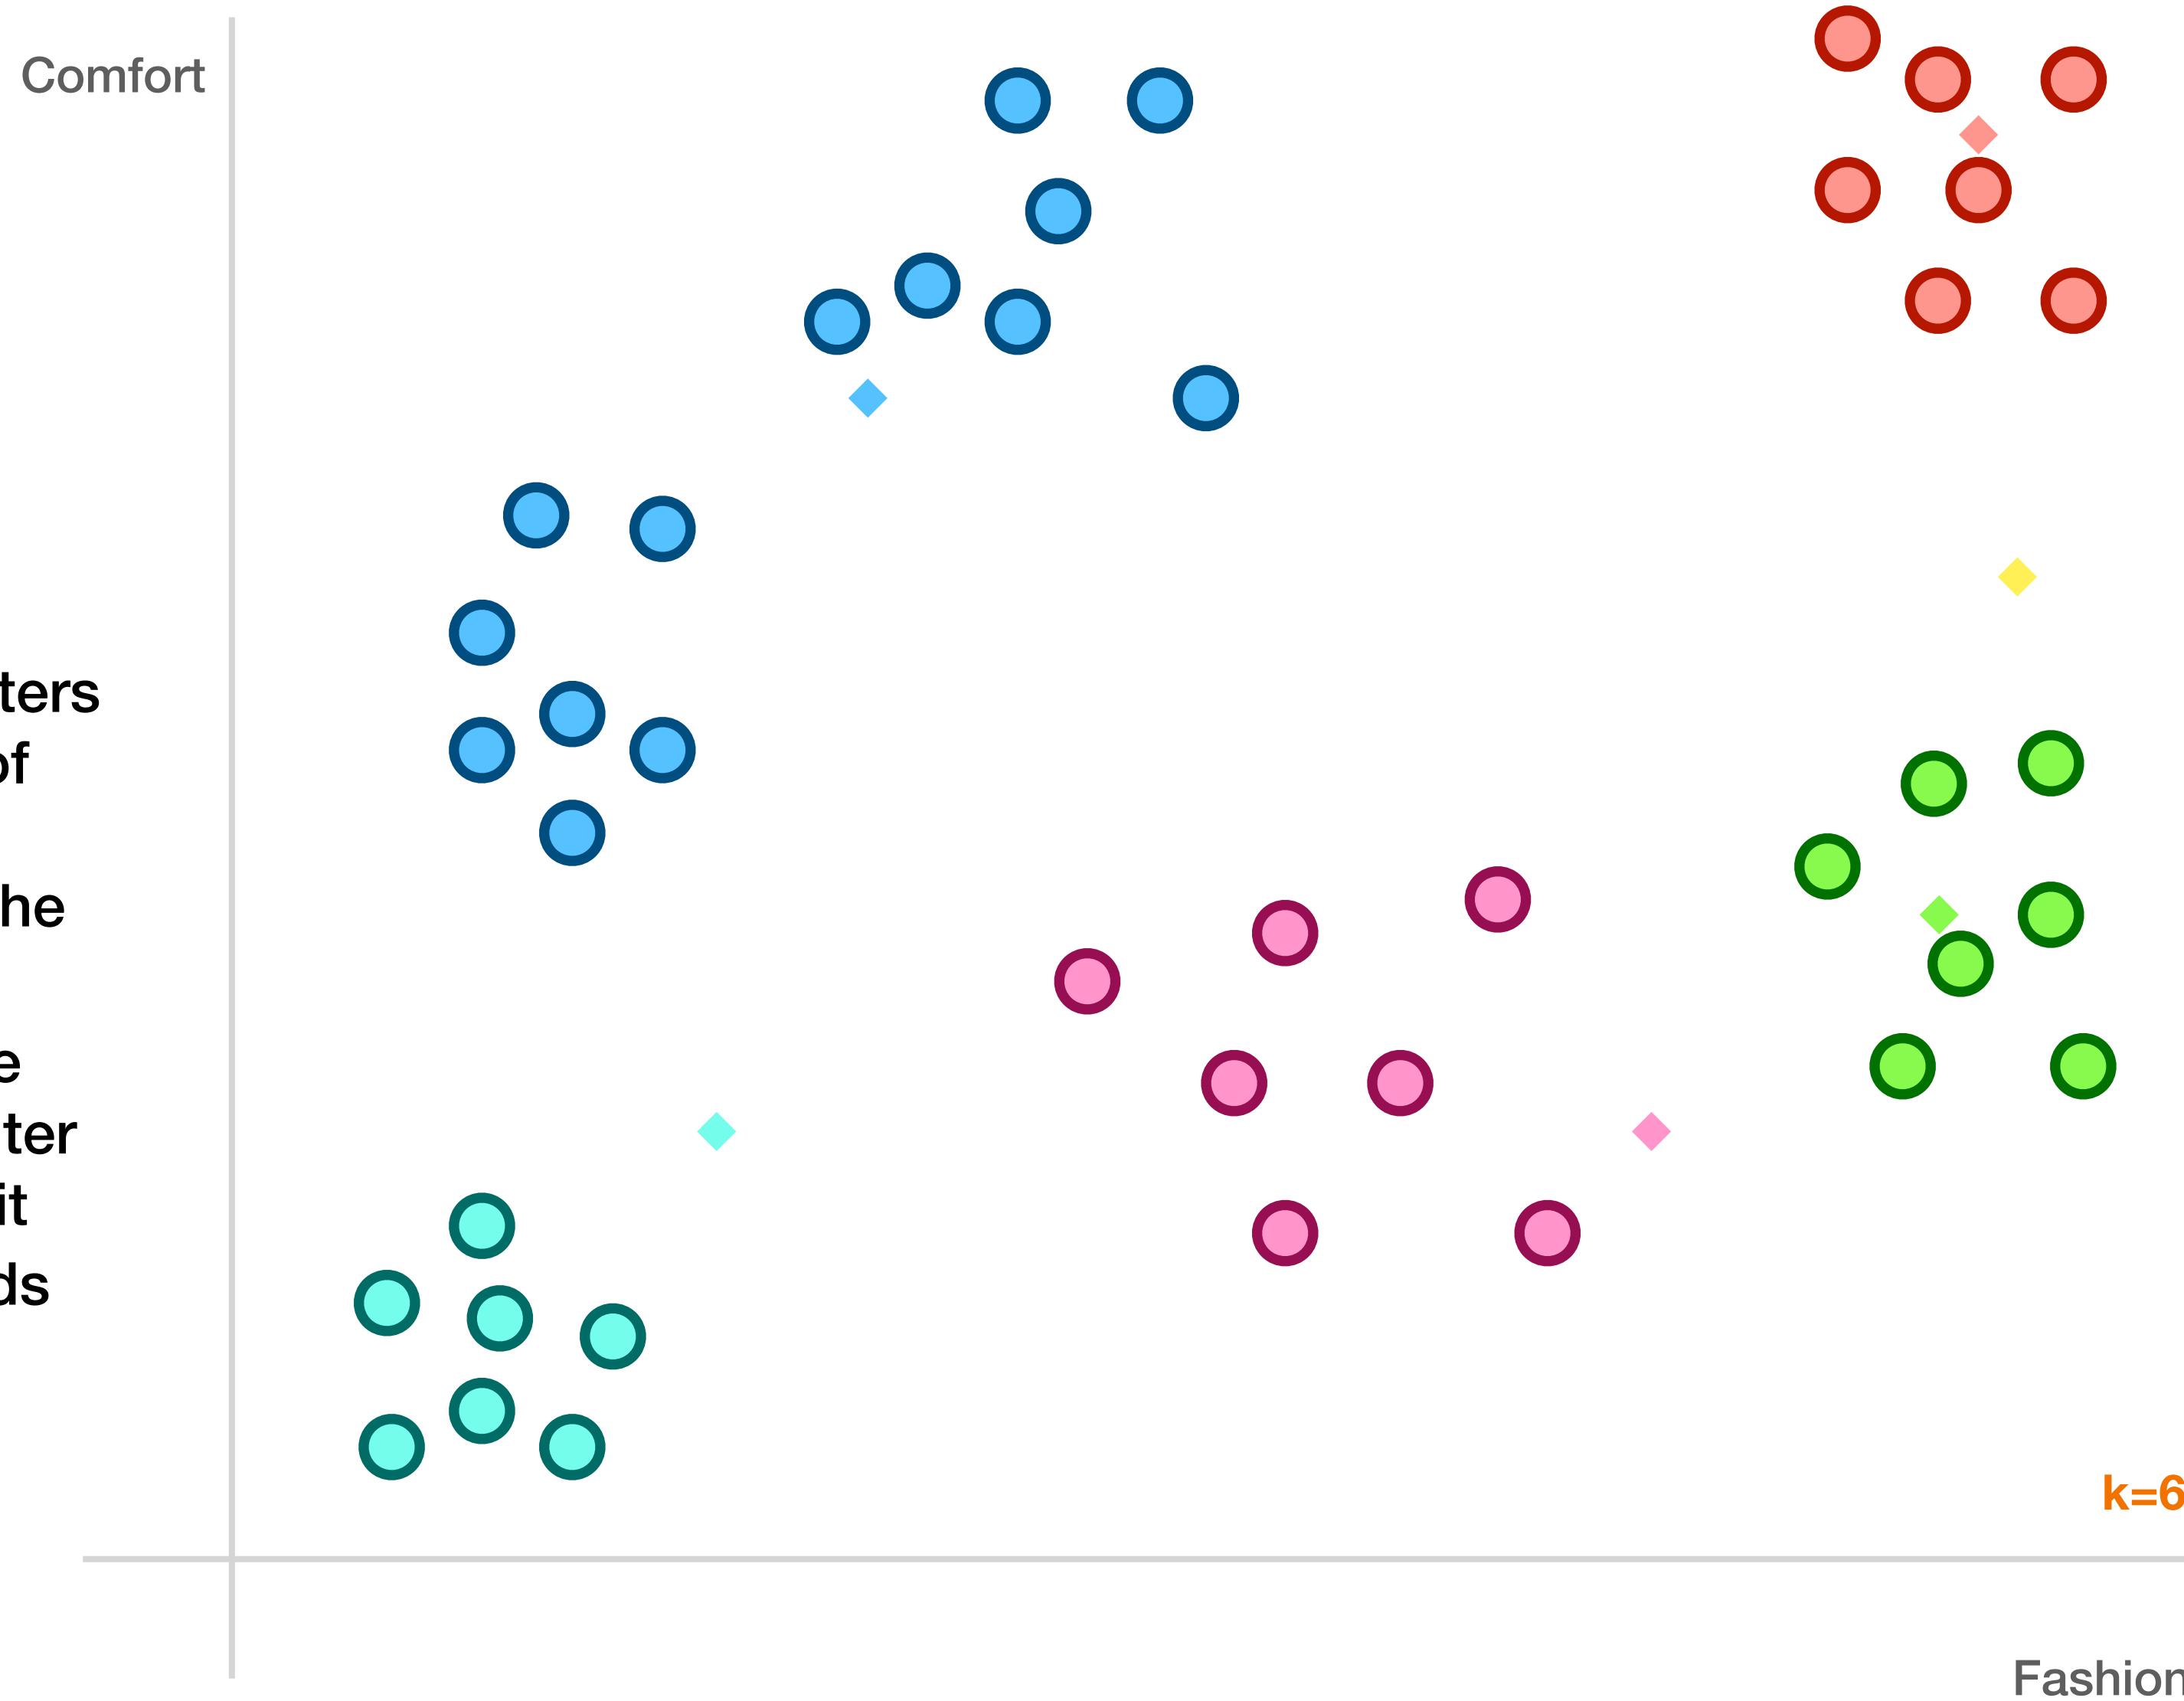
1. pick a K-number of clusters
2. randomly pick a series of "centroids"
3. assign each particle to the centroid closest to it
4. move the centroid to the weighted geometric center of samples assigned to it



1. pick a K-number of clusters
2. randomly pick a series of "centroids"
3. assign each particle to the centroid closest to it
4. move the centroid to the weighted geometric center of samples assigned to it
5. Repeat 3-4 until centroids stop moving!

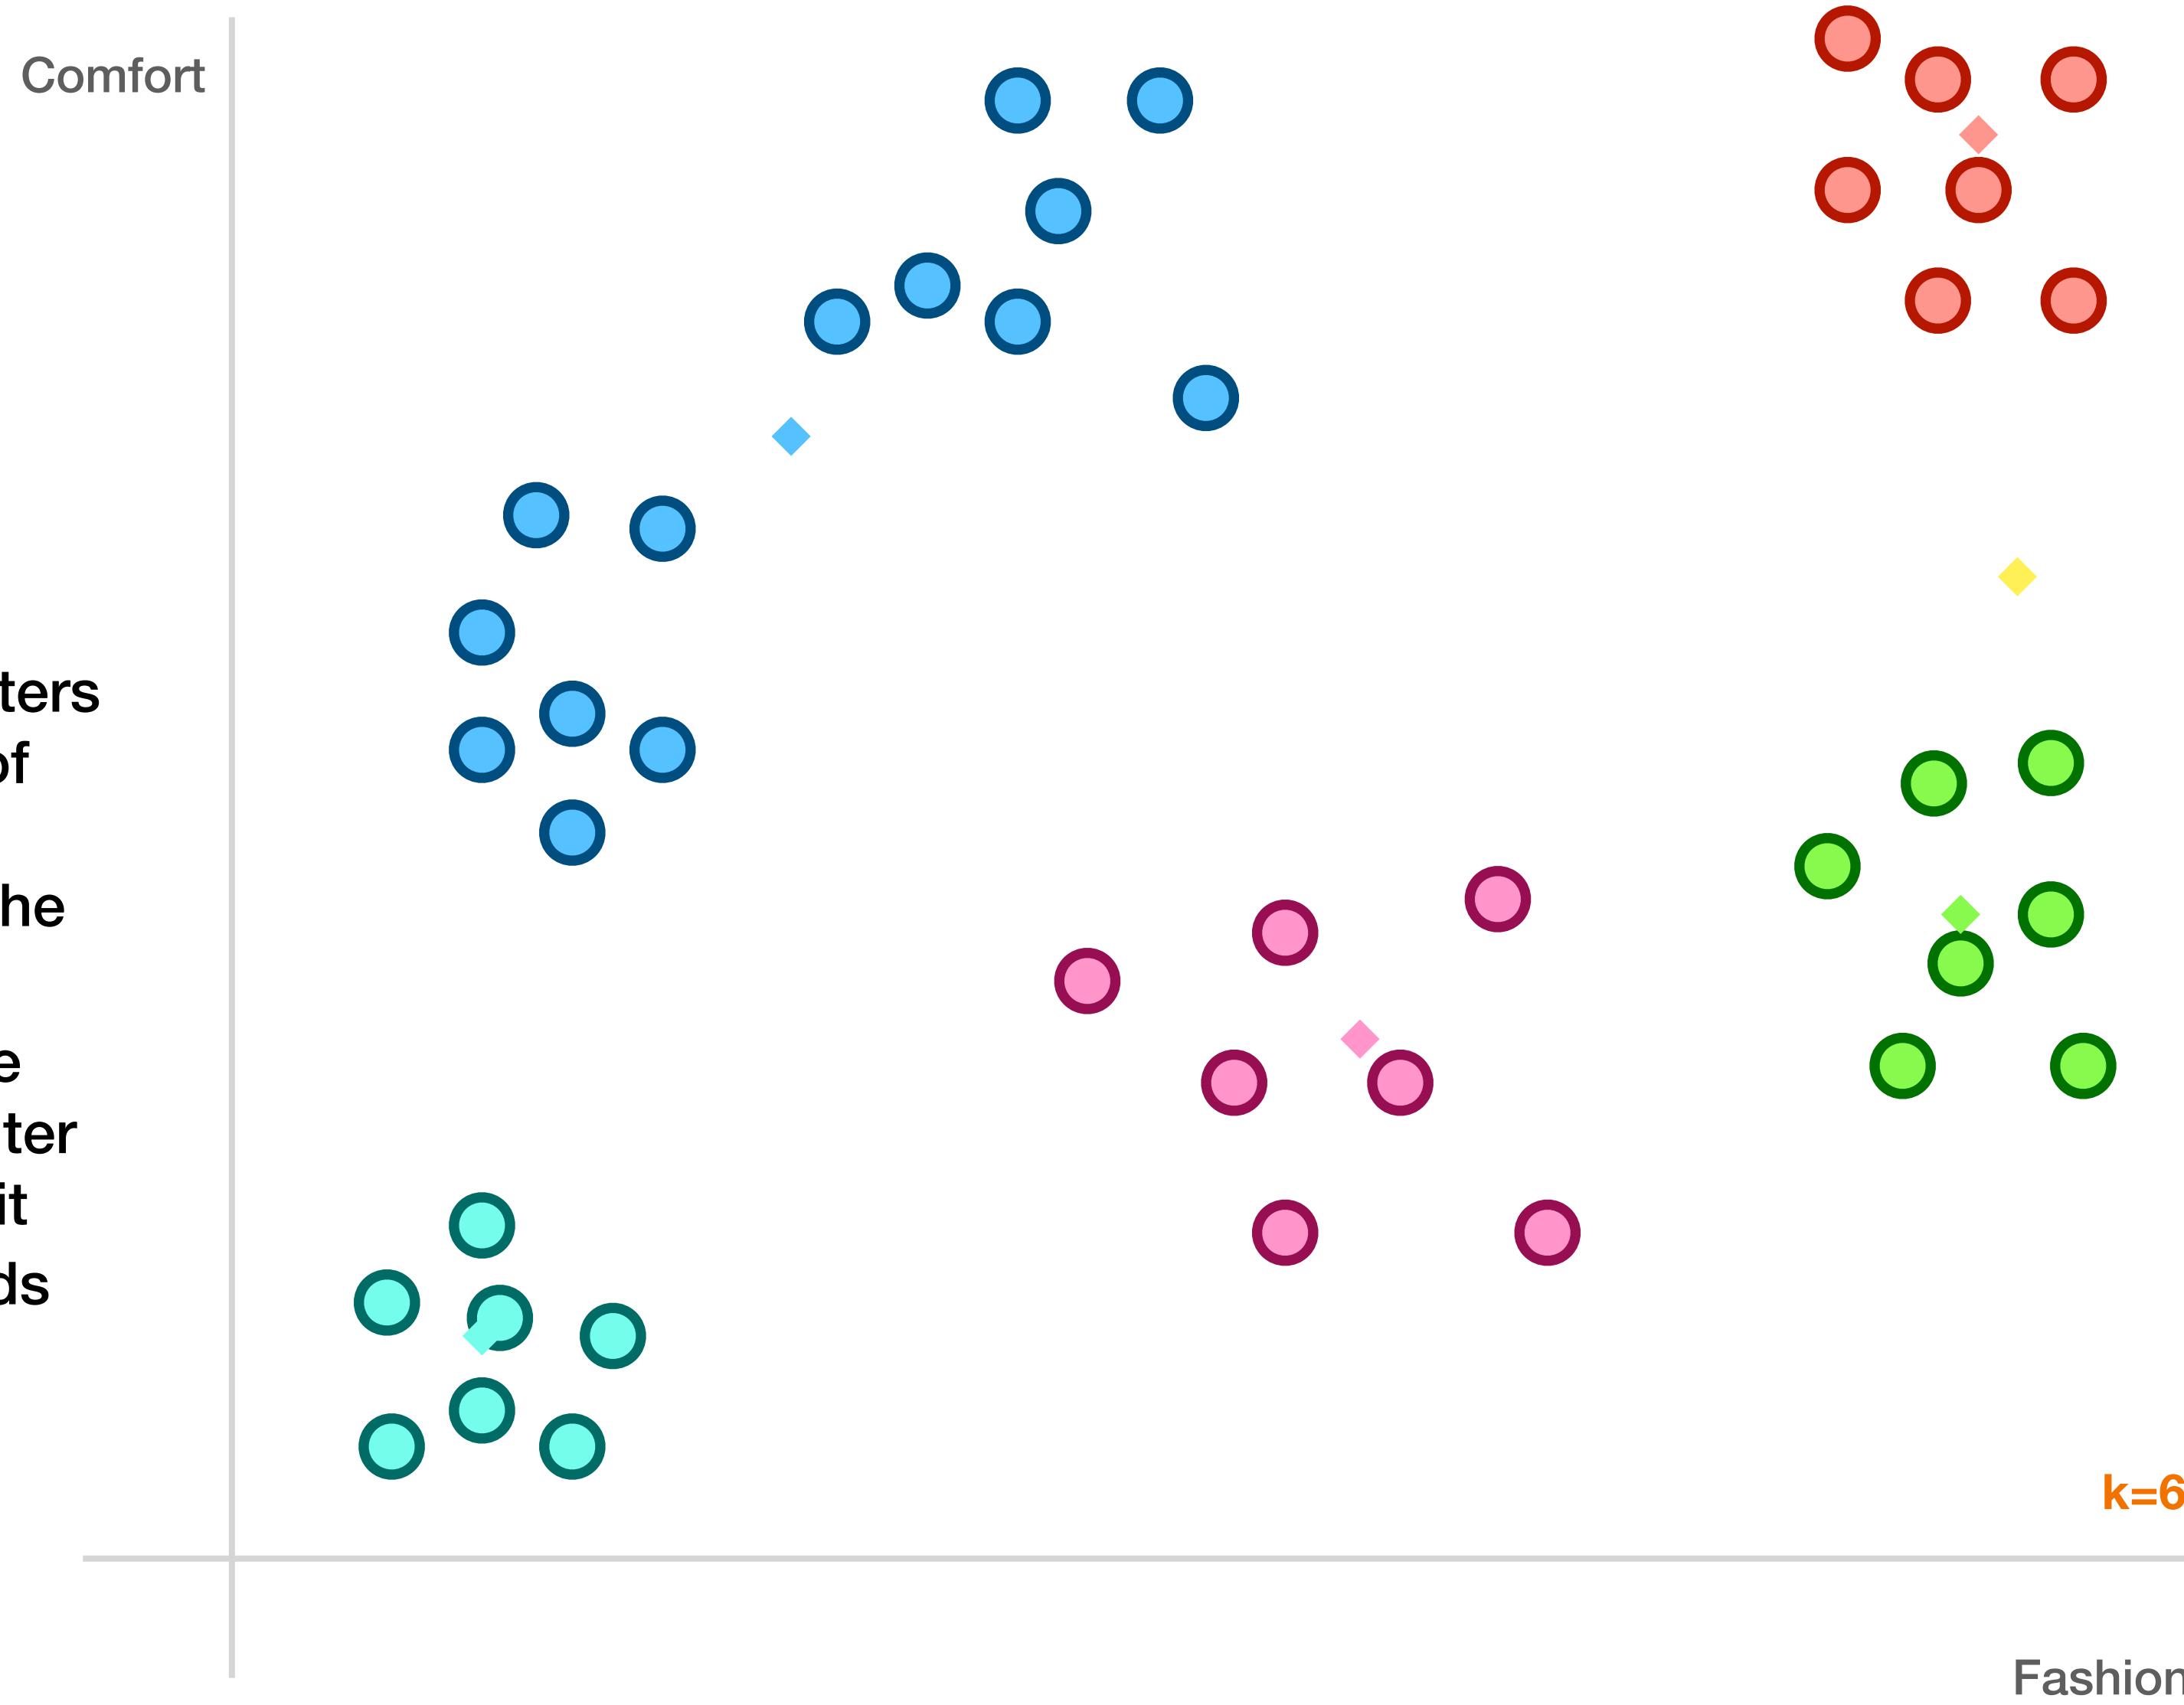


1. pick a K-number of clusters
2. randomly pick a series of "centroids"
3. assign each particle to the centroid closest to it
4. move the centroid to the weighted geometric center of samples assigned to it
5. Repeat 3-4 until centroids stop moving!



k-means clustering

1. pick a K-number of clusters
2. randomly pick a series of "centroids"
3. assign each particle to the centroid closest to it
4. move the centroid to the weighted geometric center of samples assigned to it
5. Repeat 3-4 until centroids stop moving!



Did we get back
the same clusters?
Nope. And that's OK.

**Did we get back the same clusters?
Nope. And that's OK.**

- **K-means** is an *indeterministic* algorithm—it has built-in randomness

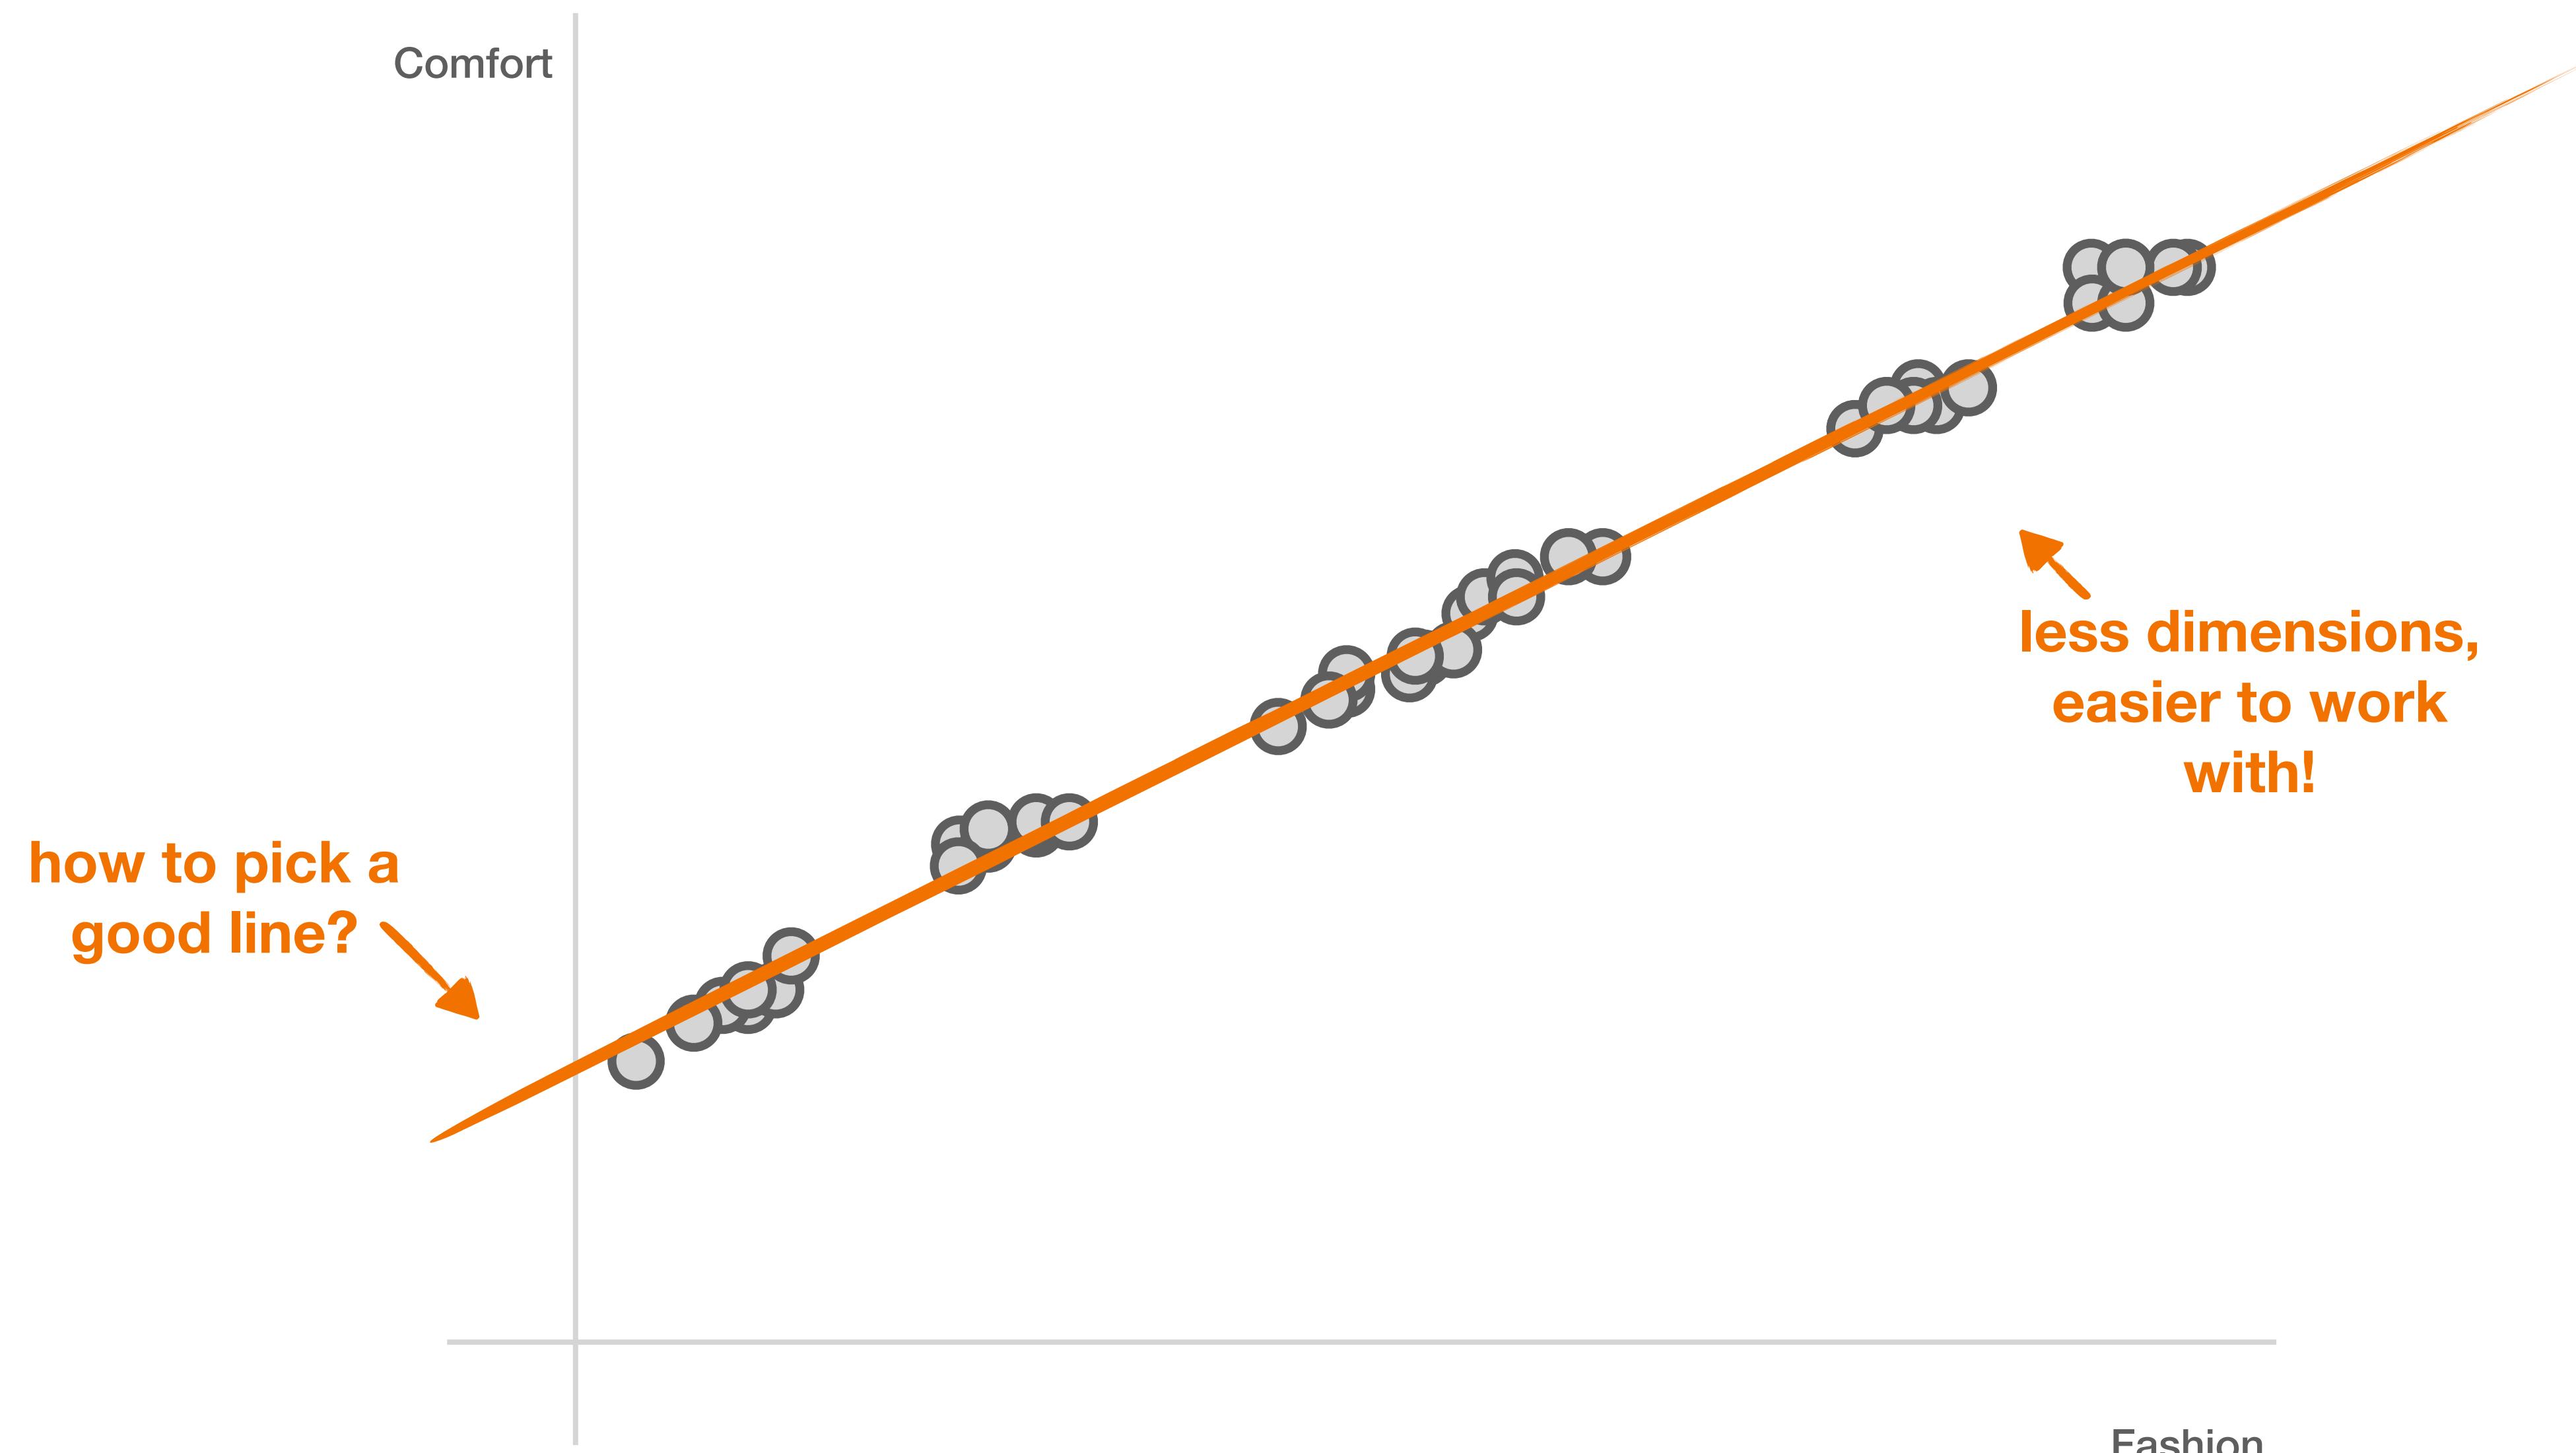
UNSUPERVISED LEARNING

Clustering

Dimension reduction

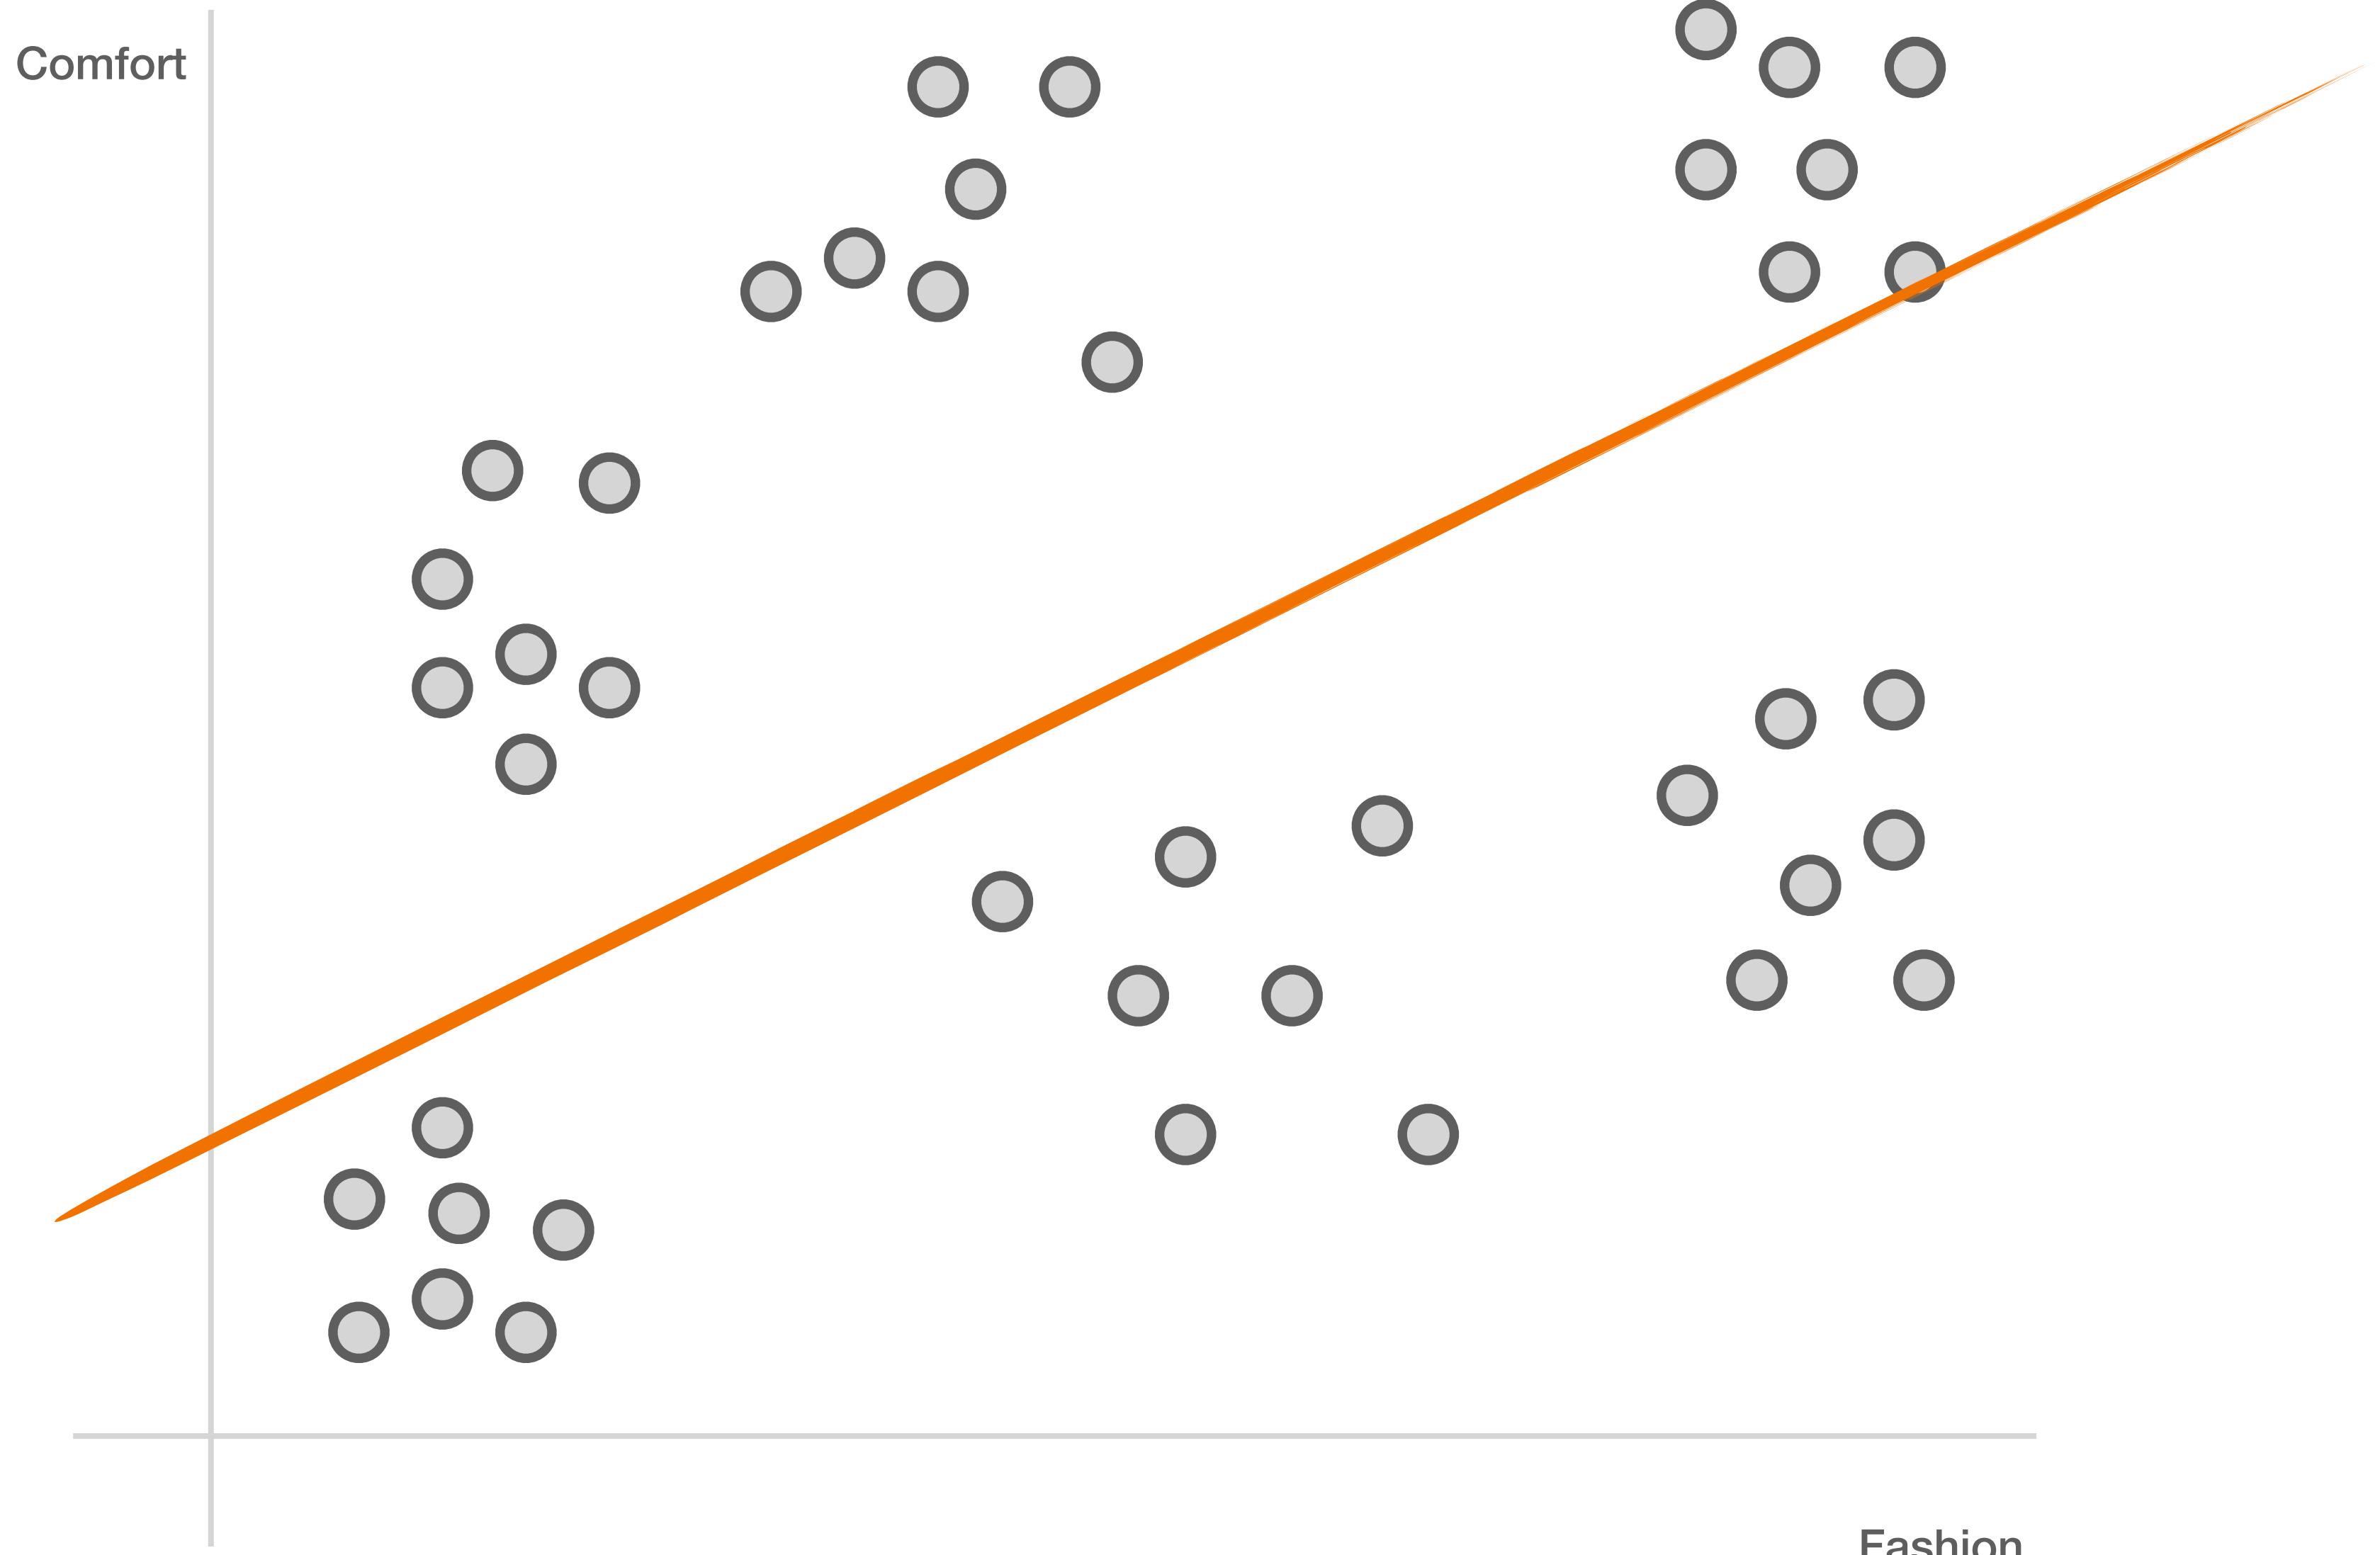
Principle Component Analysis



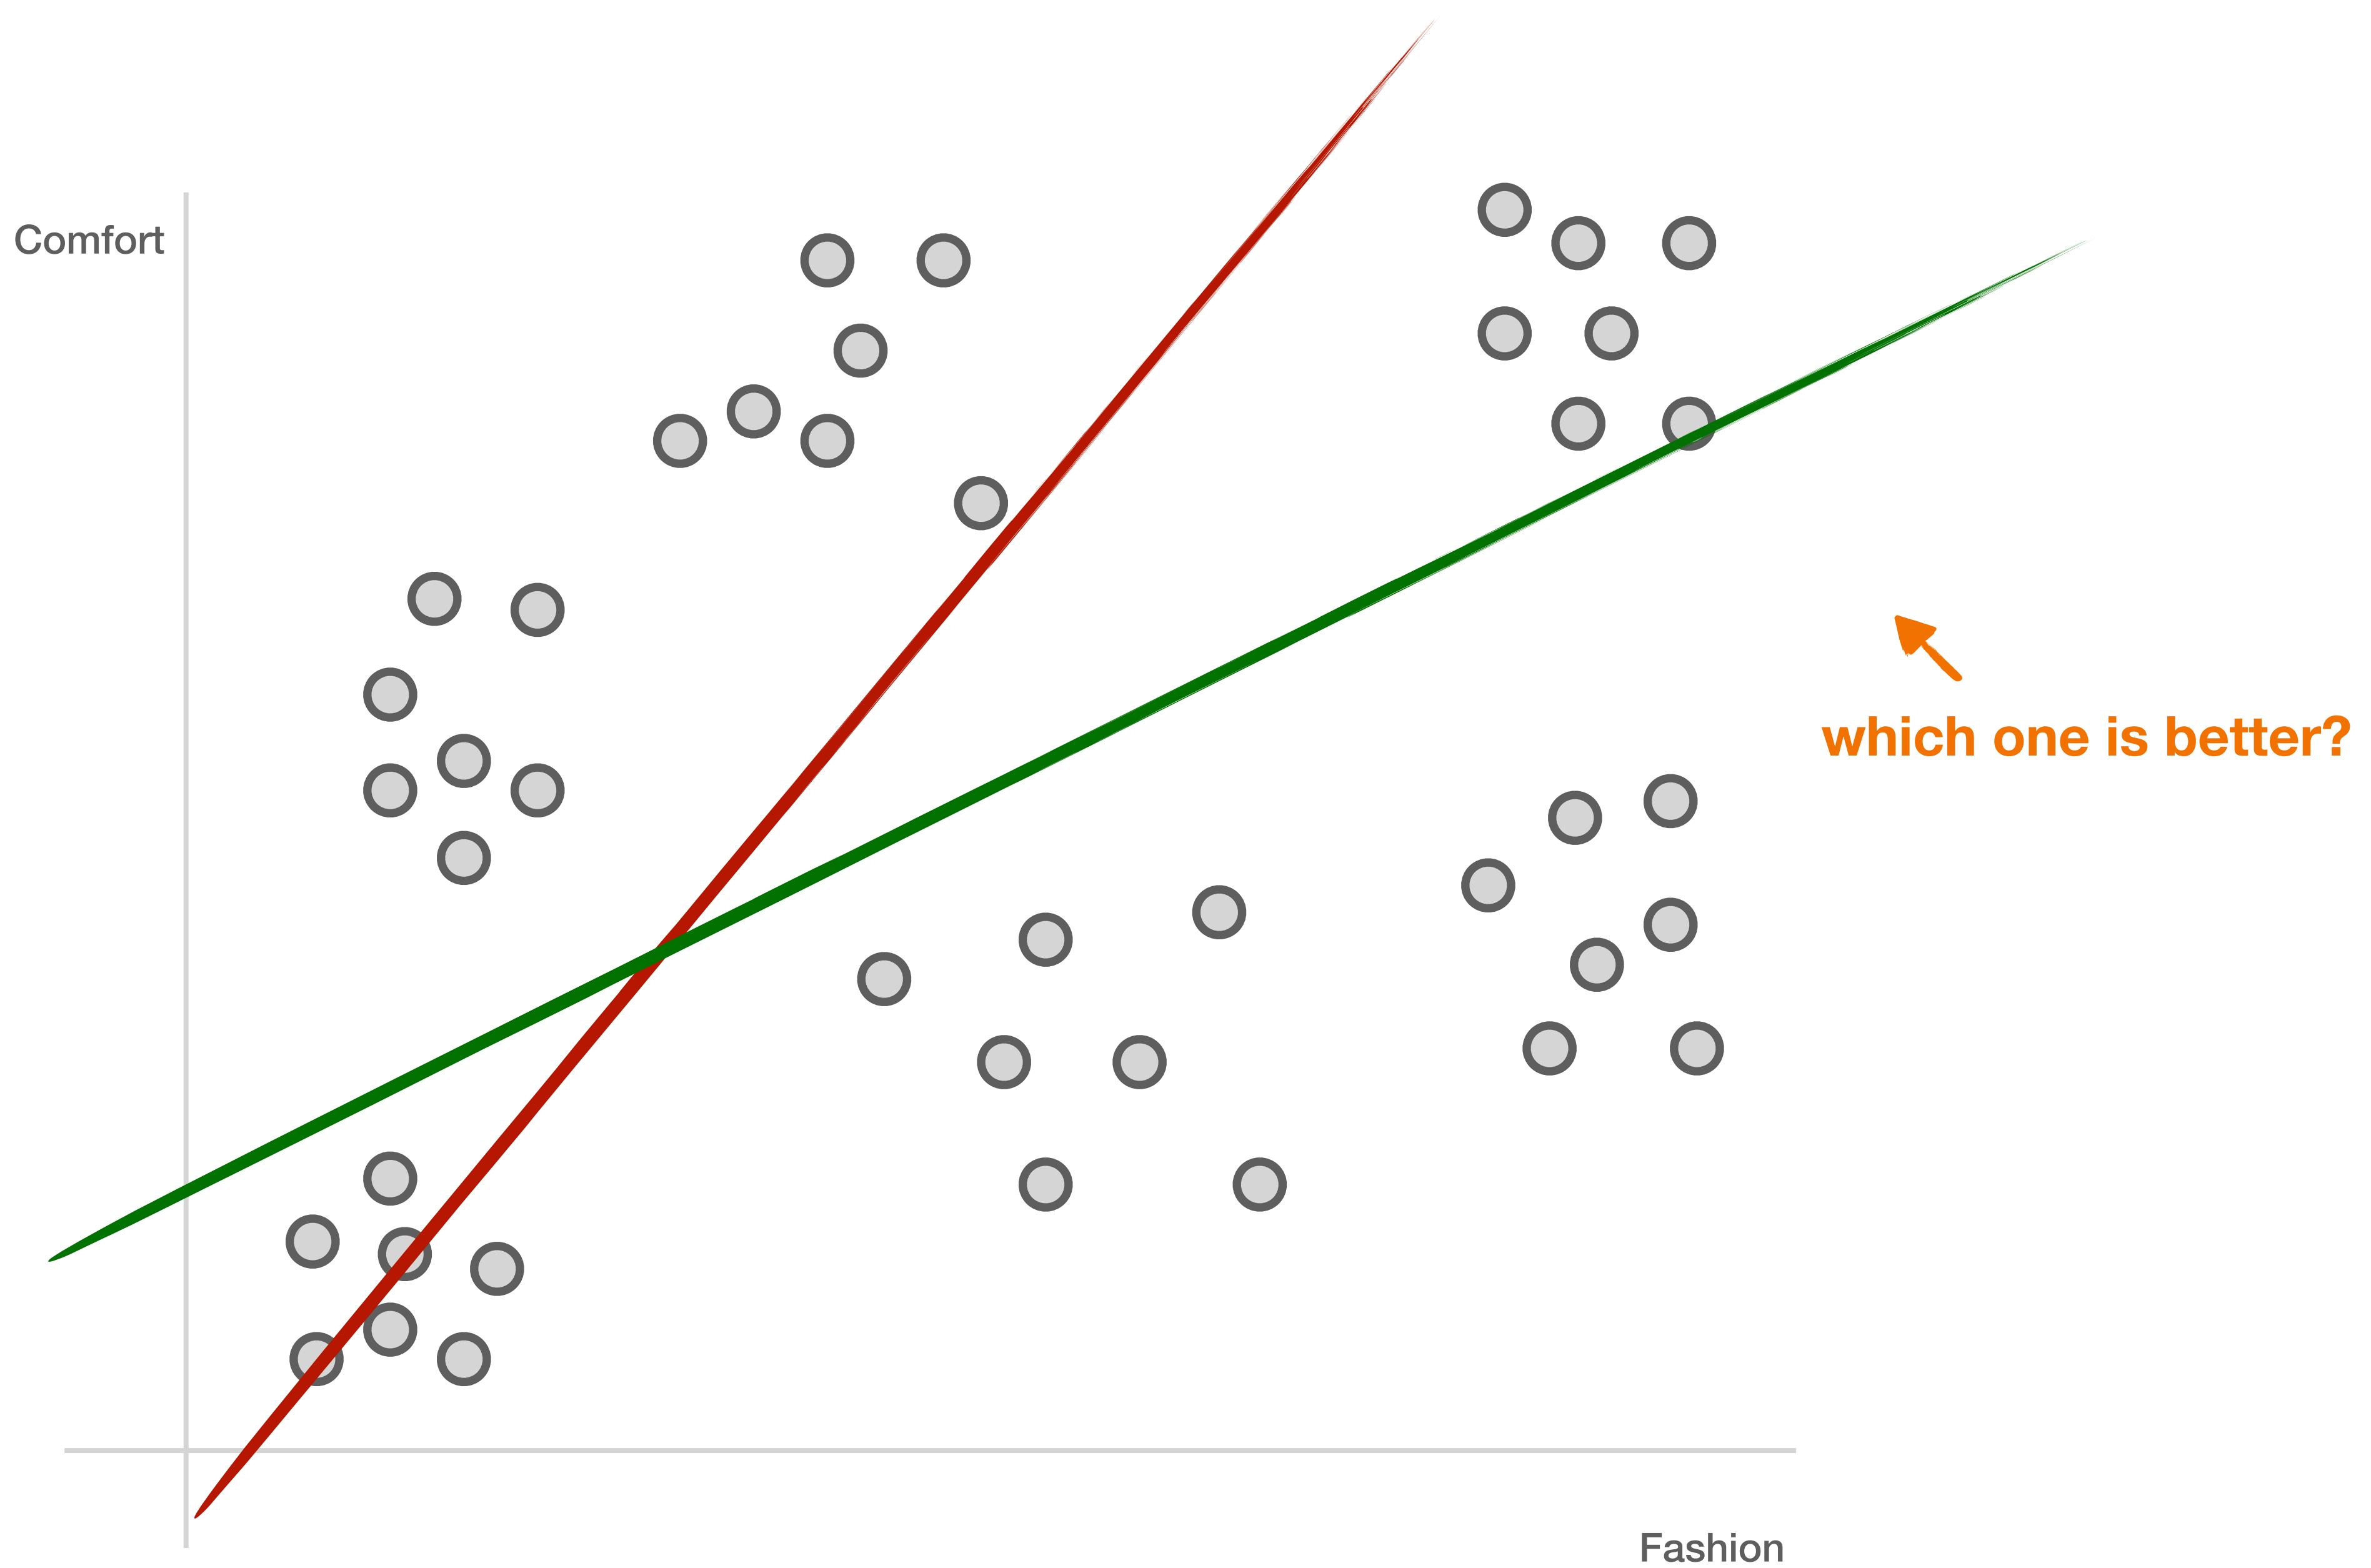


- by **projecting** the samples down to a smaller dimension, they are easier to work with.
(because the centroids have less “space” to move around)

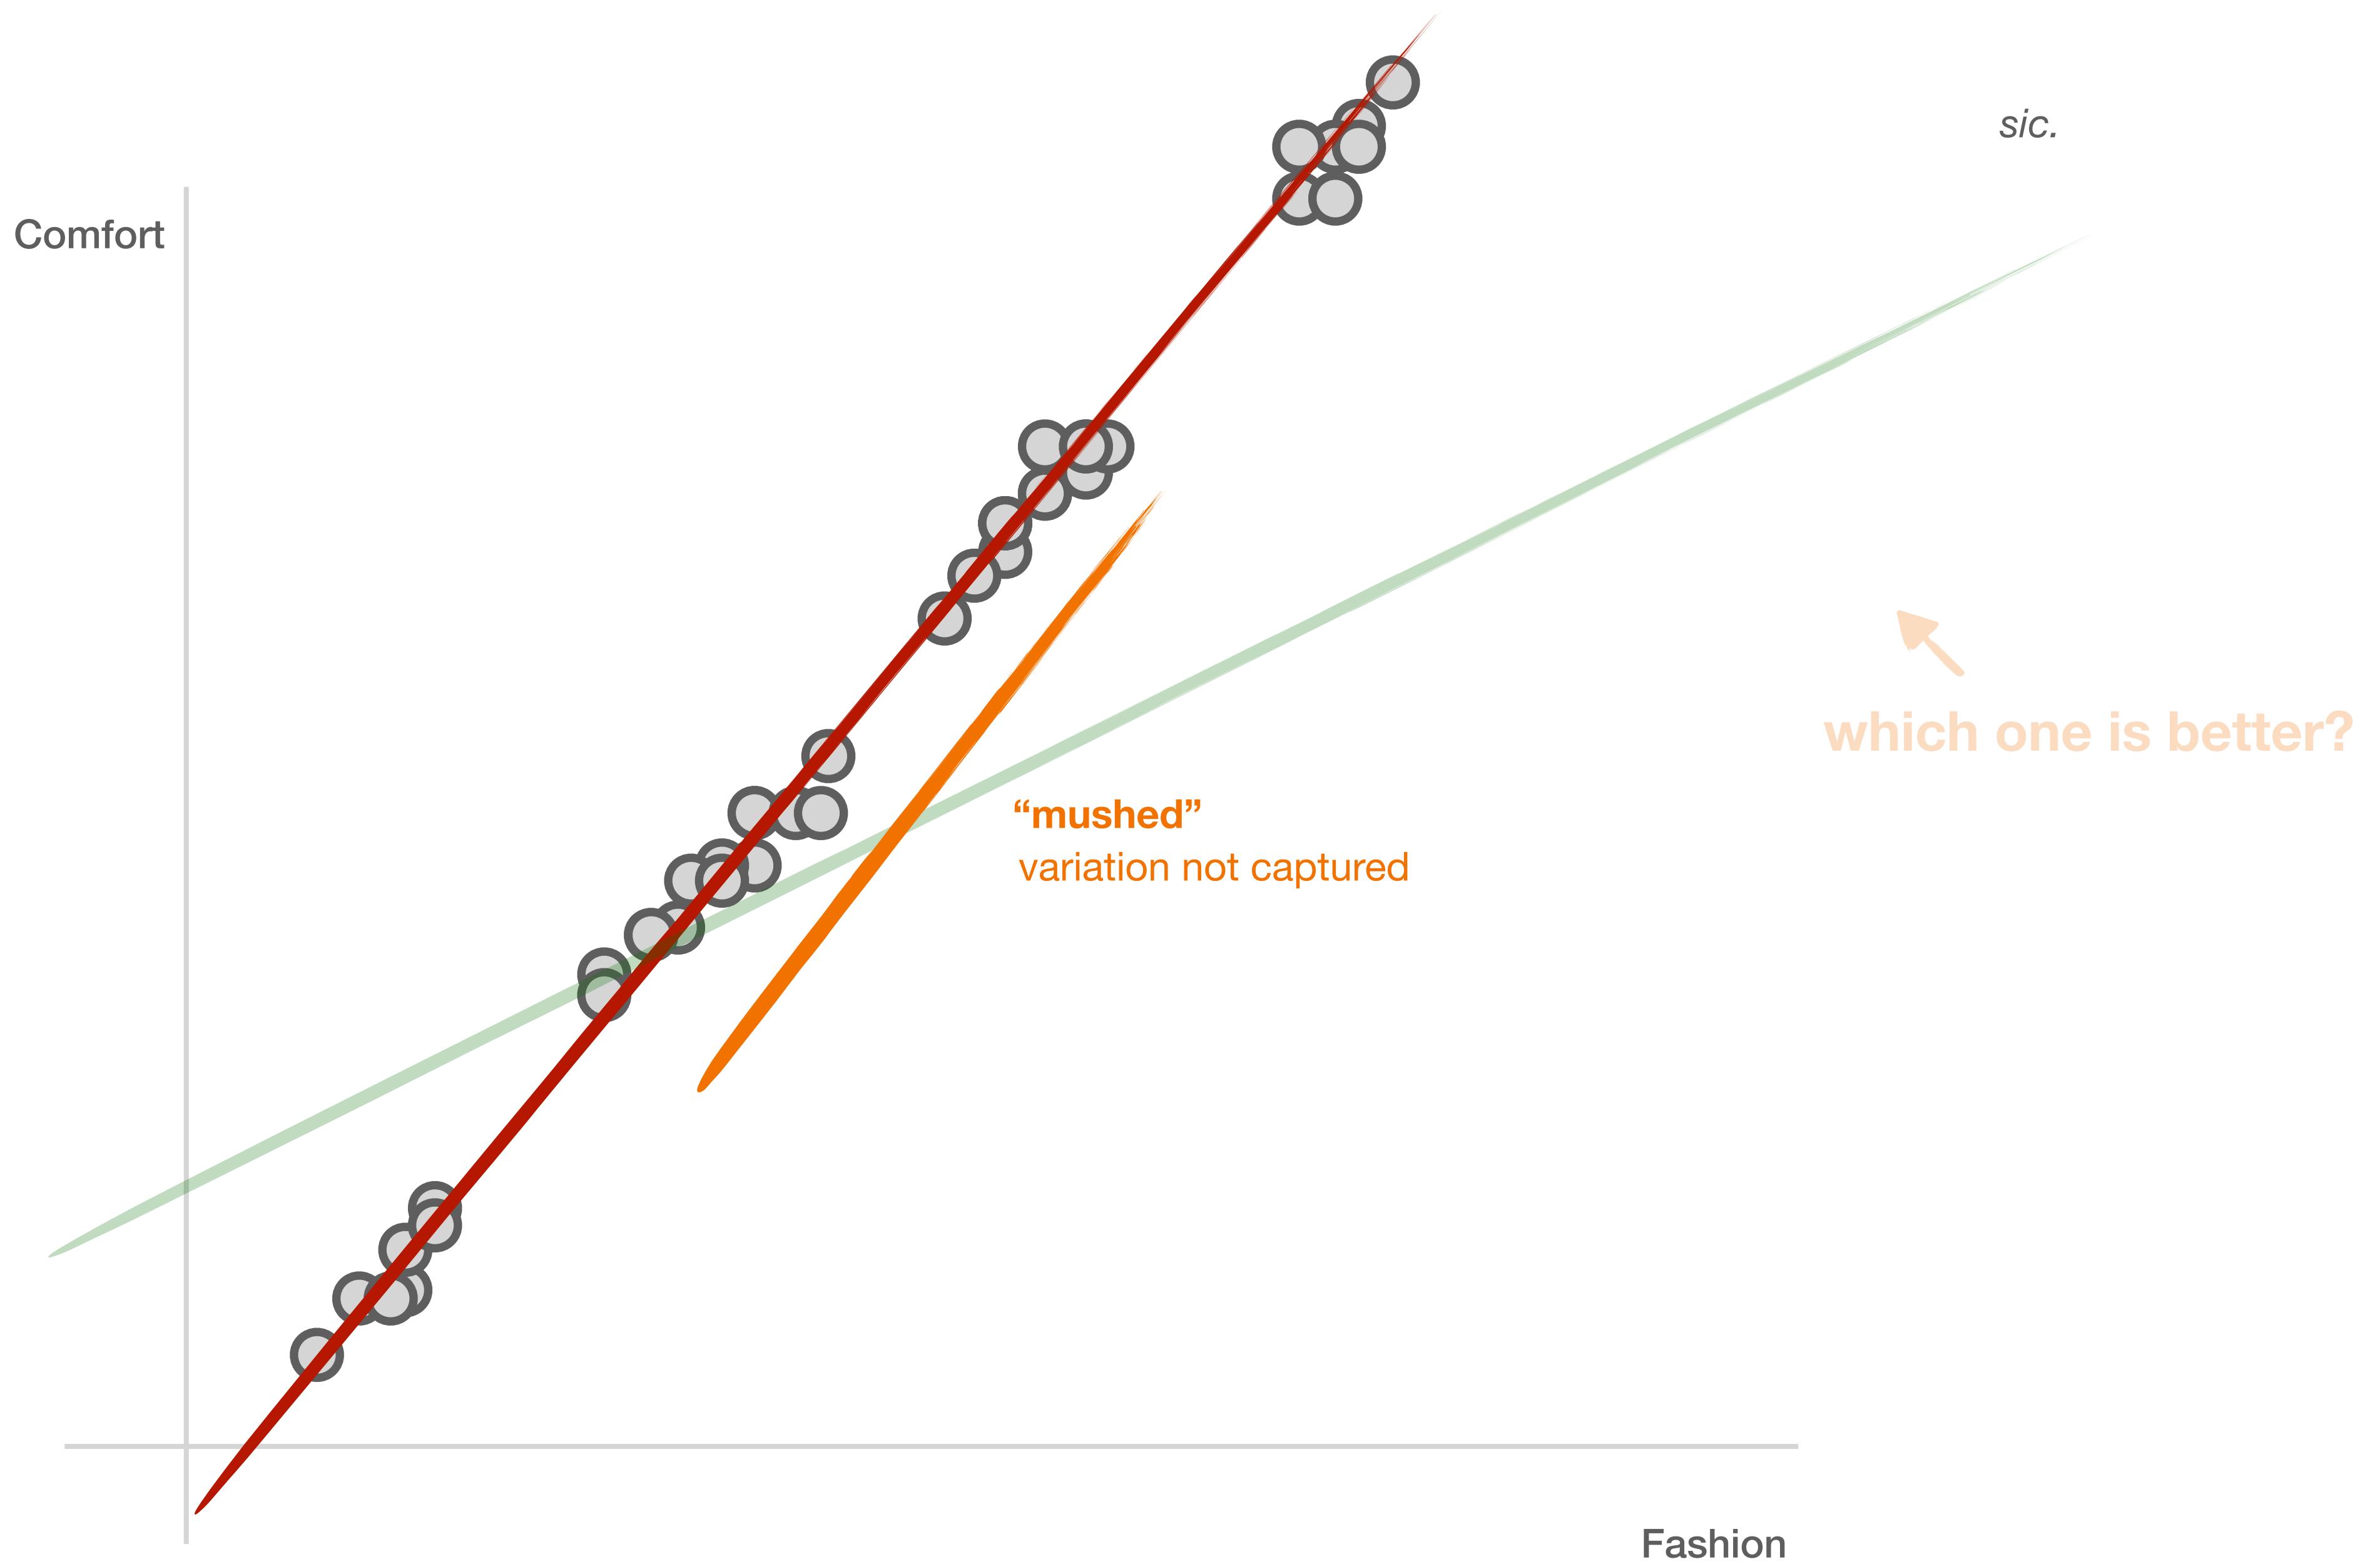
how to pick a good line?



how to pick a good line?



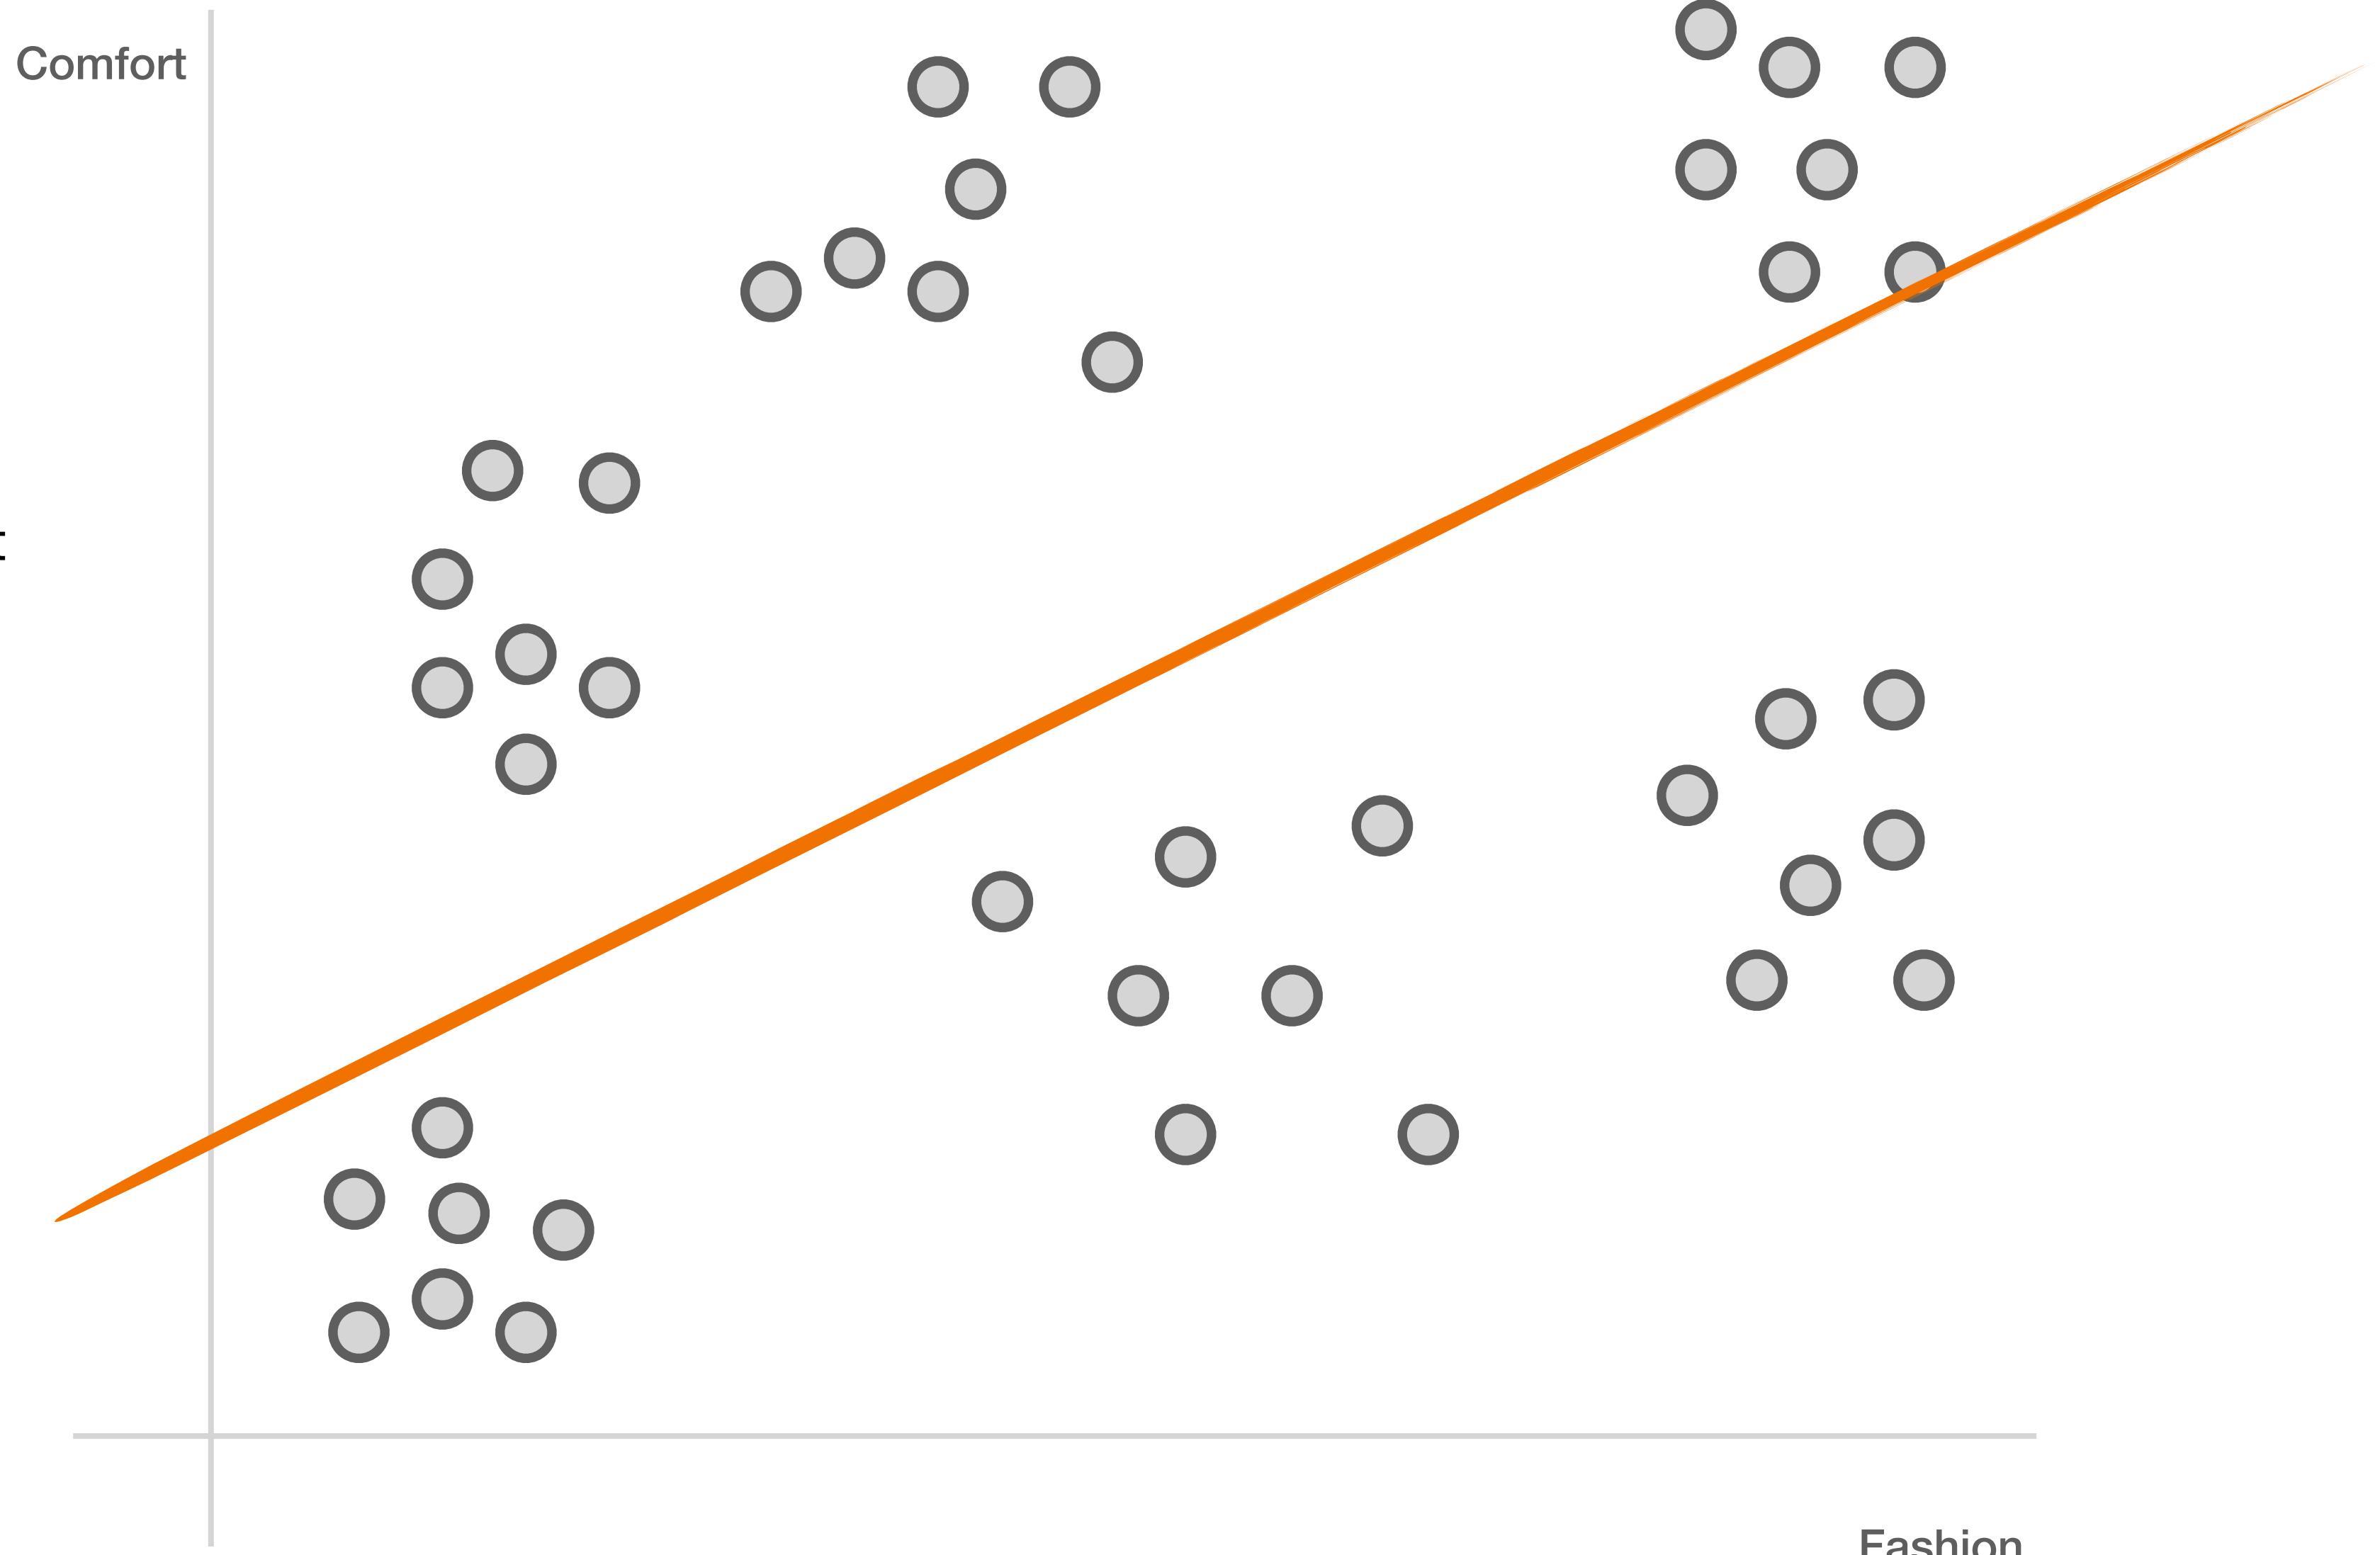
how to pick a good line?



- a good “projection” captures the **variation** in the data

how to pick a good line?

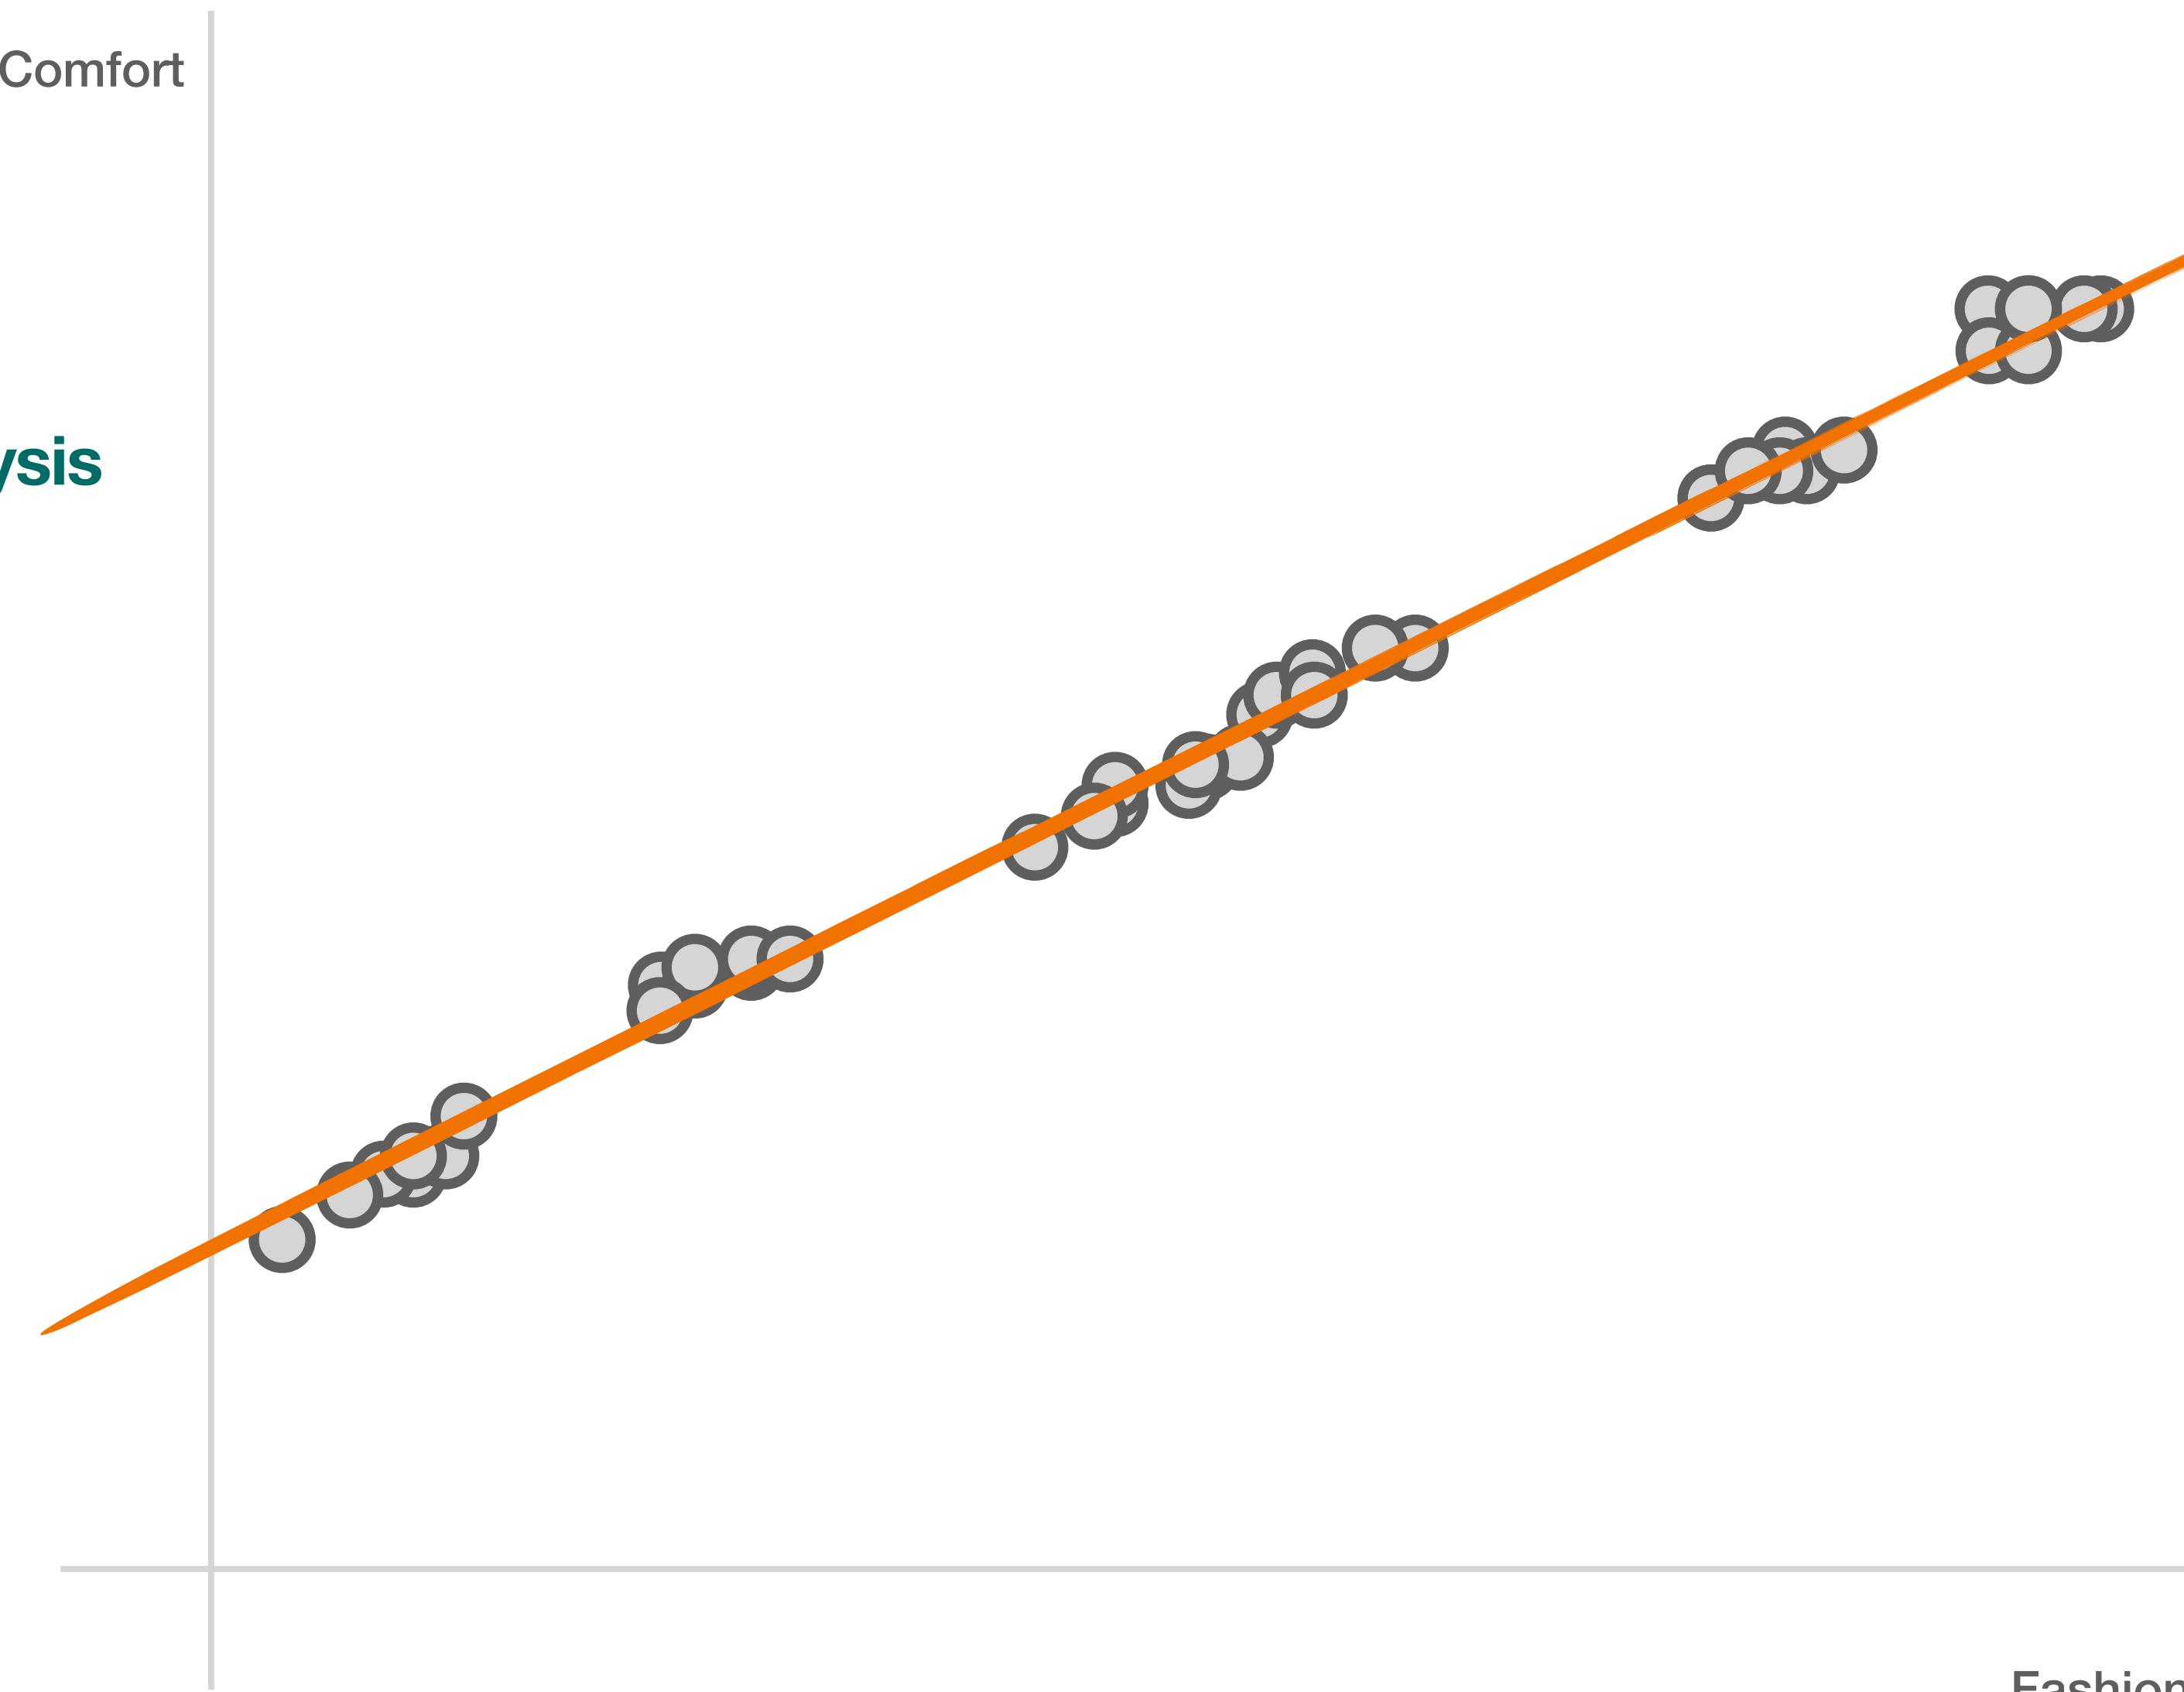
- Find a good line (**basis**) that **maximizes variation**



how to pick a good line?

principle component analysis

- Find a good line (**basis**) that **maximizes variation**
- **Project samples down**



But, can we formalize it?



Principle Component Analysis

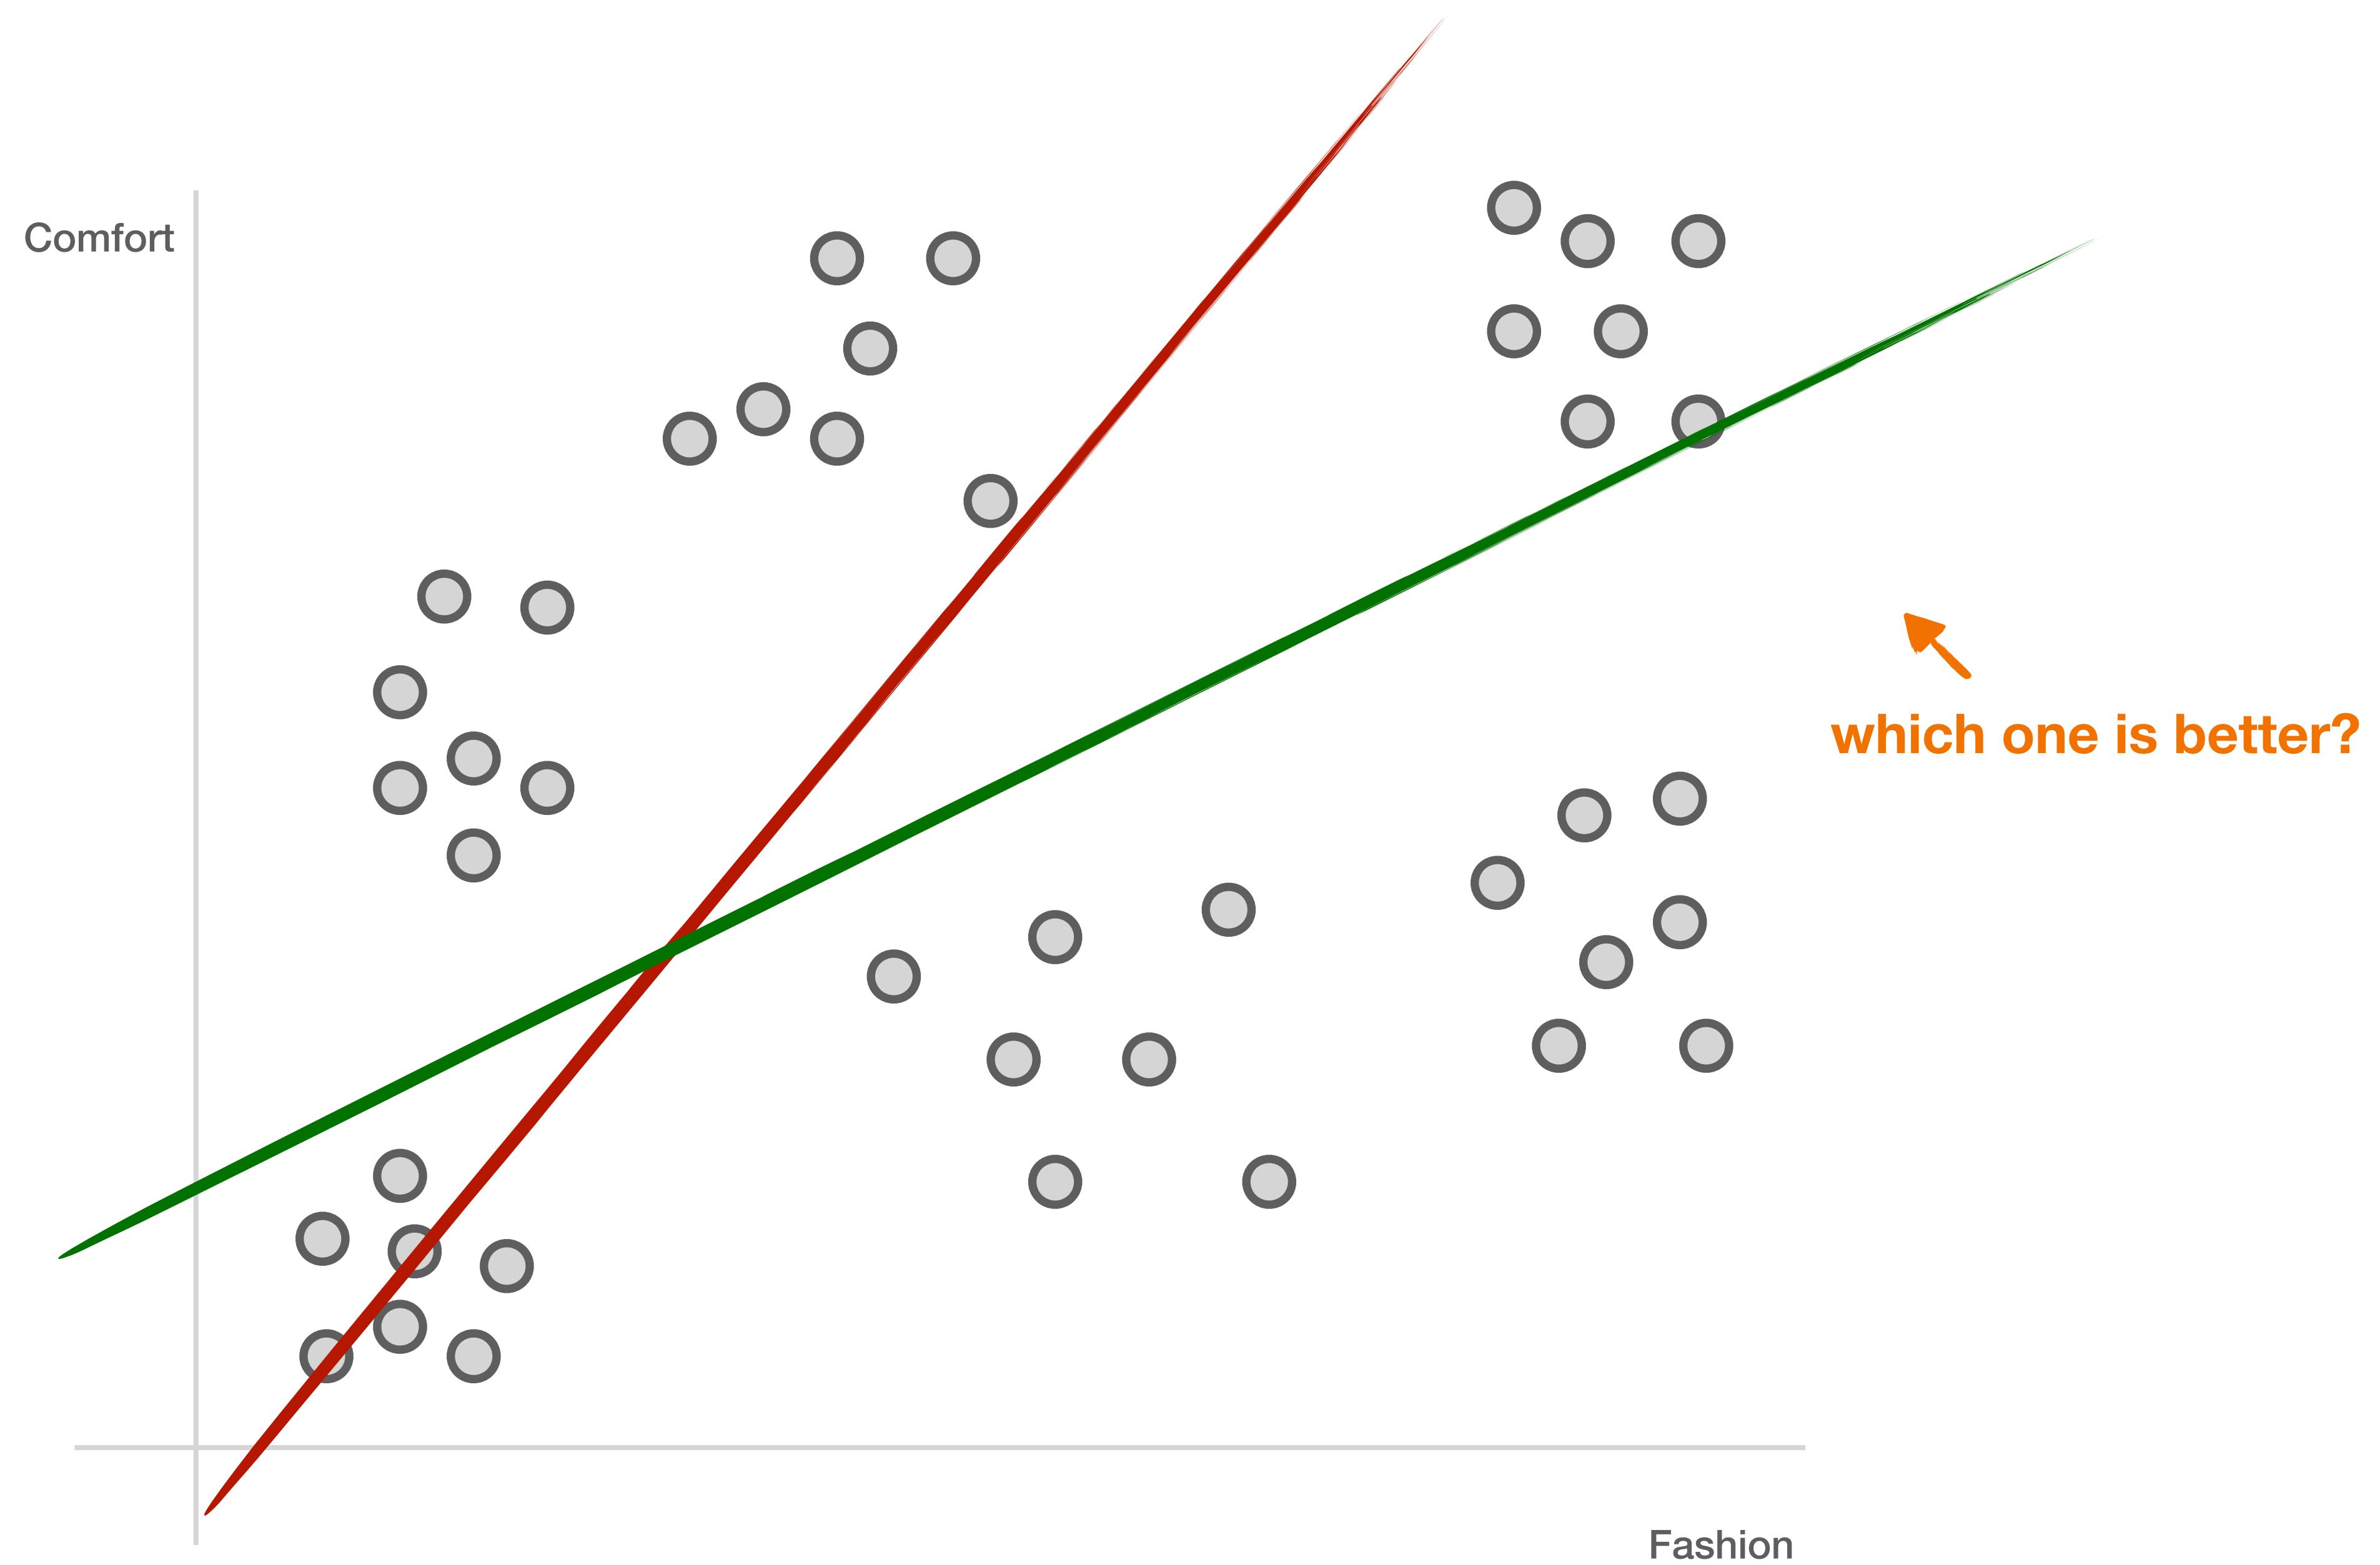
First Component: maximize the explained variance

$$\mathbf{w}_{(1)} = \arg \max_{\|\mathbf{w}\|=1} \left\{ \sum_i (\mathbf{x}_{(i)} \cdot \mathbf{w})^2 \right\}$$

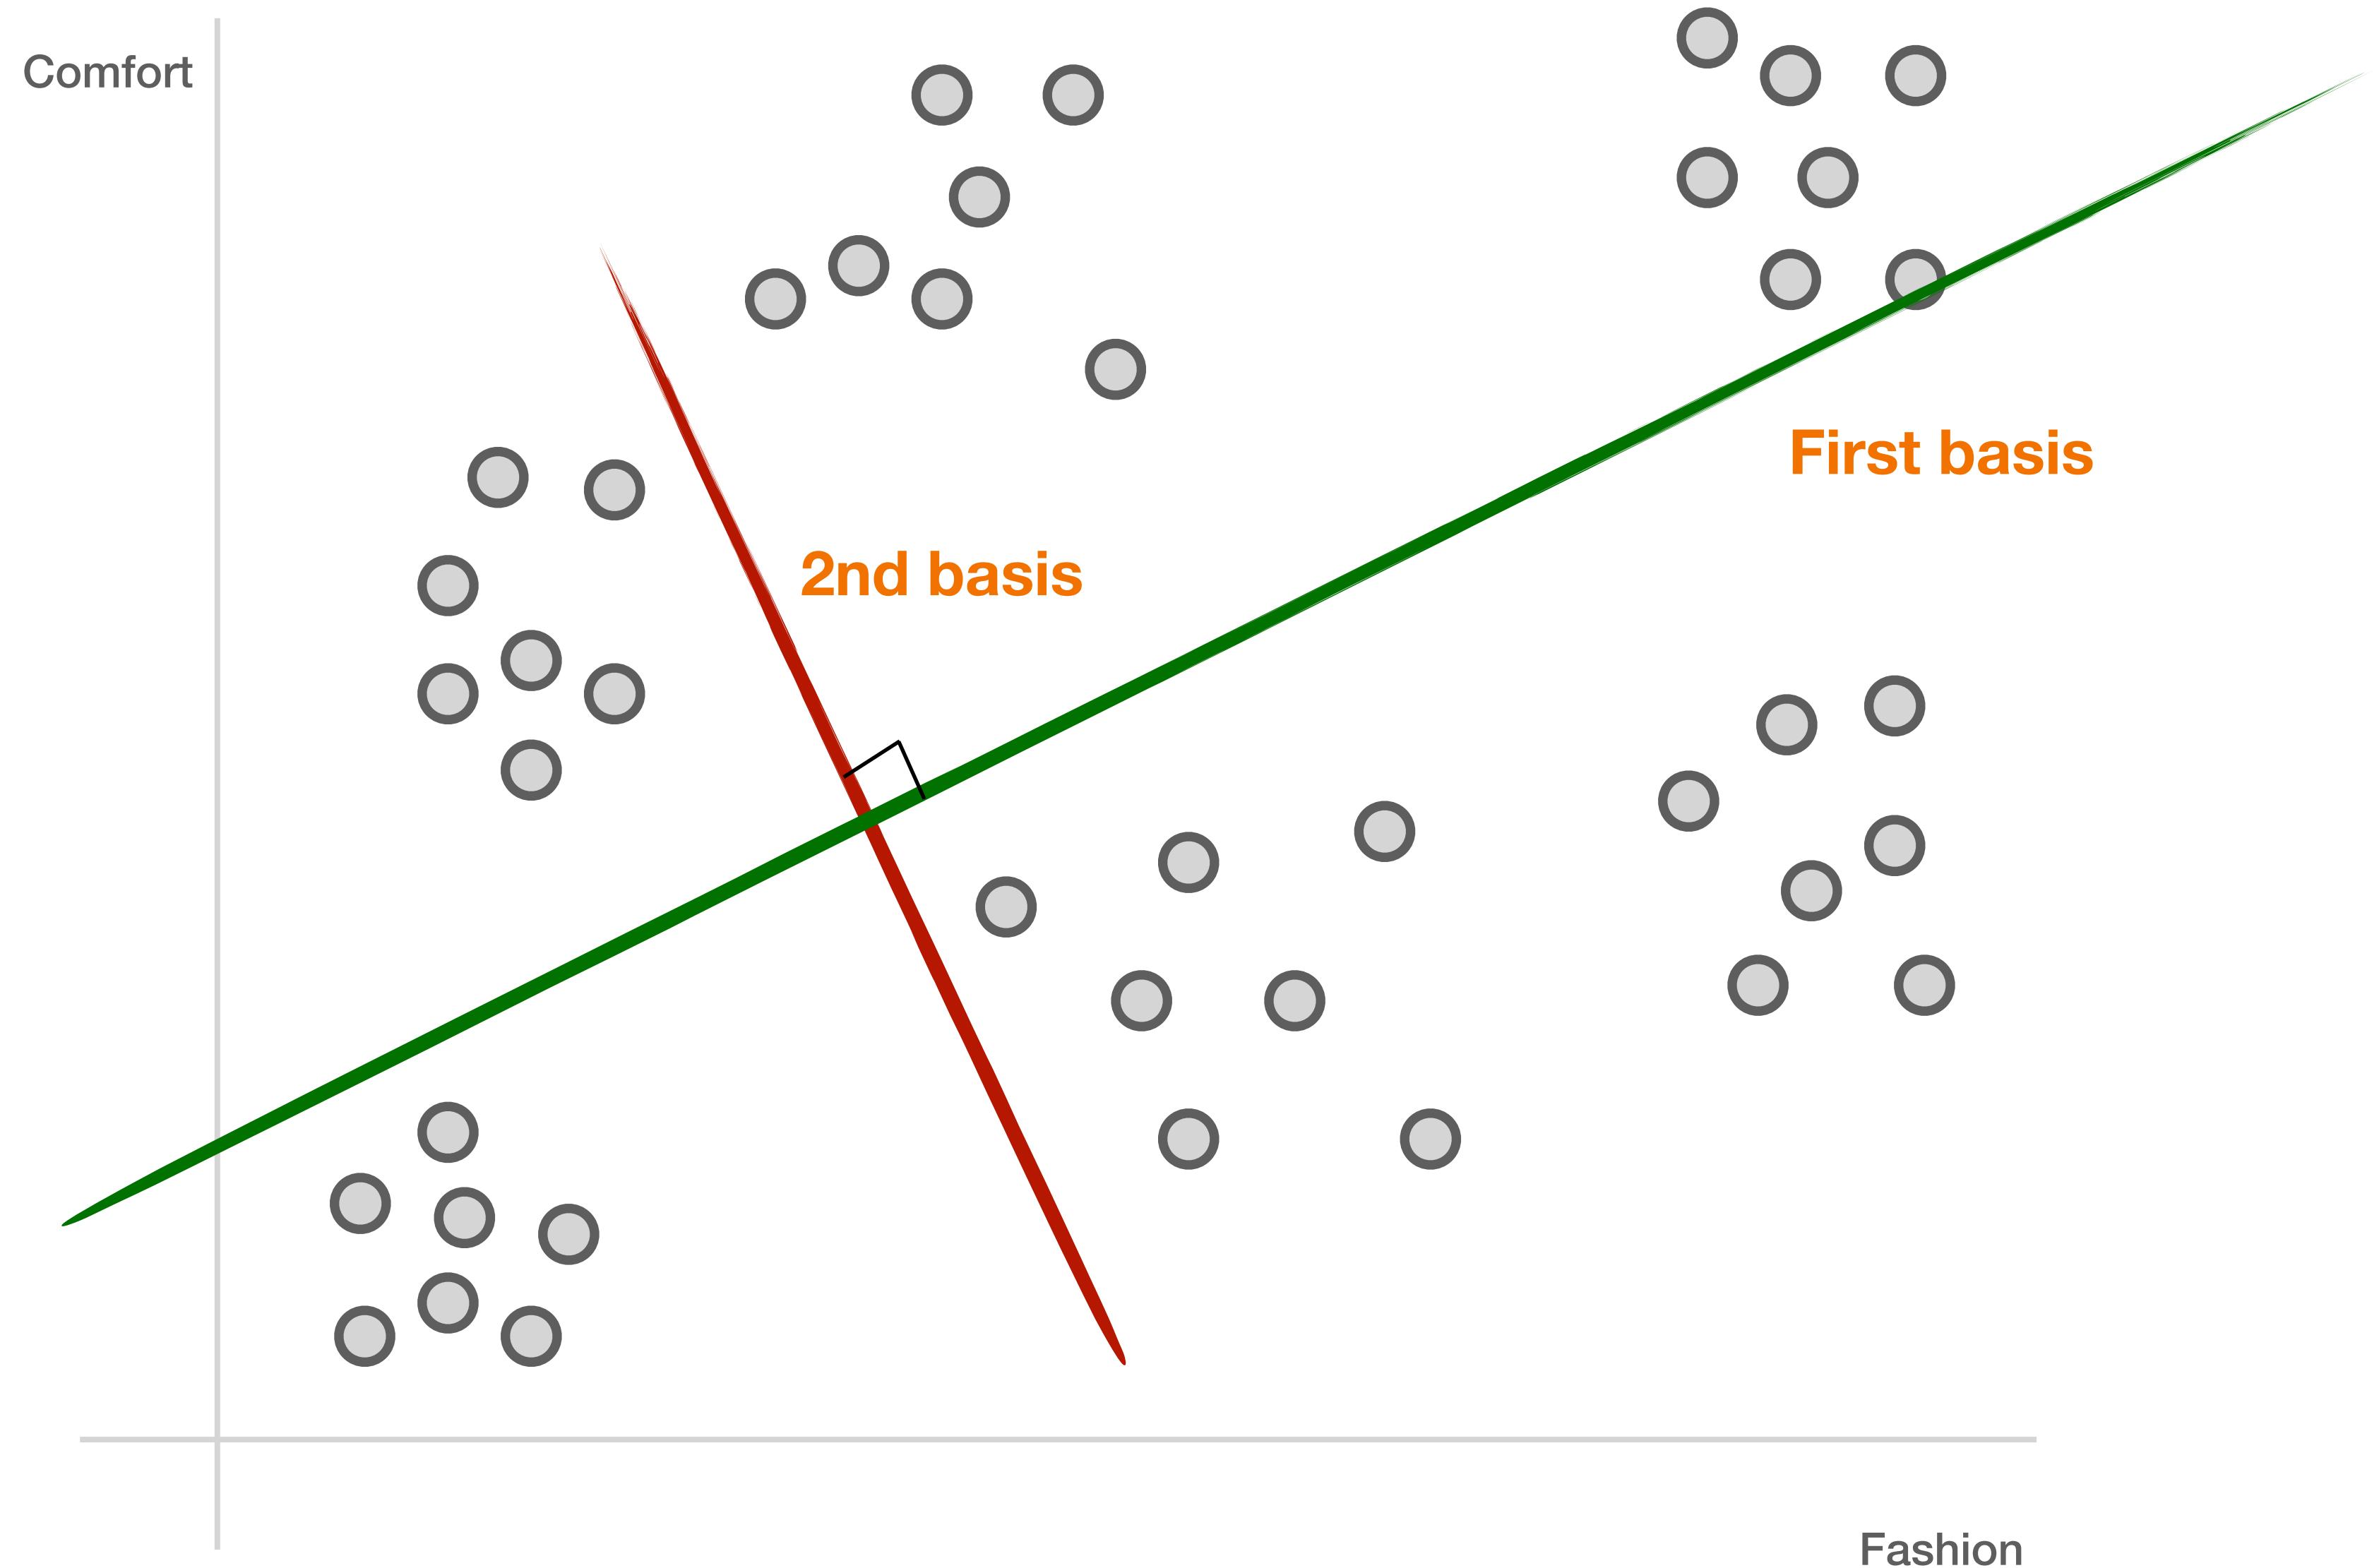
Further Components: maximize the explained variance in remainders

$$\mathbf{w}_{(k)} = \arg \max_{\|\mathbf{w}\|=1} \left\{ \left\| \hat{\mathbf{X}}_k \mathbf{w} \right\|^2 \right\} : \quad \hat{\mathbf{X}}_k = \mathbf{X} - \sum_{s=1}^{k-1} \mathbf{X} \mathbf{w}_{(s)} \mathbf{w}_{(s)}^\top$$

First Component



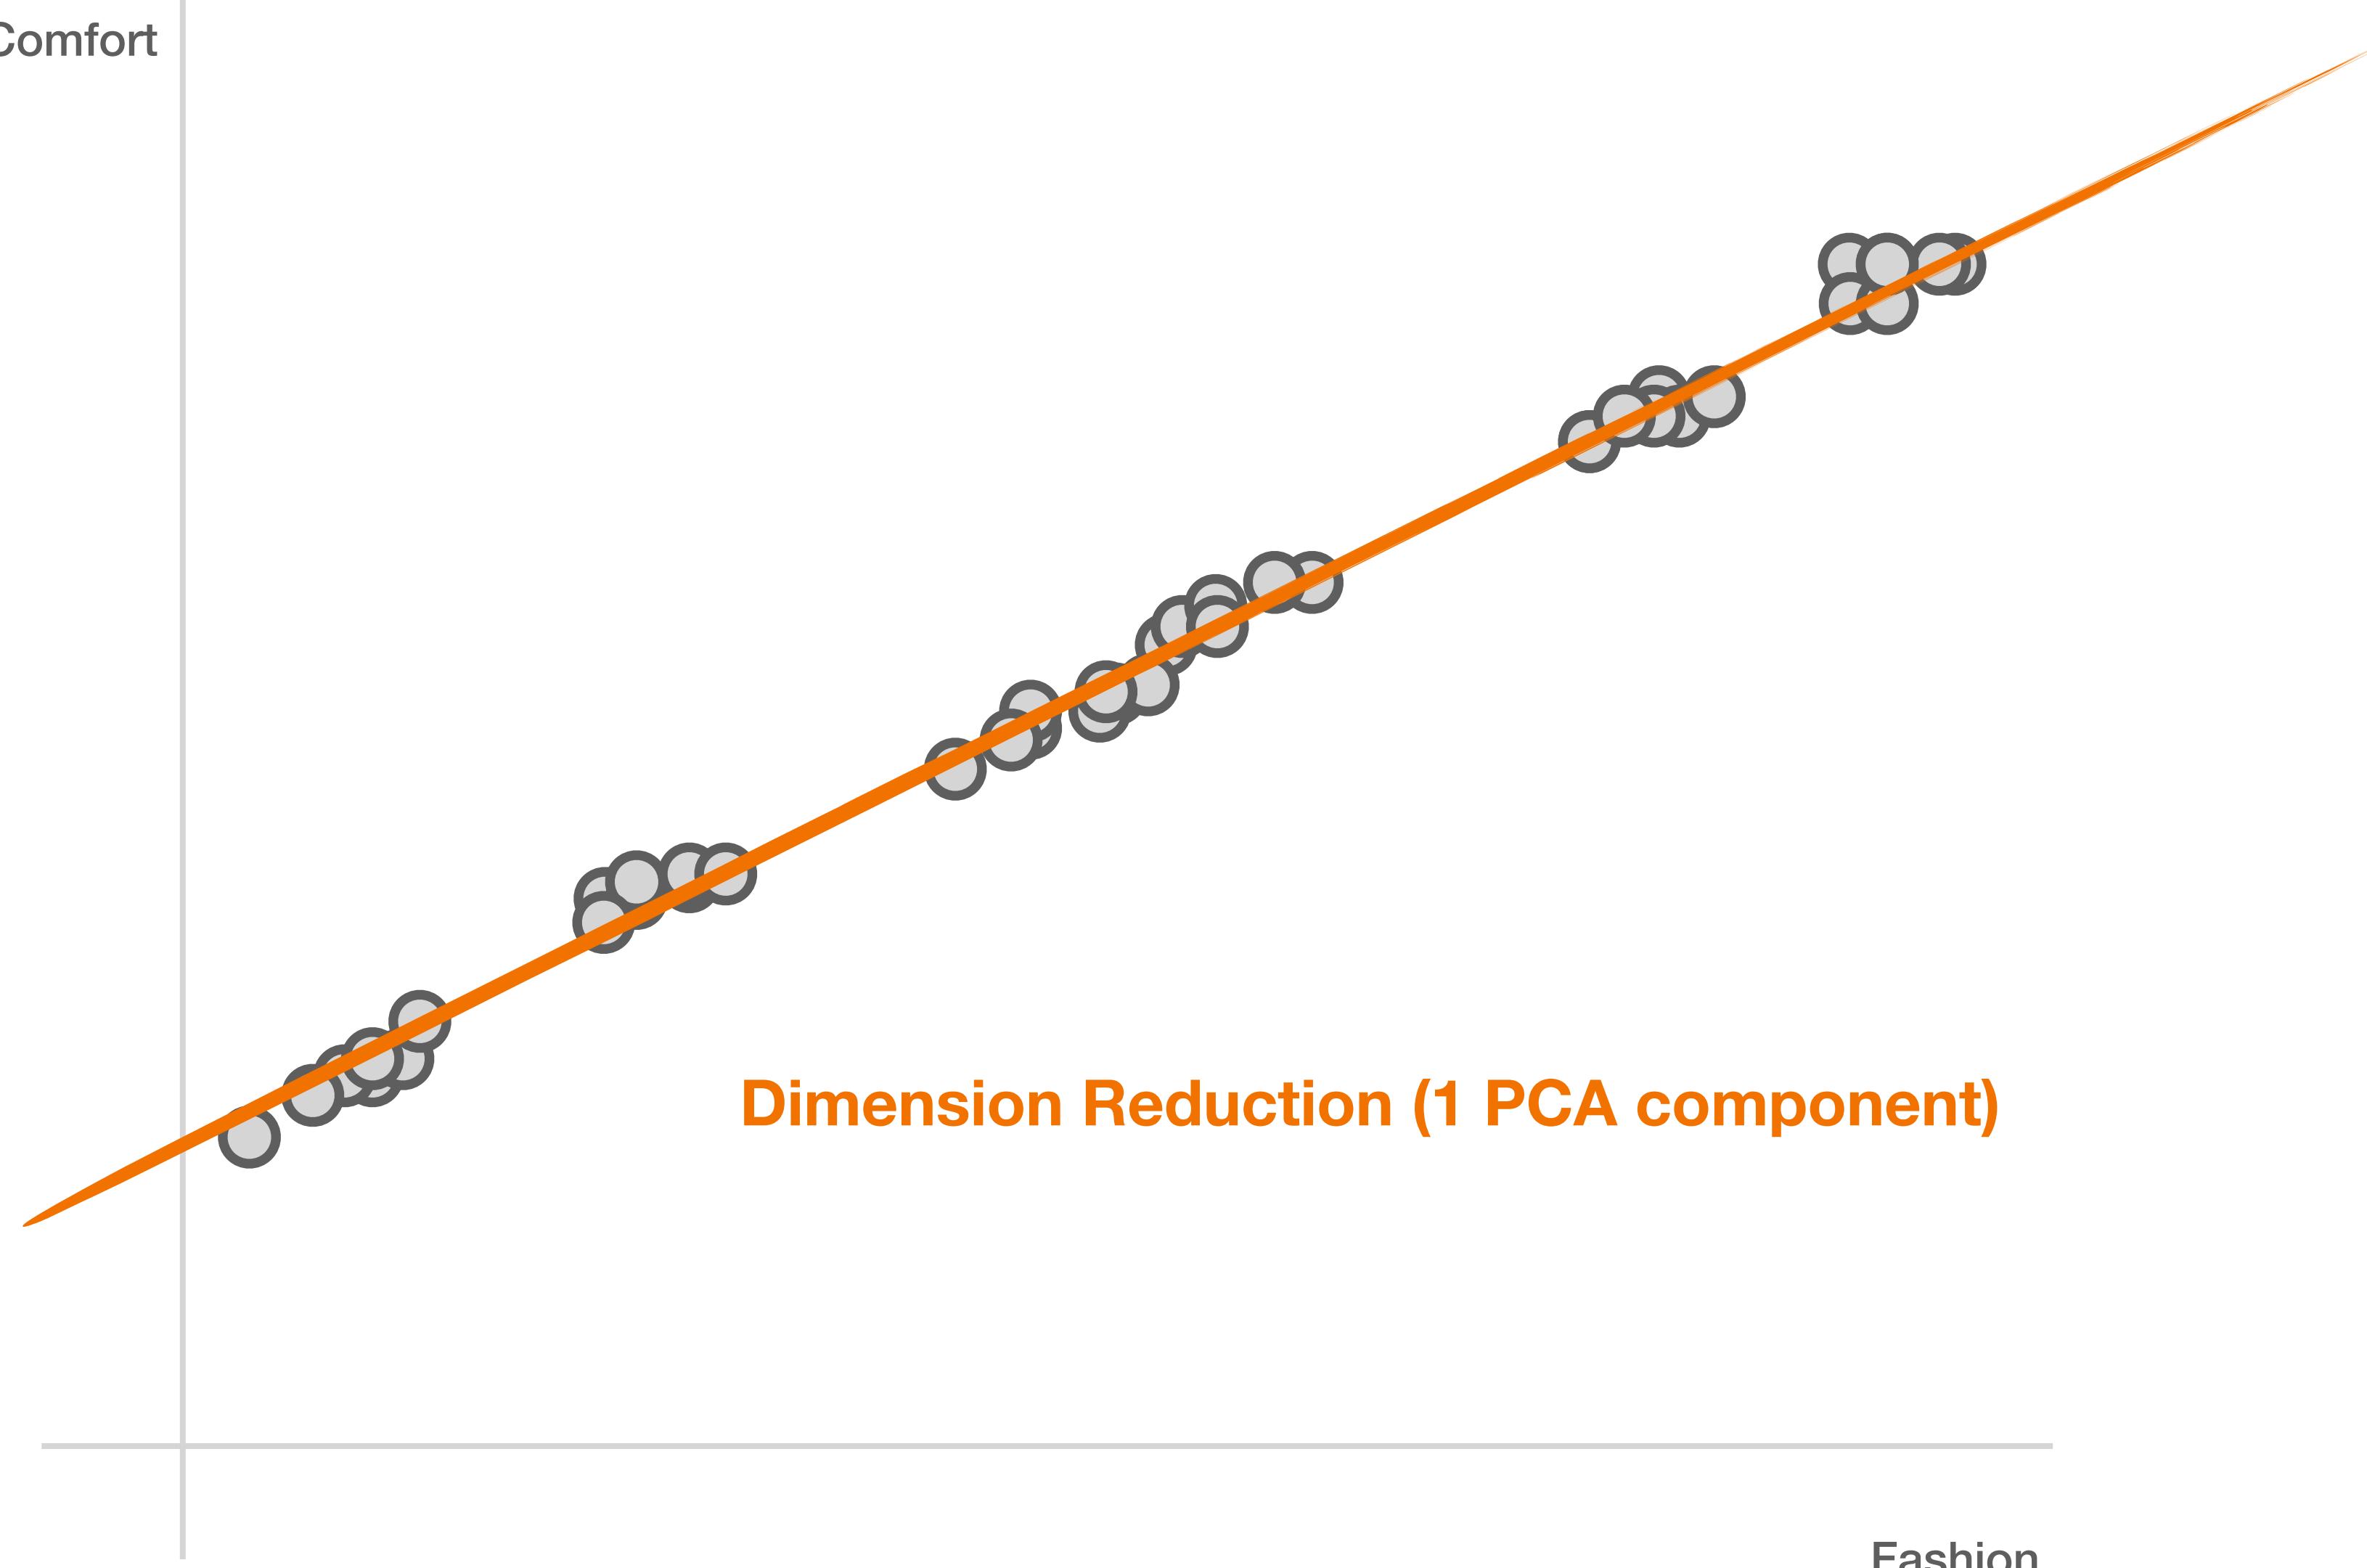
PCA



Comfort

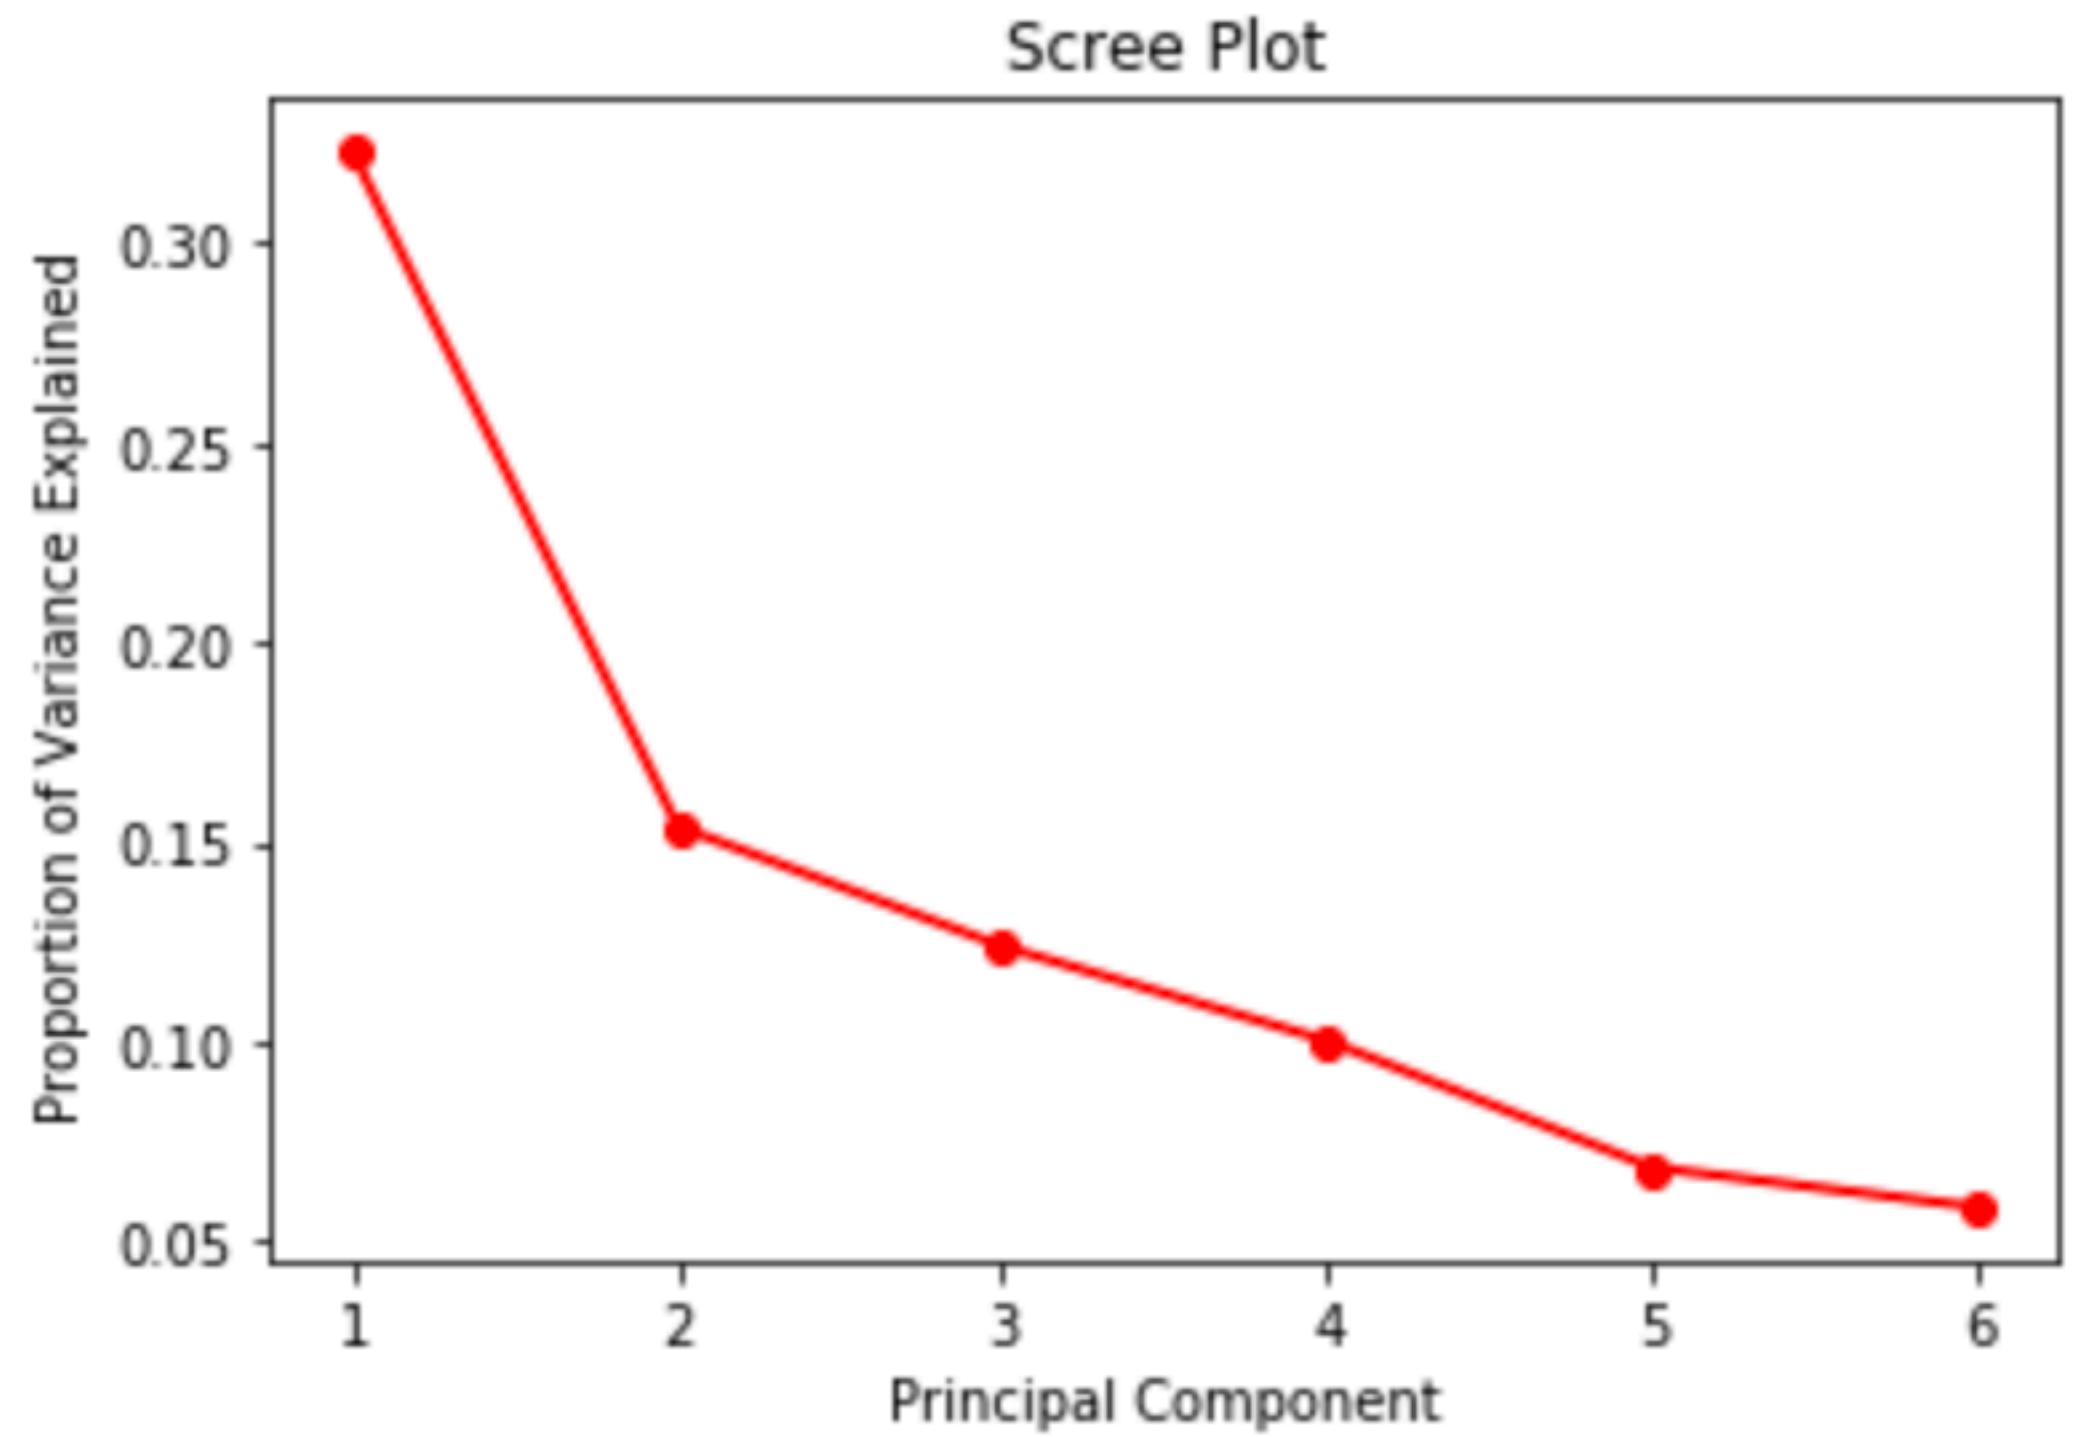
Fashion

Dimension Reduction (1 PCA component)



- Linear transformation of original features
- Dimension reduction
- Compression
- Denoise
- Lose interpretability

- How to choose k (hyperparameter)



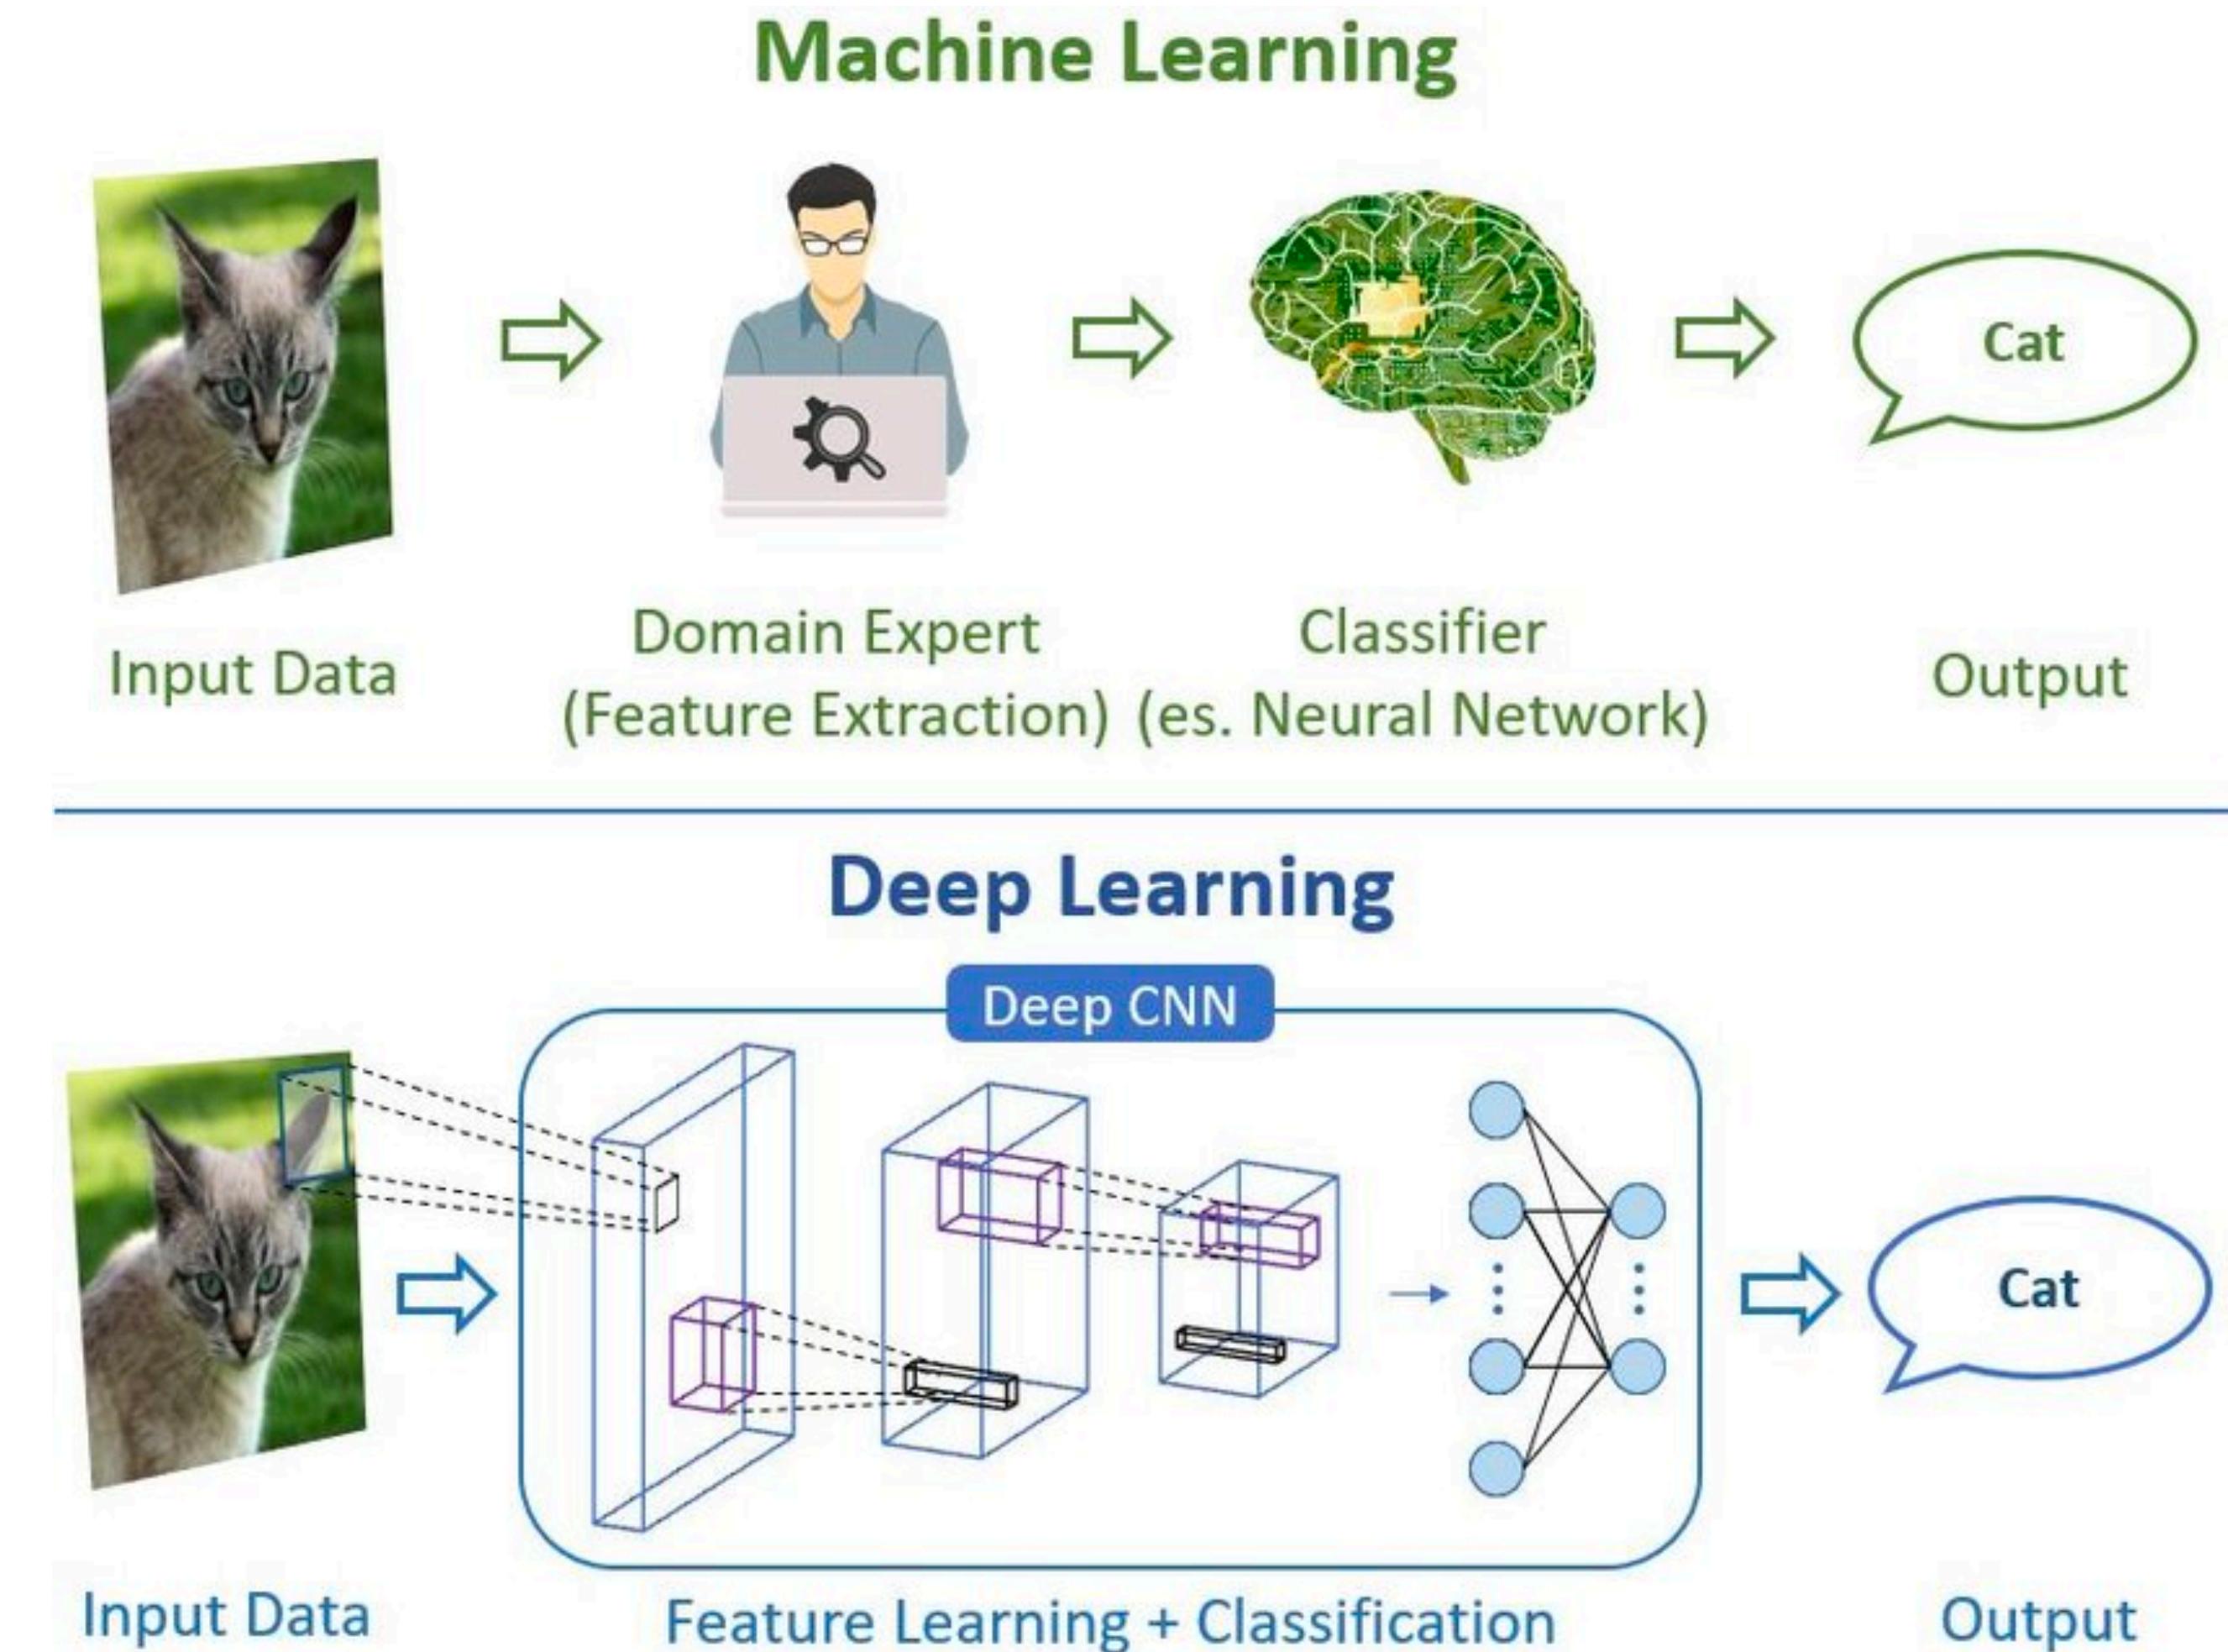
POTENTIAL PITFALLS

Things that can go wrong

- Inconsistent preprocessing (e.g., different scaling/normalization)
- Data leakage (e.g., temporal or mixing subjects)

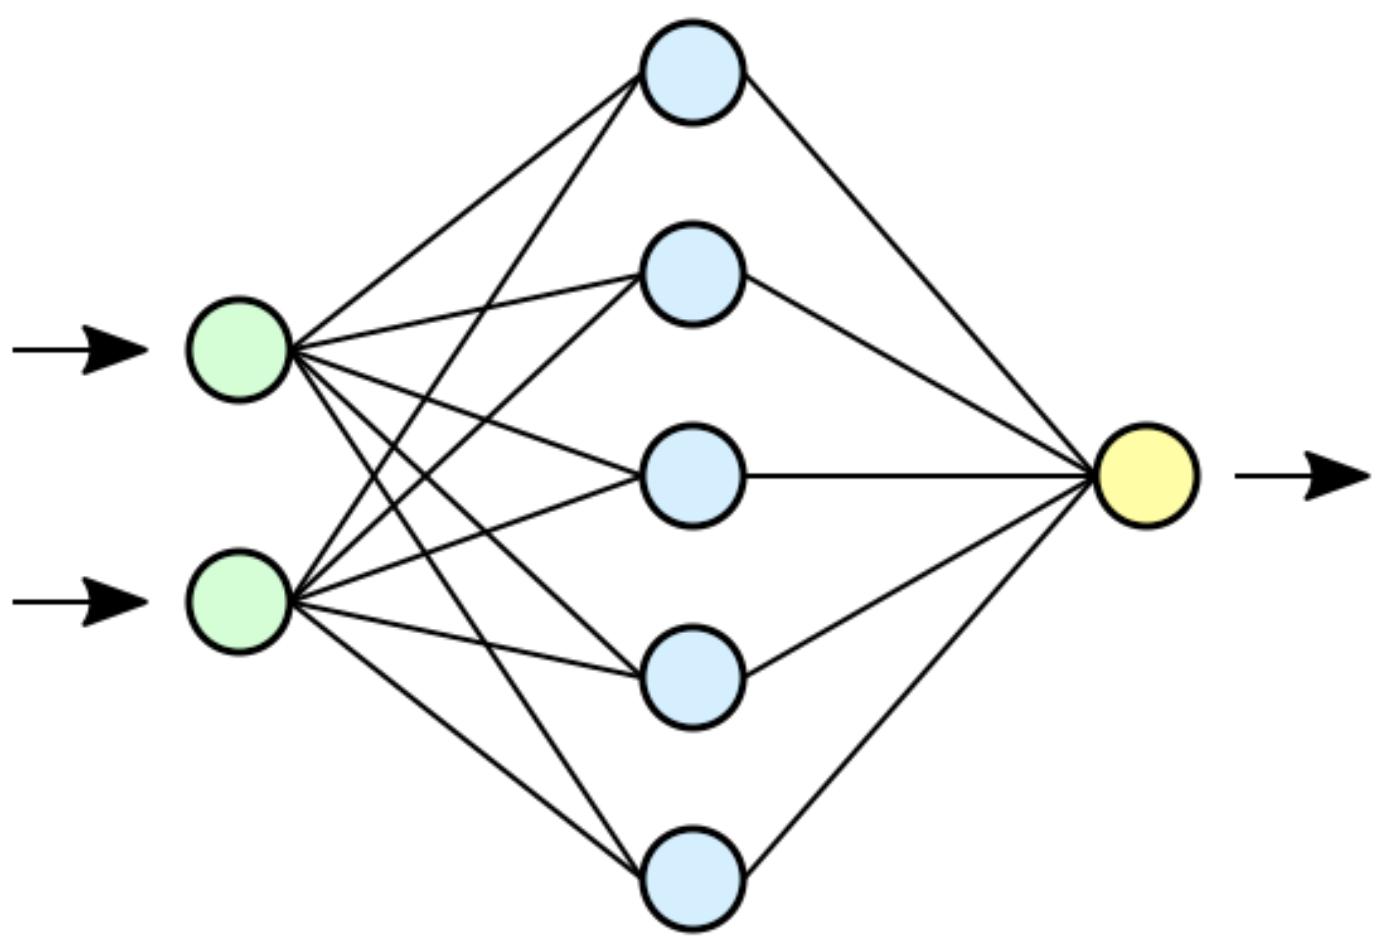
- Model is used on test data that has changed
- Selecting appropriate metrics (e.g., is 99% accuracy good enough?)
- Hidden confounders (e.g., golf is correlated with heart attacks)
- Spurious correlations (e.g., hospital ID on images)
- Performance on subgroups may be missing

Classic vs. Deep ML



Cool, so...
what's next?

Cool, so... what's next?



Neural Networks!

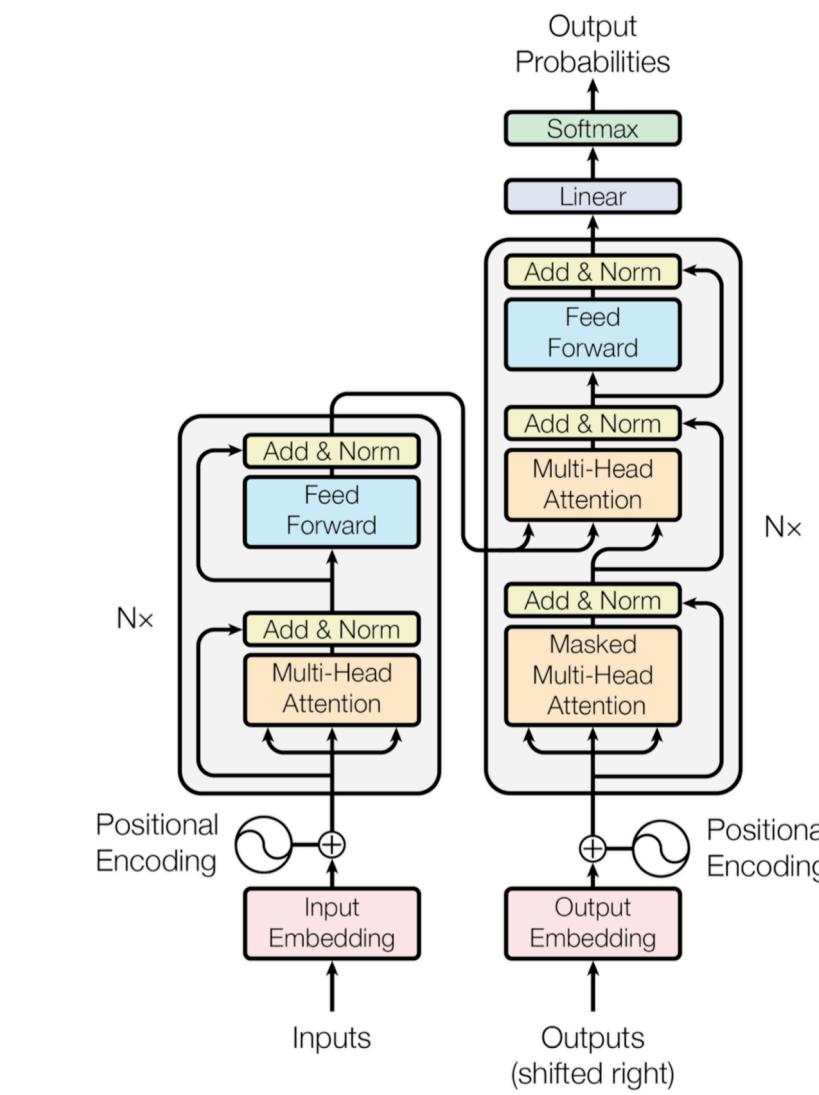
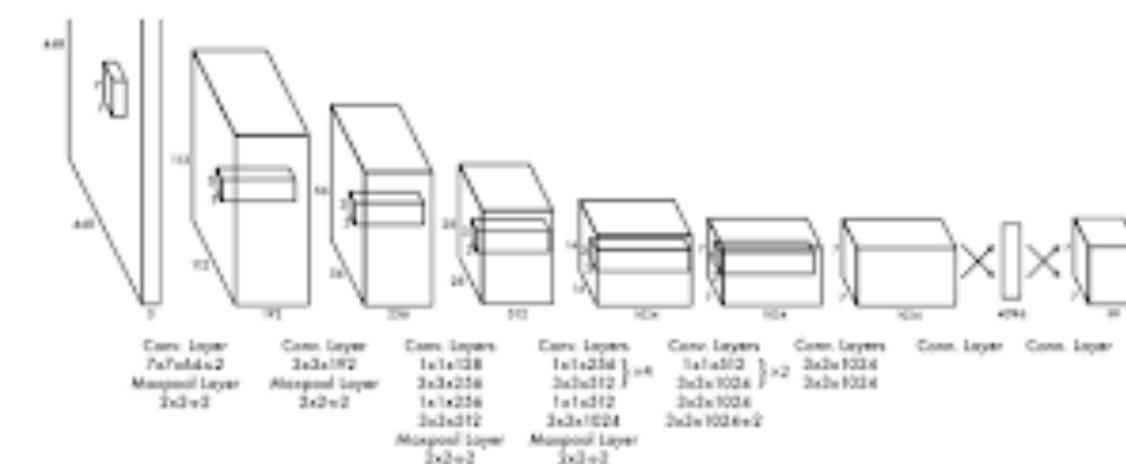


Figure 1: The Transformer - model architecture.

Natural Language Processing



Computer Vision

ENSEMBLE

Wisdom of the crowd

- Guessing the weight of a steer (Sir Francis Galton)
- Key components:
 - Base models with **diversity**
 - Infusion algorithms to integrate base models
- Bagging
- Boosting

Bagging

- Build several instances of a classifier using a subset of the original training data
 - Aggregate (Averaging) the results
 - Reduce variance
 - Ex: random forest
-
- <https://scikit-learn.org/stable/modules/ensemble.html#bagging-meta-estimator>

Boosting

- Fit a sequence of weak learners on repeatedly modified versions of the data.
- A weighted majority vote
- Ex: Adaboost, Gradient boosting

SELF-SUPERVISED LEARNING

Exploiting Unlabeled Data

A lot of unlabeled data is plentiful and cheap, e.g.,

- Document off the web
- Speech samples
- Images and video

But labeling can be expensive

Self-supervised learning



Text Corpus

Nothing is impossible.
Even the word
impossible
says I'm possible



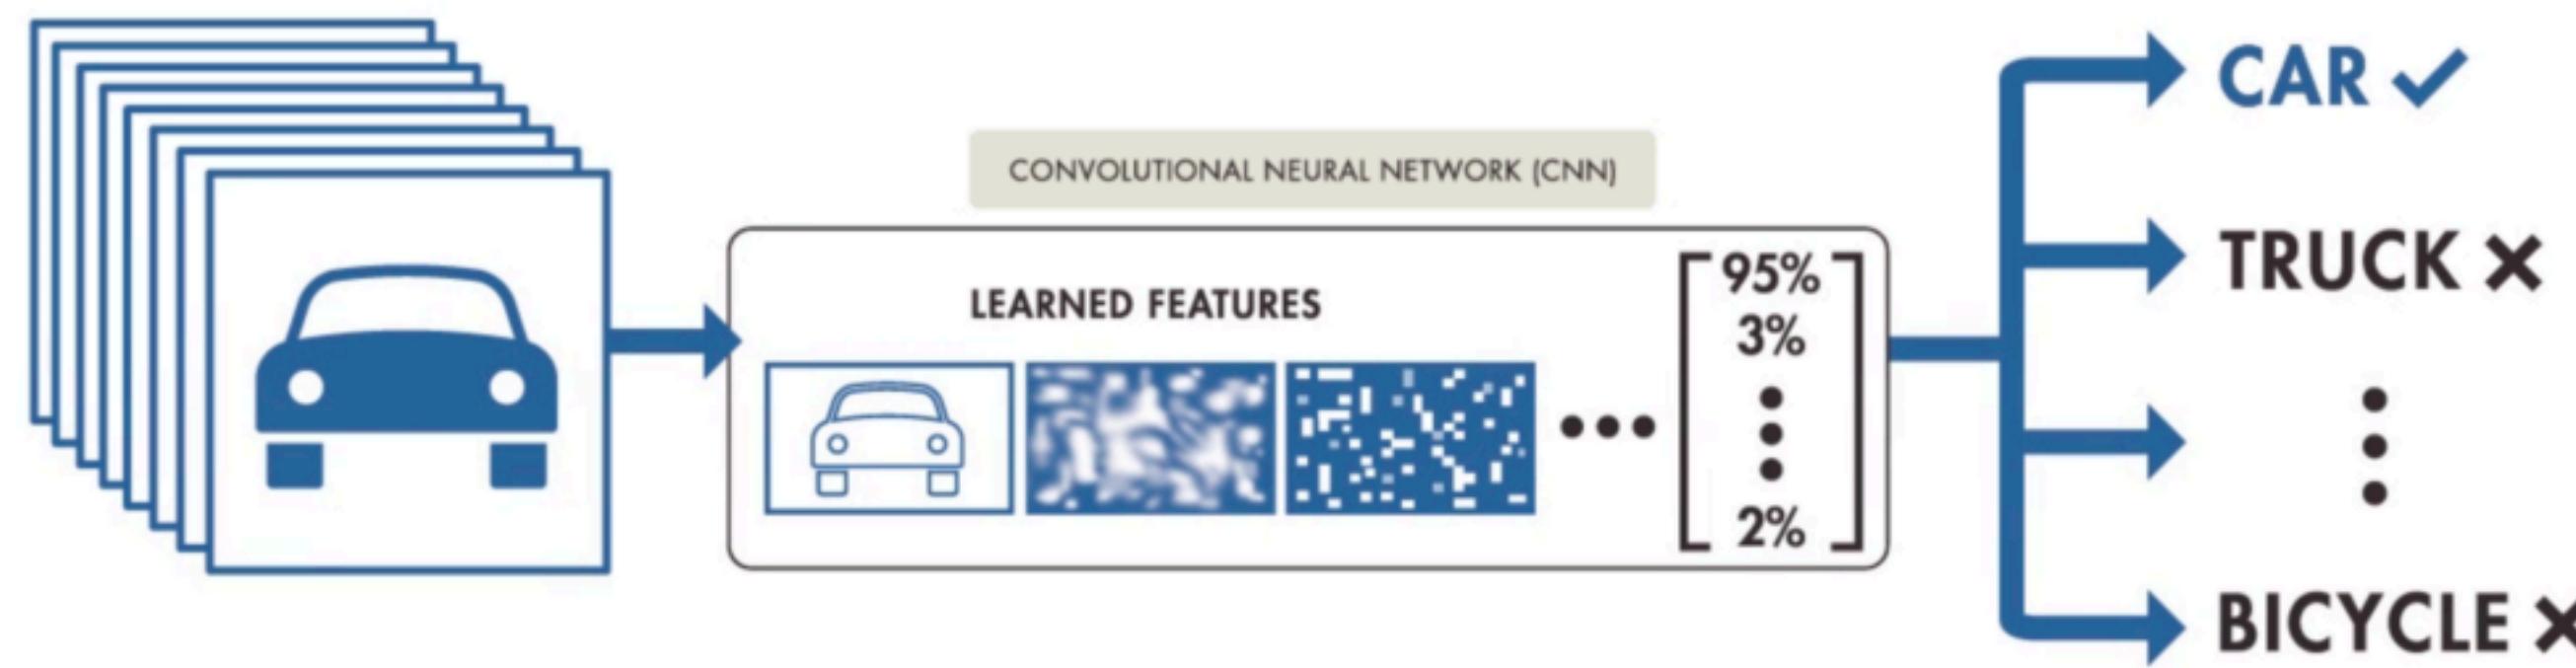
Task: Predict from past

Nothing
Nothing is
Nothing is impossible
...

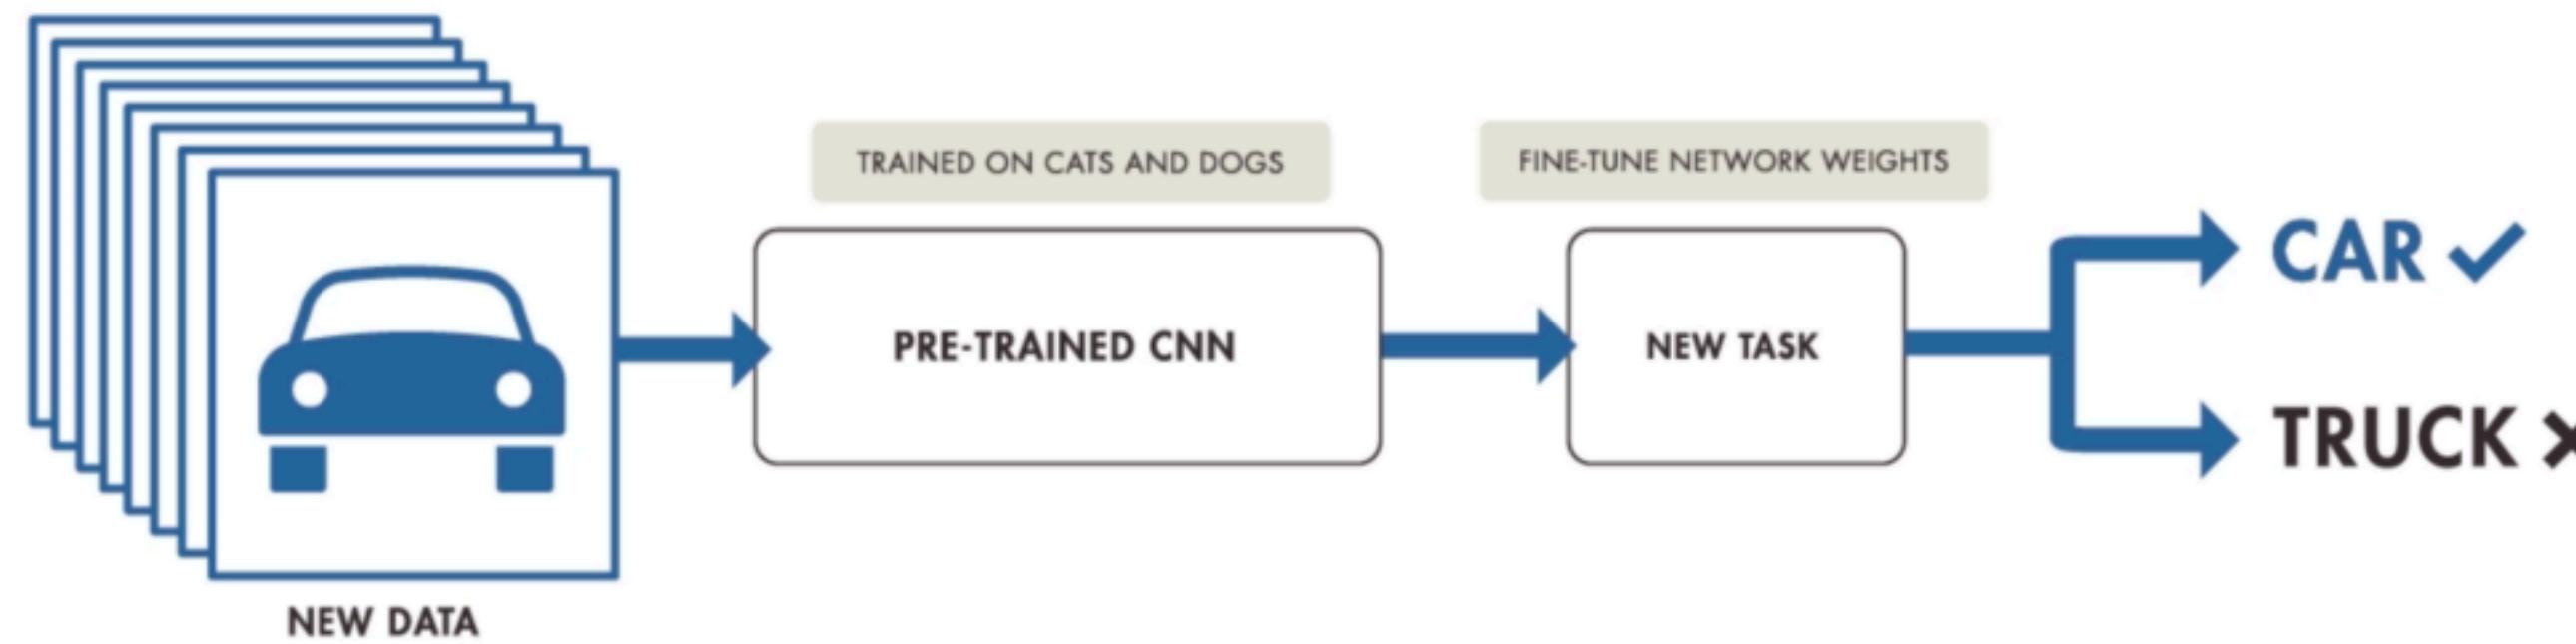
A critical step in training large language models

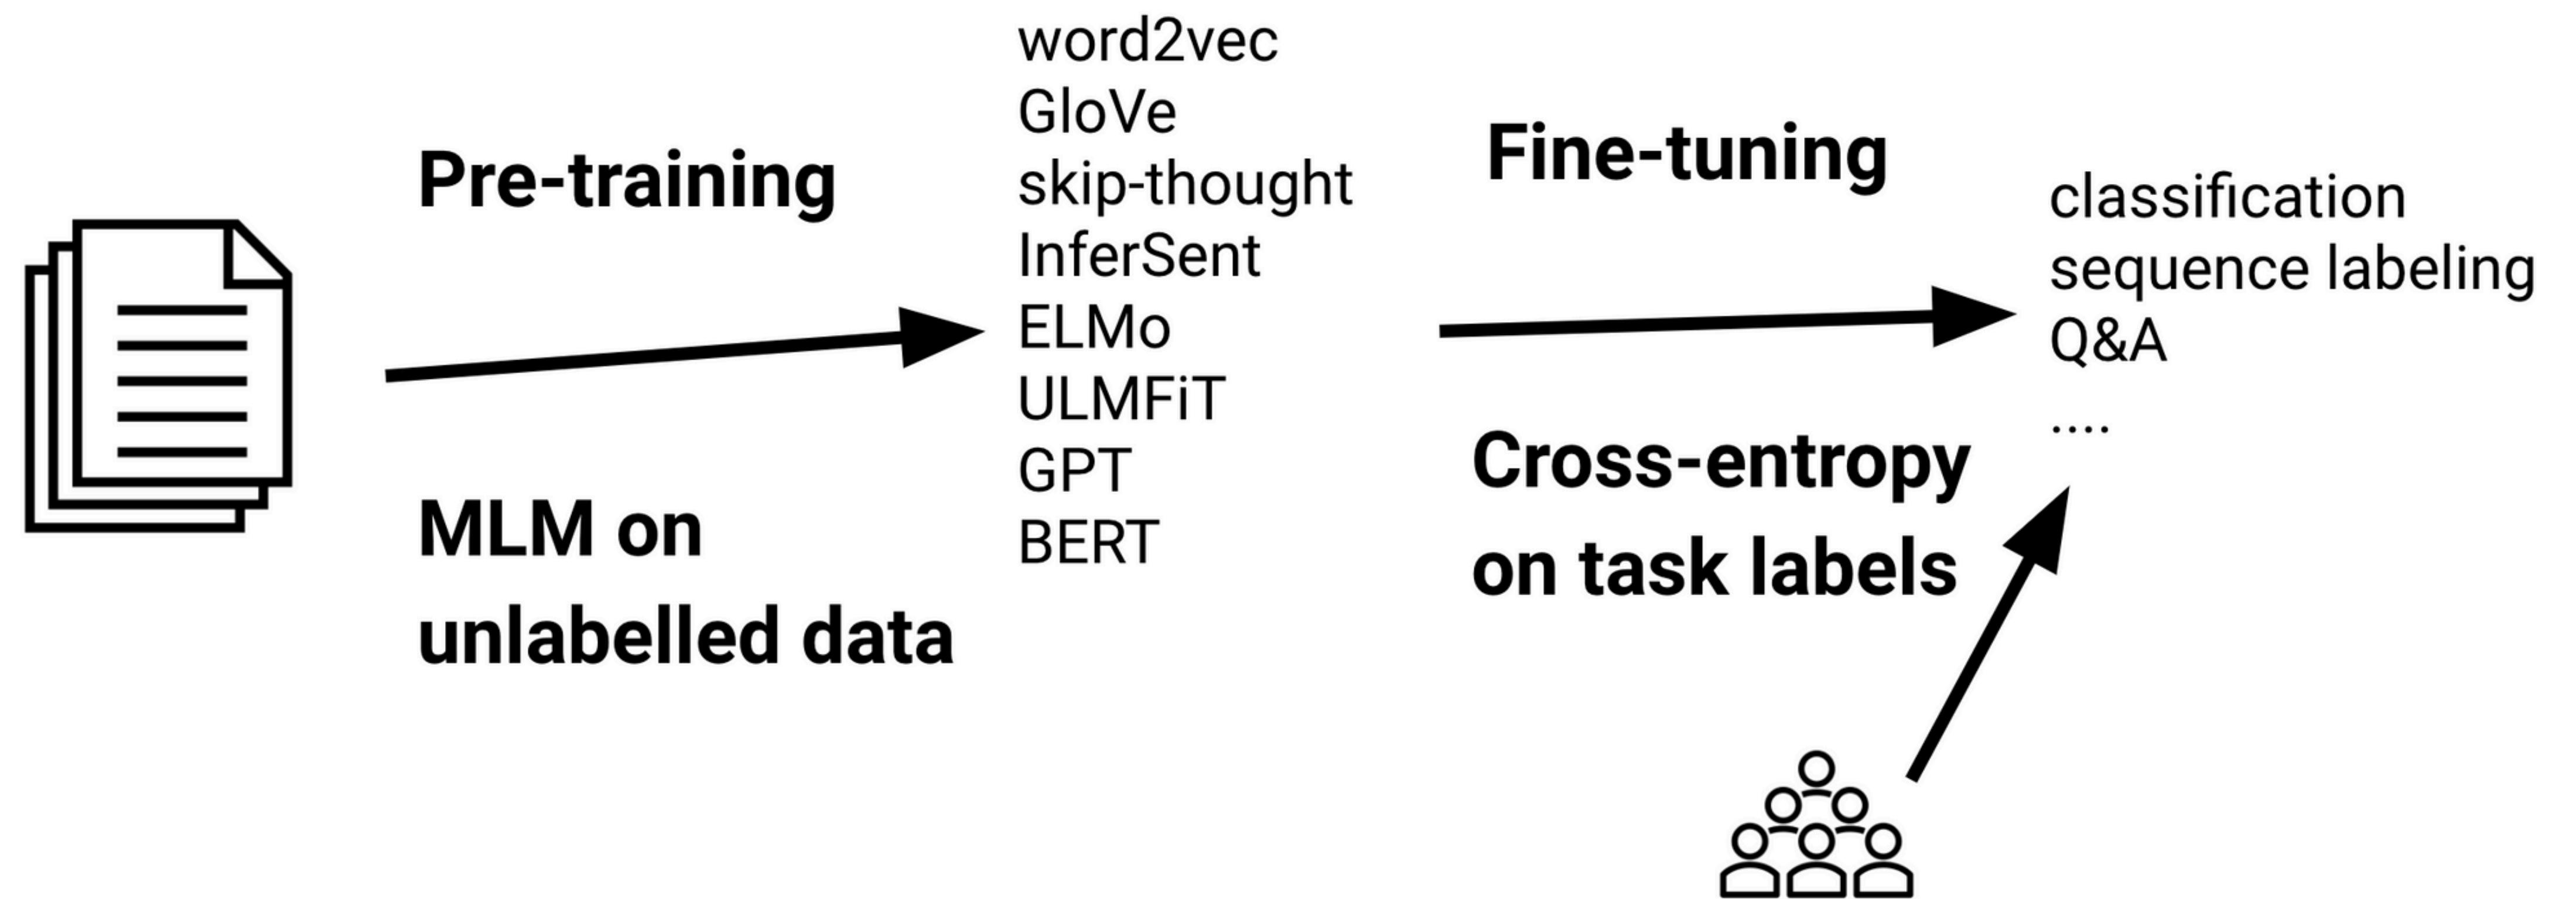
TRANSFER LEARNING

TRAINING FROM SCRATCH

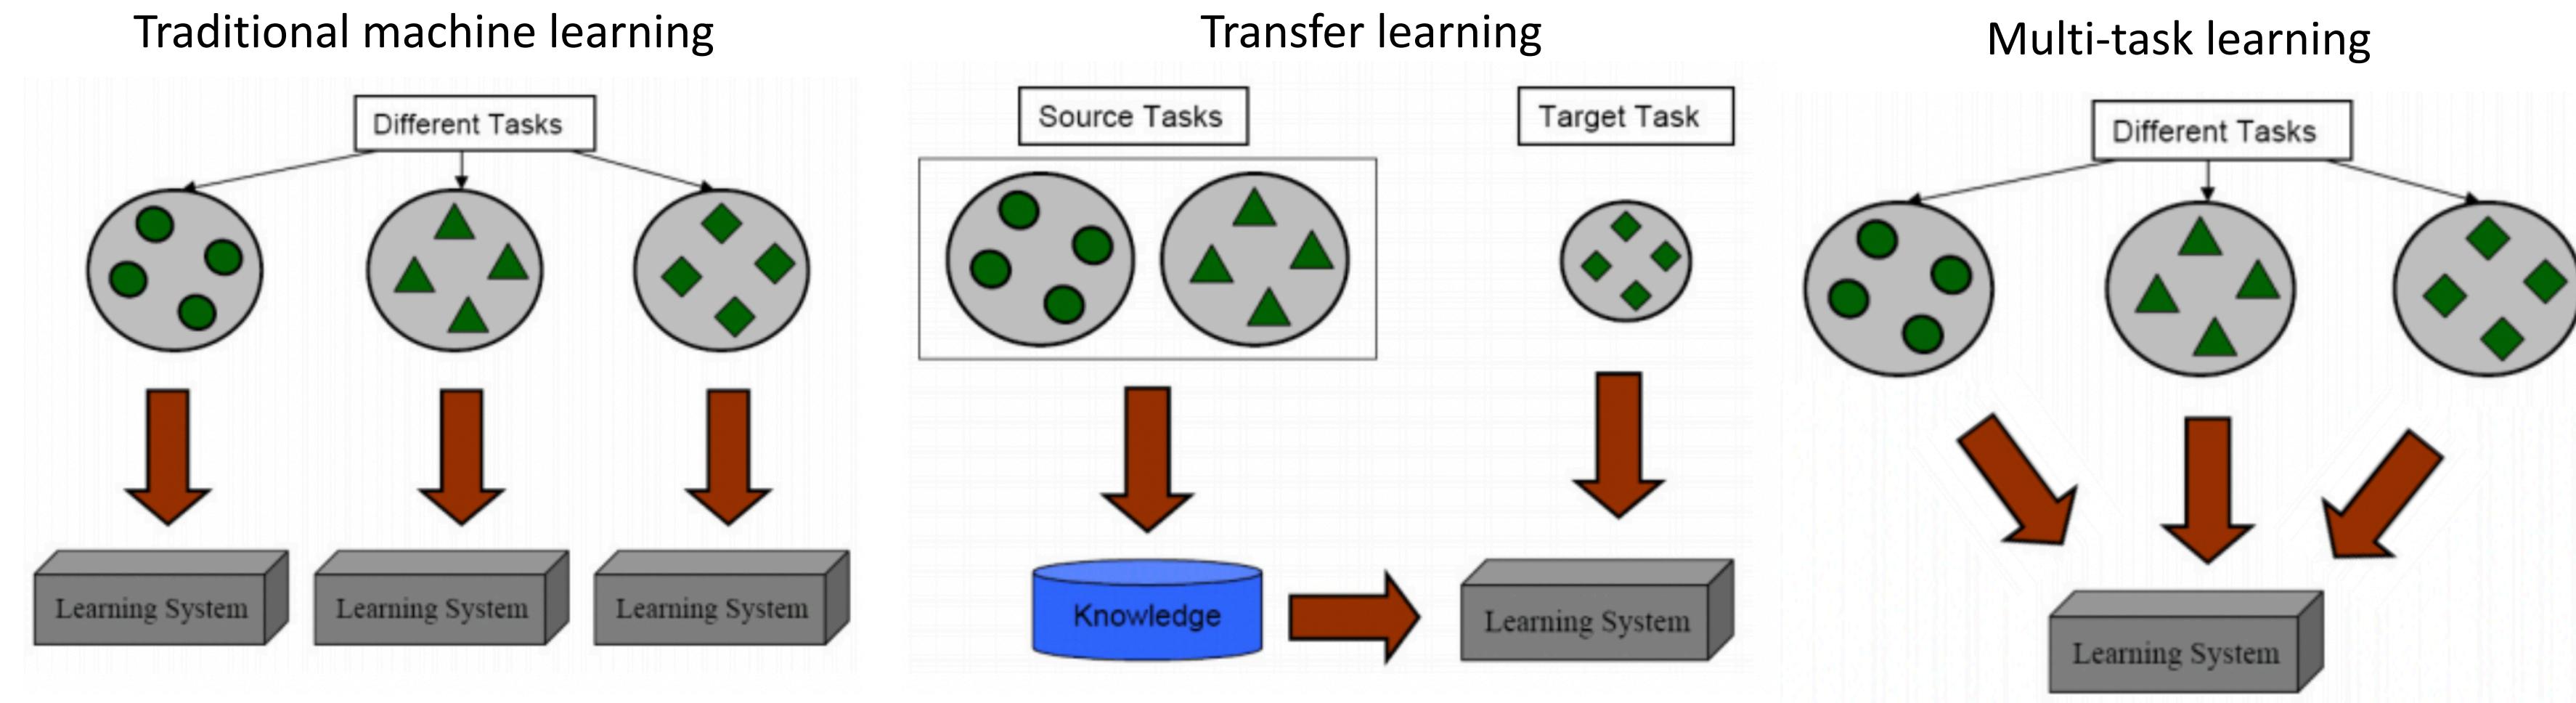


TRANSFER LEARNING



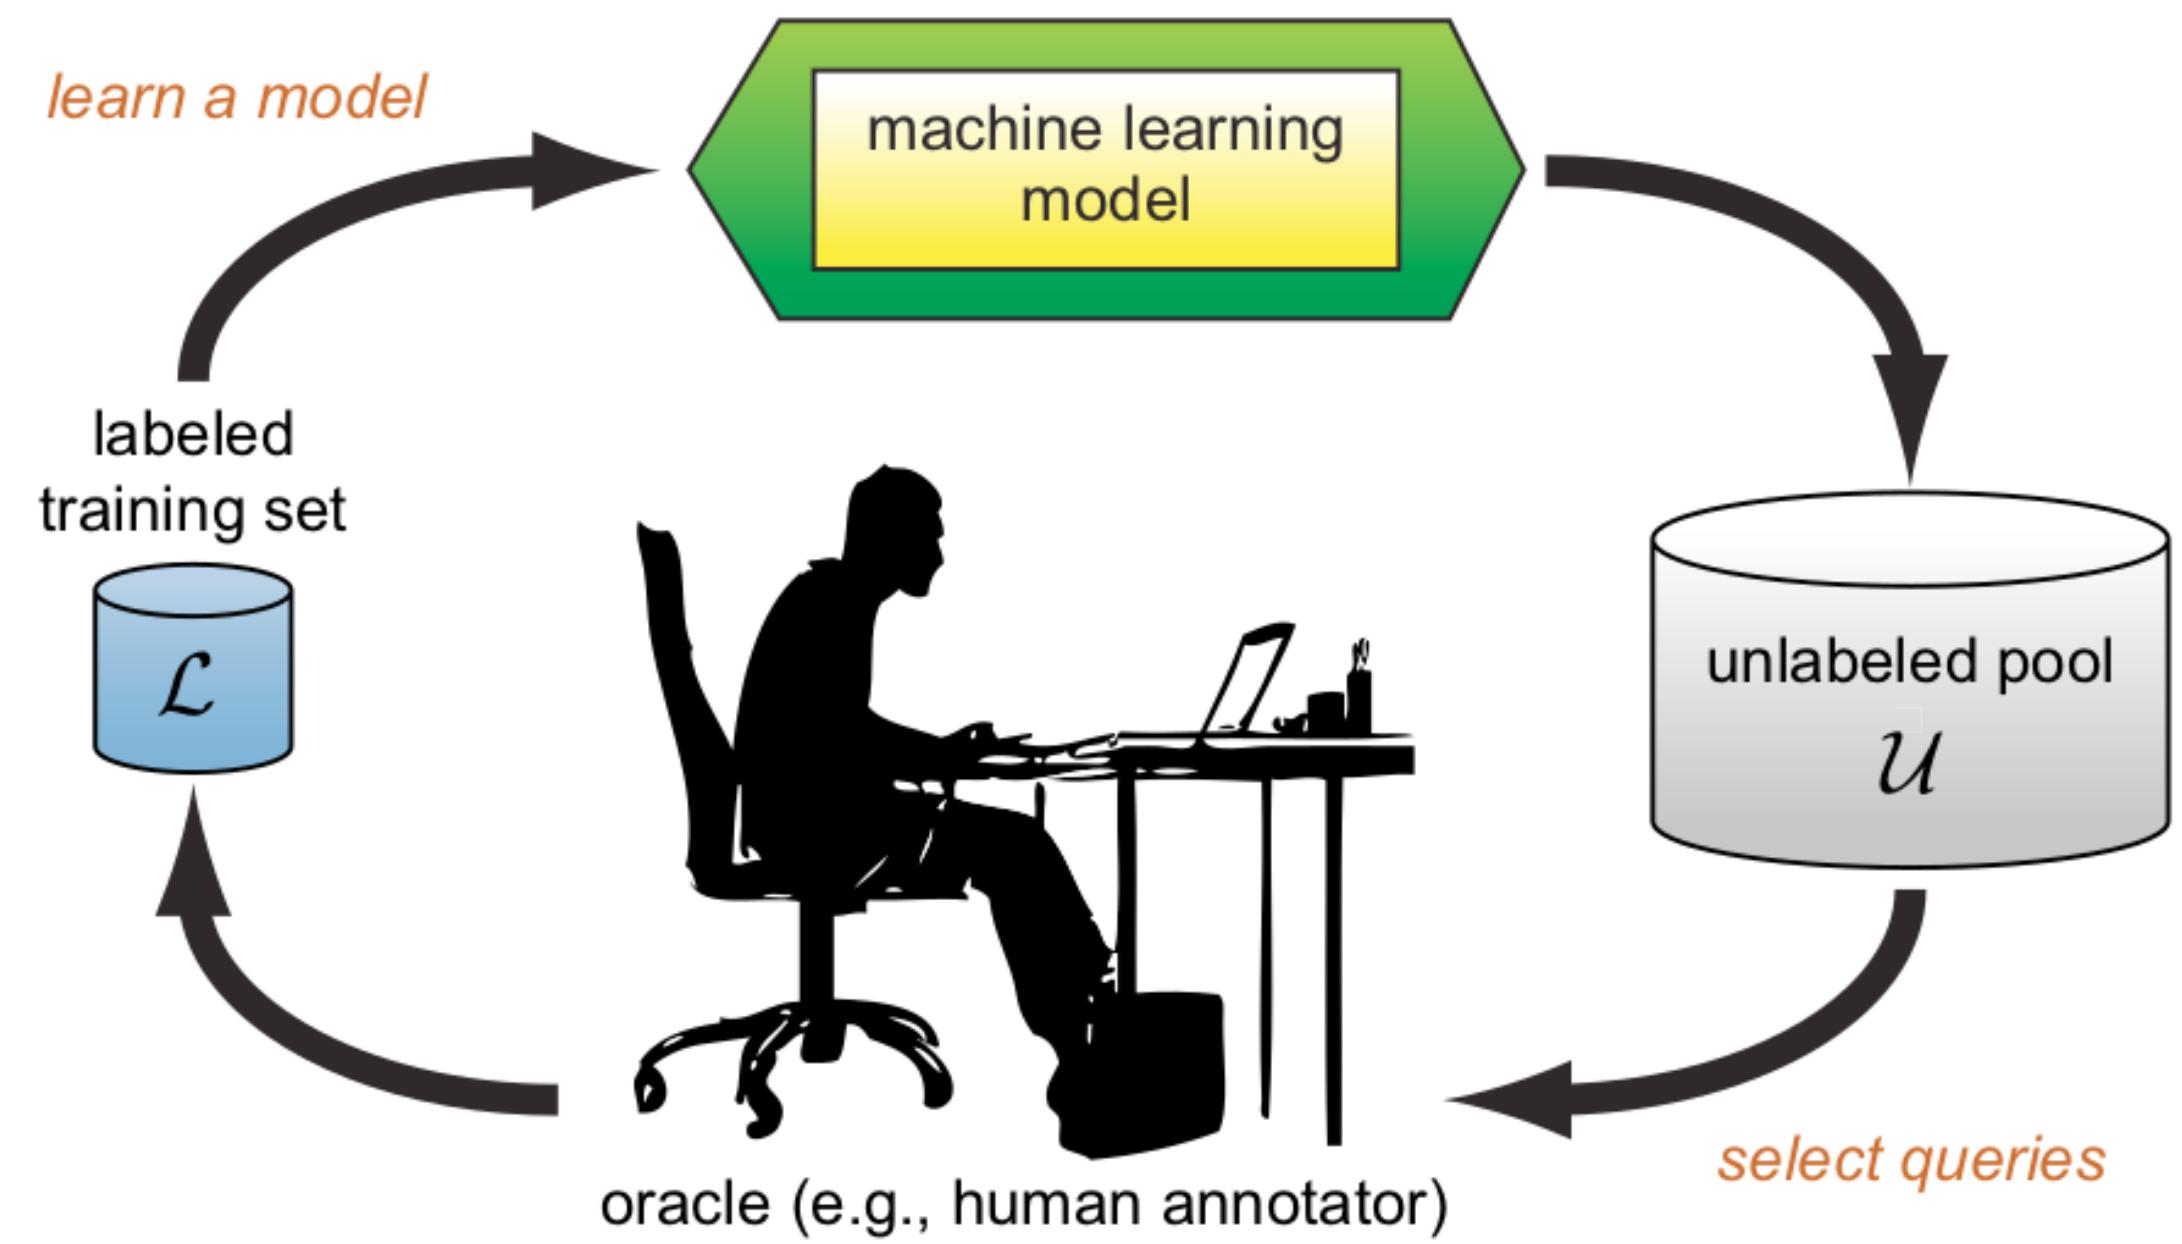


Transfer learning and multi-task learning



ACTIVE LEARNING

Active Learning, aka Query Learning



- Repeat
1. Choose unlabeled sample
 2. Annotate the chosen unlabeled sample
 3. The model trains on the labeled data set

Cheap unlabeled data
Expensive labeled data

48

Fig. The active learning cycle.

Cool, so... what's next?

The screenshot shows the scikit-learn documentation page for Support Vector Machines (SVMs). The top navigation bar includes links for Install, User Guide, API, Examples, Community, and More. The main content area features a section titled "1.4. Support Vector Machines". It describes SVMs as a set of supervised learning methods used for classification, regression, and outliers detection. It lists advantages such as effectiveness in high-dimensional spaces and memory efficiency. Disadvantages include over-fitting if features exceed samples and lack of probability estimates. A note mentions support for both dense and sparse input vectors. A sub-section "1.4.1. Classification" is shown, stating that SVC, NuSVC, and LinearSVC classes handle binary and multi-class classification. A "Toggle Menu" button is at the bottom.

This screenshot shows a different version of the SVM documentation. The top navigation bar is identical. The main content area has a slightly different layout, but the core information about SVMs, their advantages, disadvantages, and classification capabilities is present. A "Toggle Menu" button is at the bottom.

SVC and NuSVC are similar methods, but accept slightly different sets of parameters and have different mathematical formulations (see section Mathematical formulation). On the other hand, LinearSVC is another (faster) implementation of Support Vector Classification for the case of a linear kernel. Note that LinearSVC does not accept parameter kernel, as this is assumed to be linear. It also lacks some of the attributes of SVC and NuSVC, like support_.

As other classifiers, SVC, NuSVC and LinearSVC take as input two arrays: an array X of shape (n_samples, n_features) holding the training samples, and an array y of class labels (strings or integers), of shape (n_samples):

```
>>> from sklearn import svm  
>>> X = [[0, 0], [1, 1]]  
>>> y = [0, 1]  
>>> clf = svm.SVC()  
>>> clf.fit(X, y)  
SVC()
```

After being fitted, the model can then be used to predict new values:

```
>>> clf.predict([[2., 2.]])  
array([1])
```

SVMs decision function (detailed in the Mathematical formulation) depends on some subset of the training data, called the support vectors. Some properties of these support vectors can be found in attributes support_vectors_, support_, and n_support_:

```
>>> # get support vectors  
>>> clf.support_vectors_  
array([[0., 0.],  
       [1., 1.]])  
>>> # get indices of support vectors  
>>> clf.support_  
array([0, 1...])  
>>> # get number of support vectors for each class  
>>> clf.n_support_  
array([1, 1]...)
```

no seriously,
learn to read
docs



scikit-learn.org
...and general google-fu!

Cool, so... what's next?

coursera Explore ▾ What do you want to learn? 

Browse > Data Science > Machine Learning

Machine Learning Specialization

#BreakIntoAI with Machine Learning Specialization. Master fundamental AI concepts and develop practical machine learning skills in the beginner-friendly, 3-course program by AI visionary Andrew Ng

★★★★★ 4.9 9,029 ratings

 Andrew Ng +3 more instructors [TOP INSTRUCTORS](#)

[Enroll for Free](#) Starts Mar 19

Financial aid available

155,216 already enrolled



StatQuest with Josh Starmer ✓

@statquest 897K subscribers 245 videos

Statistics, Machine Learning and Data Science can sometimes be hard to understand. That's where StatQuest comes in.

towards data science

coursera Explore ▾ What do you want to learn? 

Browse > Data Science > Machine Learning

Deep Learning Specialization

Become a Machine Learning expert. Master the fundamentals of deep learning and break into AI. Recently updated with cutting-edge techniques!

★★★★★ 4.9 128,789 ratings

 Andrew Ng +2 more instructors [TOP INSTRUCTORS](#)

[Enroll for Free](#) Starts Mar 19

Financial aid available

747,307 already enrolled



**With great power
comes great responsibility.**

(and have lots of fun!)

Ethics

- Interpretability of the model
- Security
- Privacy
- Trustworthiness
- Biases
- Fairness
- Socioeconomic consequences
- Unintended consequences
- Trolley problem

?????



Thank you.





AIBridge

Samuel Ren, Henry M Gunn High School

Jailing Situ Homestead High School

Houjun Liu, The Nueva School

Xin Liu, Professor, Computer Science. UC Davis

With acknowledgments to

AI Institute for Food Systems

Saratoga Public Library